

ORB SLAM2 CUDA Guide

Written by Eran Moshe

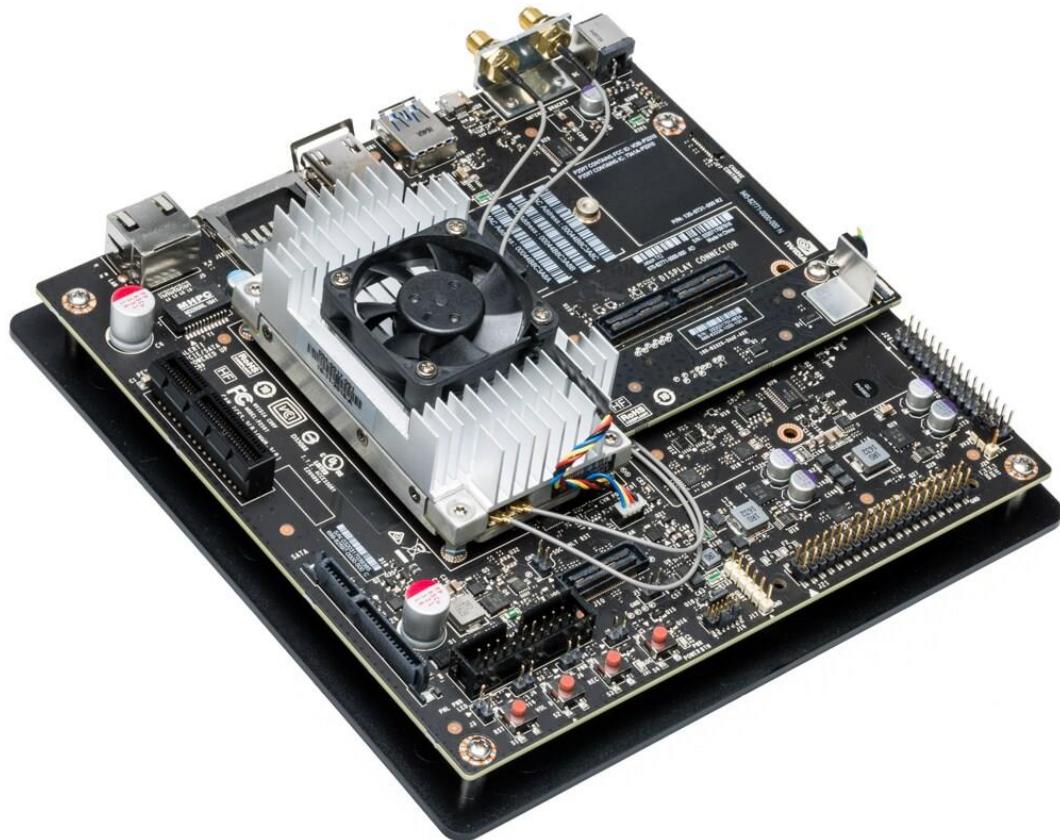
Supervision Yehonatan Mendel

[Autonomous Drones Lab](#)

Faculty of Electrical Engineering, Tel Aviv University

Version 3.0

25.05.2020



Project Description

Build CUDA Infrastructure for Parallel Programming on Nvidia Jetson TX2

In many applications in embedded devices there are more and more interactions between GPU and CPU to accelerate processes, CUDA allows to create efficiently and quickly parallel programs on GPU's (device side) which controlled by CPU (Host side). As the field of AI grows, there is a growing demand for autonomous surveillance by drones. Low latency is very important requirement for drone detection and tracking which can be reached by parallelism of operations.

The main task in the project is to build lab Cuda infrastructure to execute parallel code on Nvidia Jetson TX2. Jetson TX2 is the fastest, most power-efficient embedded AI computing device. It's built around an NVIDIA Pascal™-family GPU and loaded with 8GB of memory and 59.7GB/s of memory bandwidth.

Project stages are:

- Internet Research - Overview.
- Learning Architecture and System.
- Learning Cuda Programming.
- Ramp up Jetson TX2 Environment for Running ORB SLAM Application in Linux Environment.
- Research about Nvidia Profiling Tools to Debug Cuda Applications.

This project will help in the future to students and people which interesting in Parallel Programming study Cuda programming in easy and fast way then could implement, improve and optimize parallel applications. By using Nvidia profiling tools students and people could understand how to debug in better way application which run on CPU and especially on GPU, understand more deeply about the Parallel concepts and how to improve the Performance.



Table of Contents

Table of Contents	2
Acronyms	5
Jetson TX2 Hardware Specifications	6
Jetson Software	7
ORB SLAM Repositories	8
Sdcard Configuration	9
ORB SLAM Installation steps	10
Verify Required Tools	10
OpenCV Installation	13
ROS Installation	15
Pangolin	16
BLAS	16
LAPACK	16
Eigen3	17
Build SLAM	18
Solutions to porting issues with OpenCV 4.1	19
Performance Tip commands	26
Run SLAM application on the target	27
Run ORB SLAM with ROS	30
Run ORB SLAM without ROS	32
Configuration issue during SLAM running	34
ROS Useful commands	35
CUDA	36
Cuda Compilation Flow	37
Communication between Host and Device	39
Cuda API for transfer data	40
Cuda API for synchronization	41
ORB SLAM GPU Cuda Code	42
Main Kernel Cuda functions	43
tileCalcKeypoints_kernel()	43
calcOrb_kernel()	46
Profiling Tools	47
Comparisons Table	48
How to Install JetPack	49
SLAM APPLICATION RUNNING MODES	50
Nvprof (Nvidia profiling tool)	51
Visual profiler	54
Nsight eclipse edition	60
Network Configuration for Remote Profiling	61
GPU PCIe Configuration Check	64
How to run device SLAM application via SSH connection	66

Solutions for issues during SLAM ROS application	68
Compilation Cuda Application with Debug Symbols	70
Nsight System.....	71
SLAM application in ROS mode	71
SLAM application in Standalone Mode	77
Analyze the SLAM CUDA application with Nsight system	82
Nsight Compute	86
SLAM application in Standalone mode	86
Nsight Graphics	93
SLAM application in Standalone mode	93
Extensions	95
Enable root login over SSH.....	95

Acronyms

SLAM	Simultaneous Localization And Mapping
GPU	Graphics Processing Unit
ROS	Robots Operation System
BLAS	Basic Linear Algebra Subprograms
LAPACK	Linear Algebra Package
AI	Artificial Intelligence
FLOPS	Floating Point Operations
SDK	Software development Kit
L4T	NVidia Linux Driver Package
API	Application Programming Interface
PCL	Point Cloud Library

Jetson TX2 Hardware Specifications

Jetson TX2 is a 7.5-watt supercomputer on a module that brings true AI computing at the edge. It's built around an NVIDIA Pascal™-family GPU and loaded with 8 GB of memory and 59.7 GB/s of memory bandwidth.

Main specifications:

GPU - NVIDIA Pascal™ architecture with 256 NVIDIA CUDA cores 1.3 TFLOPS (FP16).

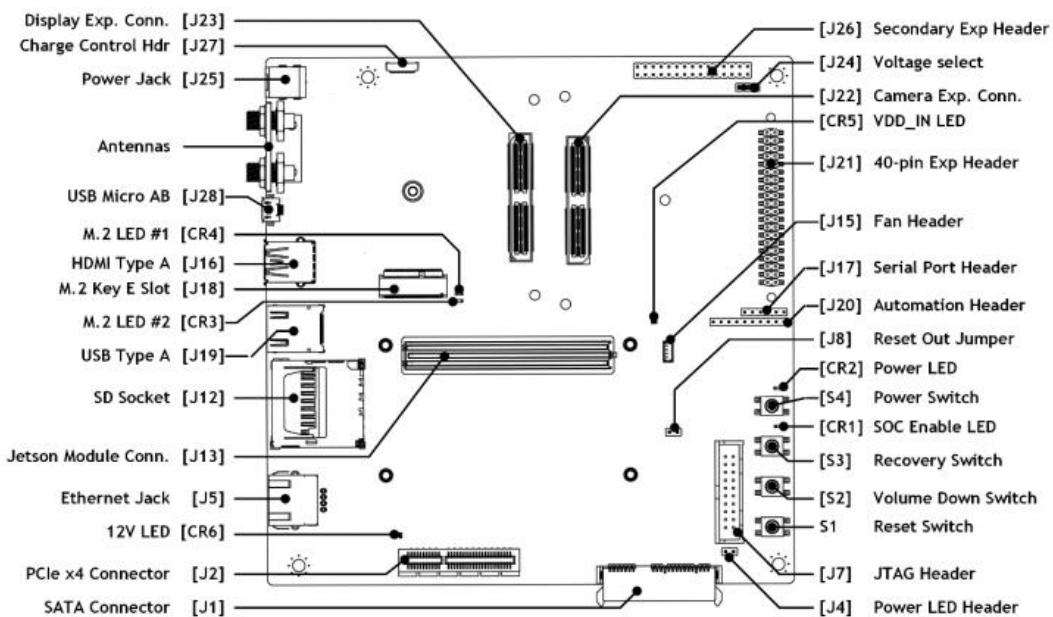
CPU - Dual-core Denver 2 64-bit CPU and quad-core ARM A57 complex

Memory - 8 GB 128-bit LPDDR4 1866MHz - 59.7 GB/s

Storage - 32 GB eMMC 5.1

Connectivity - 10/100/1000 BASE-T Ethernet

Top view of developer kit carrier board (revision C02)



HW Important link: https://elinux.org/Jetson_TX2

Jetson Software

Jetson default OS configuration is Linux Ubuntu.

JetPack bundles all of the Jetson platform software, starting with the NVIDIA® Jetson™ Linux Driver Package (L4T). L4T provides the Linux kernel, bootloader, NVIDIA drivers, flashing utilities, sample filesystem, and more for the Jetson platform.

NVIDIA JetPack SDK is a comprehensive resource for building AI applications. It includes L4T together with accelerated software libraries, APIs, sample applications, developer tools, and documentation.

NVIDIA L4T Supported features:

- Kernel version 4.9
- Support for 64-bit user space and runtime libraries
- Vulkan Support
- V4L2 media-controller driver support for camera sensors (bypassing ISP)
- libargus provides low-level frame-synchronous API for camera applications
 - RAW output CSI cameras needing ISP can be used with either libargus or GStreamer plugin
- Media APIs:
 - OpenGL 4.6 Beta
 - OpenGL ES 3.2
 - OpenGL ES path extensions
 - EGL 1.5 with EGLImage
- X Resize, Rotate and Reflect Extension (RandR) 1.4
- X11 Support

SW Important link: <https://developer.nvidia.com/embedded/develop/software>

L4T link: <https://docs.nvidia.com/jetson/l4t/index.html>

JetPack Roadmap link: <https://developer.nvidia.com/embedded/develop/roadmap>

ORB SLAM Repositories

Index	Web Page Link	Description
1	https://github.com/hoangthien94/ORB_SLAM2_CUDA	Optimized for ROS 18-22 fps
2	https://github.com/connorsoohoo/ORB-SLAM2-GPU-RGBD	Optimized for RGBD 18-20 fps
3	https://github.com/yunchih/ORB-SLAM2-GPU2016-final	GPU enhancements with Cuda ~14 fps
4	https://github.com/raulmur/ORB_SLAM2	Original ORB SLAM

In this document I used the first one which including Cuda Optimizations.

Sdcard Configuration

Jetson TX2 contain eMMC chip with storage capacity of 32 GB.
It is recommended to use external SD card for expand the memory.

For new Sdcard install the exfat support:

```
sudo apt-get install exfat-fuse exfat-utils
```

Then reinsert the card.

If it does not mount automatically, try:

manual mount:

```
sudo mount.exfat /dev/mmcblk0p1 /your/mount/point
```

ORB SLAM Installation steps

This configuration is on the target (Jetson tx2)

Verify Required Tools

- 1) Check cuda version

```
/usr/local/cuda*/bin/nvcc --version
```

Example:

```
yoni3@yoni3-GB-BRT7-H-8550:~/ORB_SLAM2_CUDA$ /usr/local/cuda*/bin/nvcc --version
nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2019 NVIDIA Corporation
Built on Mon_Mar_11_22:07:30_CDT_2019
Cuda compilation tools, release 10.0, V10.0.326
yoni3@yoni3-GB-BRT7-H-8550:~/ORB_SLAM2_CUDA$
```

2) Check OpenCV installation/version/configuration (run with python):

```
python
import cv2
print cv2.getBuildInformation()
```

Example:

```
>>> import cv2
>>>
>>>
>>> print cv2.getBuildInformation()

General configuration for OpenCV 3.3.1 =====
Version control: 3.3.1-2-g31ccdfc11

Platform:
Timestamp: 2019-02-06T09:45:58Z
Host: Linux 4.9.140-tegra aarch64
CMake: 2.8.12.2
CMake generator: Unix Makefiles
CMake build tool: /usr/bin/make
Configuration: Release

CPU/HW features:
Baseline: NEON FP16
required: NEON
disabled: VFPV3

C/C++:
Built as dynamic libs?: YES
C++11: YES
C++ Compiler: /usr/bin/c++ (ver 7.3.0)
C++ flags (Release): -fsigned-char -W -Wall -Werror=return-type -Werror=non-virtual-dtor -Werror=address -Werror=sequence-point -Wformat -Werror=format-security -Wmissing-declarations -Wundef -Winit-self -Wpointer-arith -Wshadow -Wsign-promo -Wuninitialized -Winit-self -Wno-narrowing -Wno-delete-non-virtual-dtor -Wno-comment -Wno-implicit-fallthrough -fdiagnostics-show-option -pthread -fomit-frame-pointer -ffunction-sections -fvisibility-hidden -fvisibility-inlines-hidden -O3 -DNDEBUG -DNDEBUG
C++ flags (Debug): -fsigned-char -W -Wall -Werror=return-type -Werror=non-virtual-dtor -Werror=address -Werror=sequence-point -Wformat -Werror=format-security -Wmissing-declarations -Wundef -Winit-self -Wpointer-arith -Wshadow -Wsign-promo -Wuninitialized -Winit-self -Wno-narrowing -Wno-delete-non-virtual-dtor -Wno-comment -Wno-implicit-fallthrough -fdiagnostics-show-option -pthread -fomit-frame-pointer -ffunction-sections -fvisibility-hidden -fvisibility-inlines-hidden -g -O0 -DDEBUG -D_DEBUG
C Compiler: /usr/bin/cc
C flags (Release): -fsigned-char -W -Wall -Werror=return-type -Werror=address -Werror=sequence-point -Wformat -Werror=format-security -Wmissing-declarations -Wmissing-prototypes -Wstrict-prototypes -Wundef -Winit-self -Wpointer-arith -Wshadow -Wuninitialized -Winit-self -Wno-narrowing -Wno-comment -Wno-implicit-fallthrough -fdiagnostics-show-option -pthread -fomit-frame-pointer -ffunction-sections
C flags (Debug): -fsigned-char -W -Wall -Werror=return-type -Werror=address -Werror=sequence-point -Wformat -Werror=format-security -Wmissing-declarations -Wmissing-prototypes -Wstrict-prototypes -Wundef -Winit-self -Wpointer-arith -Wshadow -Wuninitialized -Winit-self -Wno-narrowing -Wno-comment -Wno-implicit-fallthrough -fdiagnostics-show-option -pthread -fomit-frame-pointer -ffunction-sections -fvisibility-hidden -g -O0 -DDEBUG -D_DEBUG
Linker flags (Release):
Linker flags (Debug):
ccache: NO
Precompiled headers: NO
Extra dependencies: dl m pthread rt /usr/lib/aarch64-linux-gnu/libtbb.so
3rdparty dependencies:

OpenCV modules:
To be built: core flann imgproc ml objdetect photo video dnn imgcodecs shape videoio highgui superres ts features2d calib3d stitching videostab python2 python3
Disabled: js world
Disabled by dependency: -
Unavailable: cudaarithm cudabgsegm cudacodec cudafeatures2d cudafilters cudaimgproc cudalegacy cudaobjdetect cudaoptflow cudastereo cudawarping cuudev java viz

GUI:
QT: NO
GTK+ 2.x: YES (ver 2.24.32)
GThread : YES (ver 2.56.3)
GtkGlExt: NO
OpenGL support: NO
VTK support: NO

Media I/O:
ZLib: /usr/lib/aarch64-linux-gnu/libz.so (ver 1.2.11)
JPEG: /usr/lib/aarch64-linux-gnu/libjpeg.so (ver )
WEBP: build (ver encoder: 0x020e)
PNG: /usr/lib/aarch64-linux-gnu/libpng.so (ver 1.6.34)
TIFF: /usr/lib/aarch64-linux-gnu/libtiff.so (ver 42 - 4.0.9)
JPEG 2000: build (ver 1.900.1)
OpenEXR: NO
GDAL: NO
GDCM: NO

Video I/O:
DC1394 1.x: NO
DC1394 2.x: NO
FFMPEG: YES
avcodec: YES (ver 57.107.100)
avformat: YES (ver 57.83.100)
avutil: YES (ver 55.78.100)
swscale: YES (ver 4.8.100)
avresample: NO
GStreamer:
base: YES (ver 1.14.1)
video: YES (ver 1.14.1)
app: YES (ver 1.14.1)
```

```

riff:           YES (ver 1.14.1)
pbutils:        YES (ver 1.14.1)
OpenNI:          NO
OpenNI PrimeSensor Modules: NO
OpenNI2:         NO
PvAPI:          NO
GigEVisionSDK:  NO
Aravis SDK:     NO
UniCap:         NO
UniCap ucil:   NO/YES
V4L/V4L2:       NO/YES
XIMEA:          NO
Xine:           NO
Intel Media SDK: NO
gPhoto2:         NO

Parallel framework: TBB (ver 2017.0 interface 9107)

Trace:          YES ()

Other third-party libraries:
Use Intel IPP:  NO
Use Intel IPP IW: NO
Use VA:          NO
Use Intel VA-API/OpenCL: NO
Use Lapack:      NO
Use Eigen:       YES (ver 3.3.4)
Use Cuda:        NO
Use OpenCL:      NO
Use OpenVX:      NO
Use custom HAL: YES (carotene (ver 0.0.1))

Python 2:
Interpreter:   /usr/bin/python2.7 (ver 2.7.15)
Libraries:     /usr/lib/aarch64-linux-gnu/libpython2.7.so (ver 2.7.15rc1)
numpy:          /usr/lib/python2.7/dist-packages/numpy/core/include (ver 1.13.3)
packages path: lib/python2.7/dist-packages

Python 3:
Interpreter:   /usr/bin/python3 (ver 3.6.7)
Libraries:     /usr/lib/aarch64-linux-gnu/libpython3.6m.so (ver 3.6.7)
numpy:          /usr/lib/python3/dist-packages/numpy/core/include (ver 1.13.3)
packages path: lib/python3.6/dist-packages

Python (for build): /usr/bin/python2.7

Java:
ant:           NO
JNI:           NO
Java wrappers: NO
Java tests:    NO

Matlab:         Matlab not found or implicitly disabled

Documentation:
Doxygen:       NO

Tests and samples:
Tests:          YES
Performance tests: YES
C/C++ Examples: YES

Install path:   /usr
cvconfig.h is in: /home/nvidia/build_opencv/build
-----
```

OpenCV Installation

This tutorial explain how to install OpenCV version 4.1.1

Check the current version with command:

Example:

```
slamgpu@slamgpu-desktop:~$ opencv_version  
3.3.1
```

To upgrade OpenCV to 4.1.1 run script below:

```
#!/bin/bash  
#  
# Copyright (c) 2018, NVIDIA CORPORATION. All rights reserved.  
#  
# NVIDIA Corporation and its licensors retain all intellectual property  
# and proprietary rights in and to this software, related documentation  
# and any modifications thereto. Any use, reproduction, disclosure or  
# distribution of this software and related documentation without an express  
# license agreement from NVIDIA Corporation is strictly prohibited.  
#  
if [ "$#" -ne 1 ]; then  
    echo "Usage: $0 <Install Folder>"  
    exit  
fi  
folder="$1"  
user="nvidia"  
passwd="nvidia"  
  
echo "*** Remove other OpenCV first"  
sudo sudo apt-get purge "libopencv"  
  
echo "*** Install requirement"  
sudo apt-get update  
sudo apt-get install -y build-essential cmake git libgdk2.0-dev pkg-config libavcodec-dev libavformat-dev libswscale-dev  
sudo apt-get install -y libgstreamer1.0-dev libgstreamer-plugins-base1.0-dev  
sudo apt-get install -y python2.7-dev python3.6-dev python-dev python-numpy python3-numpy  
sudo apt-get install -y libtbb2 libtbb-dev libjpeg-dev libpng-dev libtiff-dev libdc1394-22-dev  
sudo apt-get install -y libv4l-dev v4l-utils qv4l2 v4l2ucp  
sudo apt-get install -y curl  
sudo apt-get update  
  
echo "*** Download opencv-4.1.1"  
cd $folder  
curl -L https://github.com/opencv/opencv/archive/4.1.1.zip -o opencv-4.1.1.zip  
curl -L https://github.com/opencv/opencv_contrib/archive/4.1.1.zip -o opencv_contrib-4.1.1.zip  
unzip opencv-4.1.1.zip  
unzip opencv_contrib-4.1.1.zip  
cd opencv-4.1.1/  
  
echo "*** Apply patch"  
sed -i 's/include <Eigen/Core>/include <eigen3/Eigen/Core>/g' modules/core/include/opencv2/core/private.hpp  
  
echo "*** Building..."  
mkdir release  
cd release/  
cmake -D WITH_CUDA=ON -D CUDA_ARCH_BIN="5.3,6.2,7.2" -D CUDA_ARCH_PTX="" -D OPENCV_EXTRA_MODULES_PATH=../../opencv_contrib-4.1.1/modules -D WITH_GSTREAMER=ON -D WITH_LIBV4L=ON -D BUILD_opencv_python2=ON -D BUILD_opencv_python3=ON -D BUILD_TESTS=OFF -D BUILD_PERF_TESTS=OFF -D BUILD_EXAMPLES=OFF -D CMAKE_BUILD_TYPE=RELEASE -D CMAKE_INSTALL_PREFIX=/usr/local ..  
make -j3  
sudo make install  
echo 'export PYTHONPATH=$PYTHONPATH:$PWD/python_loader/' >> ~/.bashrc  
source ~/.bashrc  
  
echo "*** Install opencv-4.1.1 successfully"  
echo "*** Bye. :)"
```

with folder to download and compile the OpenCV package.

You can go to the repository updated script for installing openCV on jetsons:

<https://github.com/AastaNV/JEP/tree/master/script>

Note that OpenCV installation can take several hours.

Tip: You can edit the script to build in parallel with **make -j 4**

In case of compilation errors, Try run installation again to continue, or switch back to make without parallelism.

ROS Installation

ROS provides libraries and tools to help software developers create robot applications. It provides hardware abstraction, device drivers, libraries, visualizers, message-passing, package management, and more.

[Installation guide](http://wiki.ros.org/melodic/Installation/Ubuntu): <http://wiki.ros.org/melodic/Installation/Ubuntu>

Command for installation

```
sudo apt-get install devscripts build-essential
```

```
sudo apt install ros-melodic-desktop
```

```
sudo apt-get source ros-kinetic-opencv3
```

```
sudo echo "deb-src http://packages.ros.org/ros/ubuntu xenial main" >>
/etc/apt/sources.list.d/ros-Latest.list
```

```
sudo apt-get update
```

```
sudo apt-get source ros-kinetic-opencv3
```

if you see the error :

```
Err:15 http://packages.ros.org/ros/ubuntu xenial InRelease
```

```
The following signatures couldn't be verified because the public key is not available:
```

```
NO_PUBKEY F42ED6FBAB17C654
```

```
Reading package lists... Done
```

```
W: GPG error: http://packages.ros.org/ros/ubuntu xenial InRelease: The following signatures
couldn't be verified because the public key is not available: NO_PUBKEY F42ED6FBAB17C654
```

```
E: The repository 'http://packages.ros.org/ros/ubuntu xenial InRelease' is not signed.
```

```
N: Updating from such a repository can't be done securely, and is therefore disabled by default.
```

```
N: See apt-secure(8) manpage for repository creation and user configuration details.
```

You can fix by run the command:

```
curl http://repo.ros2.org/repos.key | sudo apt-key add -
```

5) Dependancies installtion for ORB-SLAM2

Pangolin

Pangolin is a lightweight portable rapid development library for managing OpenGL display / interaction and abstracting video input.

You can download and install by instructions in the link:

<https://github.com/stevenlovegrove/Pangolin>.

Commands for install with command line:

```
git clone https://github.com/stevenlovegrove/Pangolin.git  
cd Pangolin  
mkdir build  
cd build  
cmake ..  
cmake --build .
```

BLAS

Basic Linear Algebra Subprograms are routines that provide standard building blocks for performing basic vector and matrix operations.

LAPACK

Linear Algebra Package is a standard software library for numerical linear algebra. It provides routines for solving systems of linear equations and linear least squares, eigenvalue problems, and singular value decomposition.

Commands for BLAS and LAPACK installations in Ubuntu:

```
sudo apt-get install libblas-dev  
sudo apt-get install liblapack-dev
```

Eigen3

Eigen is a C++ template library for linear algebra: matrices, vectors, numerical solvers, and related algorithms.

Installation instructions can be found at: <http://eigen.tuxfamily.org> .

For ORB SLAM required at least version 3.1.0.

You can install by the following commands:

```
git clone https://gitlab.com/libeigen/eigen.git  
cd eigen  
mkdir build  
cd build  
cmake ..  
make  
sudo make install
```

PCL for ROS

Point Cloud Library ROS interface stack. PCL-ROS is the preferred bridge for 3D applications involving n-D Point Clouds and 3D geometry processing in ROS.

You can install by run the following commands:

```
sudo apt-get install libopenni2-dev  
sudo apt-get install ros-melodic-pcl-ros  
sudo apt-get install python-vtk - not available
```

Build SLAM

Clone the repo and execute the build script for normal ORB-SLAM2:

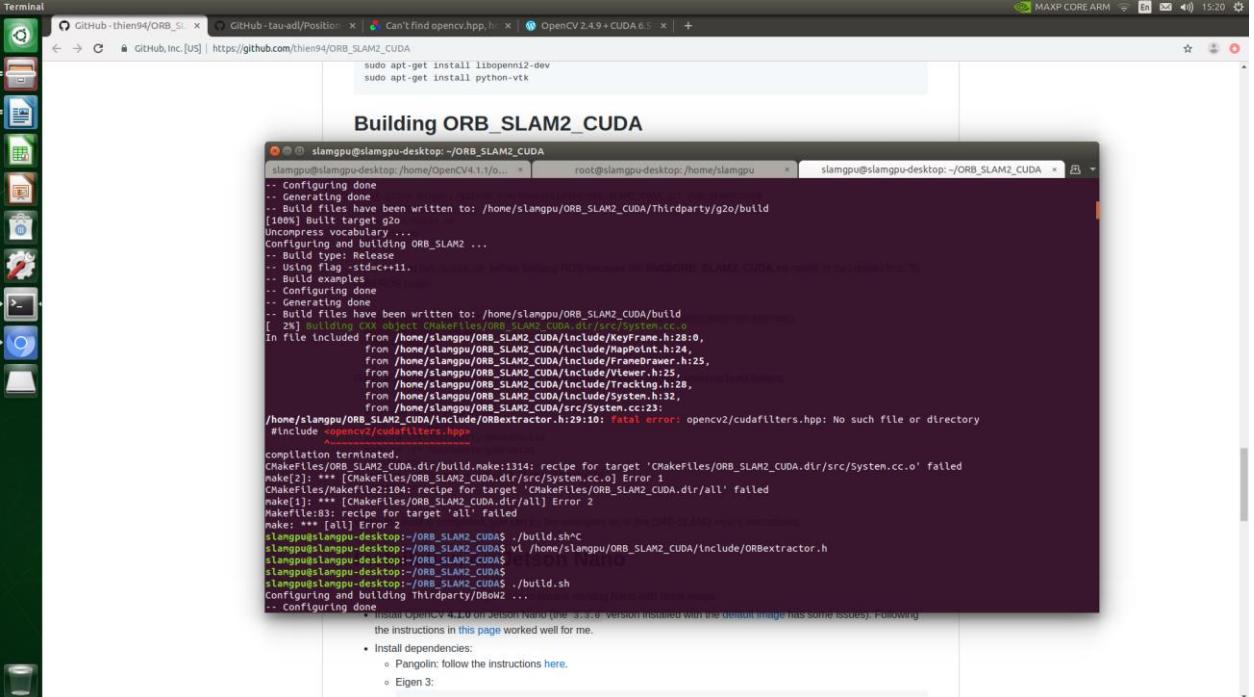
```
git clone https://github.com/hoangthien94/ORB_SLAM2_CUDA.git ORB_SLAM2_CUDA  
cd ORB_SLAM2_CUDA  
chmod +x build.sh  
./build.sh
```

Solutions to porting issues with OpenCV 4.1

1) Issue - Include from OpenCv2 is not recognized.

In the files:

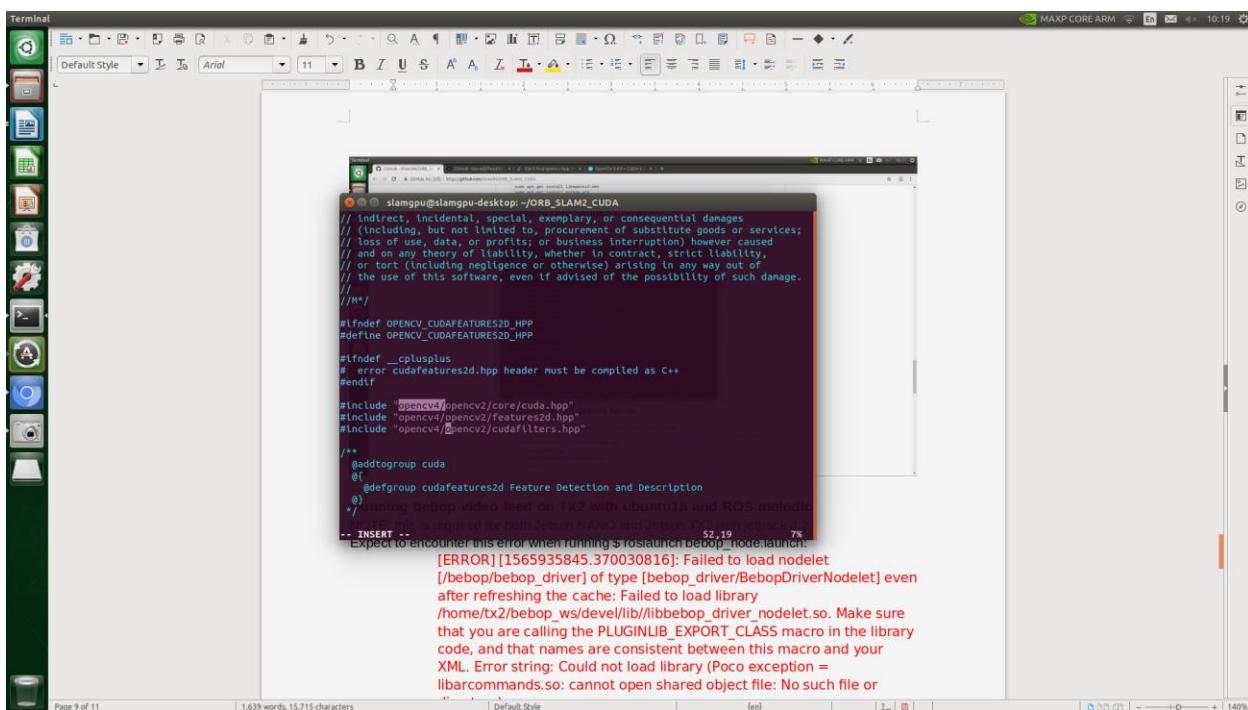
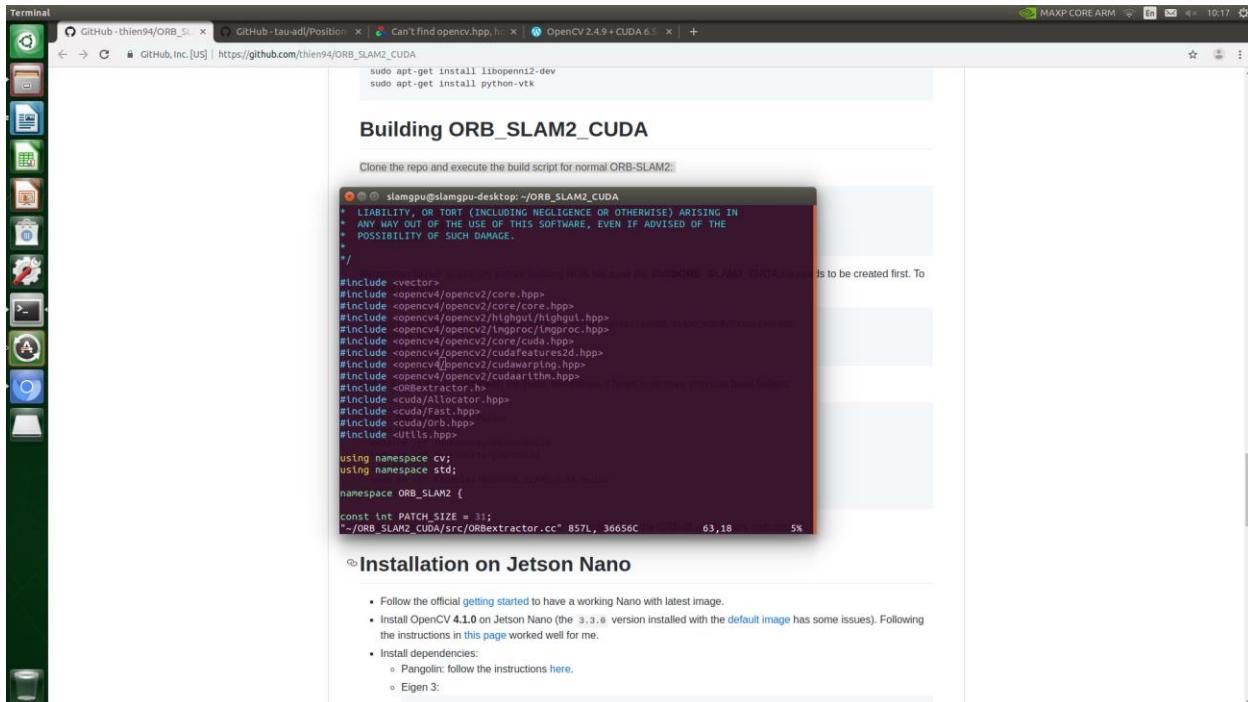
/usr/local/include/opencv4/opencv2/cudafeatures2d.hpp
/home/slamgpu/ORB_SLAM2_CUDA/src/ORBextractor.cc
/home/slamgpu/ORB_SLAM2_CUDA/include/ORBextractor.h



```
sudo apt-get install libopenni2-dev
sudo apt-get install python-vtk

Building ORB_SLAM2_CUDA
slamgpu@slamgpu-desktop:~/ORB_SLAM2_CUDA
-- Configuring done
-- Generating done
-- Build files have been written to: /home/slamgpu/ORB_SLAM2_CUDA/Thirdparty/g2o/build
[100%] Built target g2o
Uncompress vocabulary ...
Configuring and building ORB_SLAM2 ...
-- Build type: Release
-- Using compiler: g++-5
-- Build examples
-- Configuring done
-- Generating done
-- Build files have been written to: /home/slamgpu/ORB_SLAM2_CUDA/build
[ 2%] Building CXX object ORBFiles/ORB_SLAM2_CUDA/src/System.cxx
In file included from /home/slamgpu/ORB_SLAM2_CUDA/include/System.h:28,
                 from /home/slamgpu/ORB_SLAM2_CUDA/include/MapPoint.h:24,
                 from /home/slamgpu/ORB_SLAM2_CUDA/include/FrameDrawer.h:25,
                 from /home/slamgpu/ORB_SLAM2_CUDA/include/Viewer.h:25,
                 from /home/slamgpu/ORB_SLAM2_CUDA/include/Tracking.h:28,
                 from /home/slamgpu/ORB_SLAM2_CUDA/include/System.h:32,
                 from /home/slamgpu/ORB_SLAM2_CUDA/src/System.c:23:
/home/slamgpu/ORB_SLAM2_CUDA/include/ORBextractor.h:29:10: fatal error: opencv2/cudaFilters.hpp: No such file or directory
 #include <opencv2/cudaFilters.hpp>
                                         ^
compilation terminated.
CMakeFiles/ORB_SLAM2_CUDA.dir/build.make:1334: recipe for target 'CMakeFiles/ORB_SLAM2_CUDA.dir/src/System.cc.o' failed
make[1]: *** [CMakeFiles/ORB_SLAM2_CUDA.dir/src/System.cc.o] Error 1
CMakeFiles/Makefile2:104: recipe for target 'CMakeFiles/ORB_SLAM2_CUDA.dir/all' failed
make[1]: *** [CMakeFiles/ORB_SLAM2_CUDA.dir/all] Error 2
Makefile:83: recipe for target 'all' failed
make: *** [all] Error 2
slamgpu@slamgpu-desktop:~/ORB_SLAM2_CUDA$ ./build.sh^C
slamgpu@slamgpu-desktop:~/ORB_SLAM2_CUDA$ vi /home/slamgpu/ORB_SLAM2_CUDA/include/ORBextractor.h
slamgpu@slamgpu-desktop:~/ORB_SLAM2_CUDA$ slamgpu@slamgpu-desktop:~/ORB_SLAM2_CUDA$ ./build.sh
Configuring and building Thirdparty/DBOW2 ...
-- Configuring done
-- Generating done
-- Build files have been written to: /home/slamgpu/ORB_SLAM2_CUDA/Thirdparty/DBOW2
Note: OpenCV 4.1.0 was built with CUDA 9.0. Following the instructions in this page worked well for me.
• Install dependencies:
  • Pangolin: follow the instructions here.
  • Eigen: 3.
```

Solution: Add to include <opencv4/opencv2/cudaFilters.hpp> to all the includes to solve the issue.



2) Issue - library is missing - /usr/bin/ld: cannot find -lInvToolsExt

Solution: run the command

```
sudo ln -s /usr/local/cuda-10.0/lib64/libnvToolsExt.so /usr/lib/libnvToolsExt.so
```

<https://github.com/yunchih/ORB-SLAM2-GPU2016-final/issues/6>

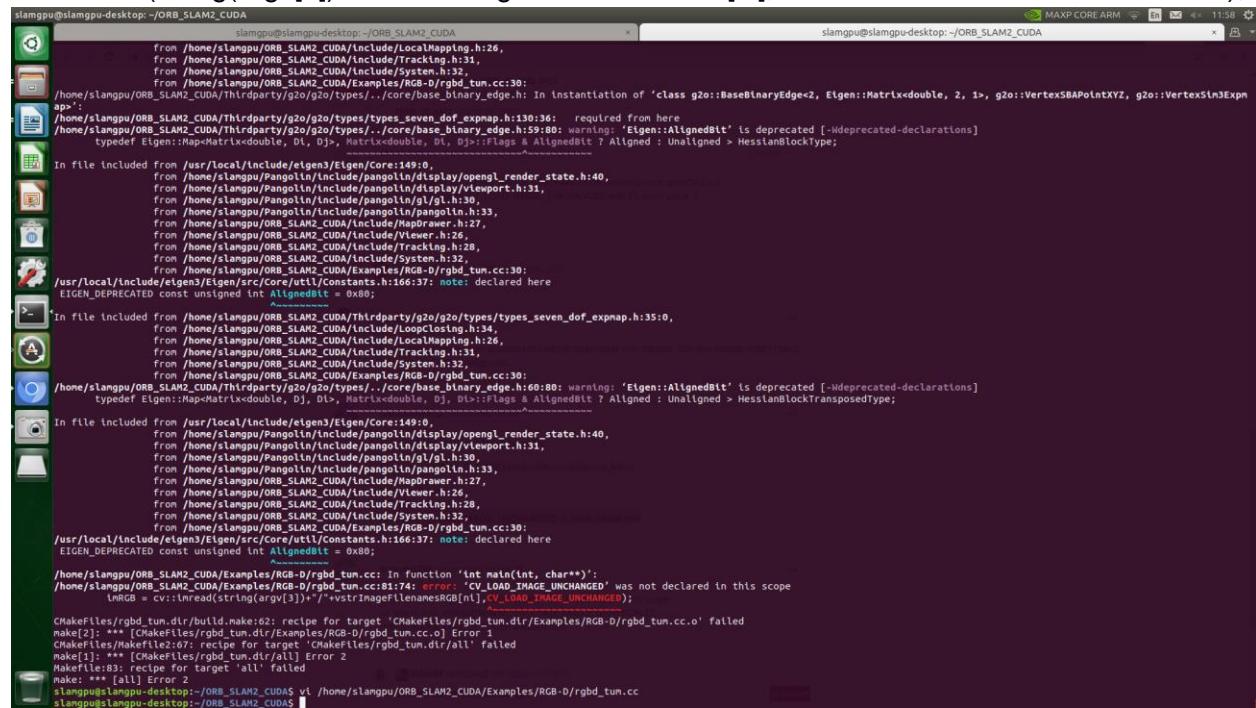
3) Issue /home/slamgpu/ORB_SLAM2_CUDA/Examples/RGB-D/rgbd_tum.cc: In function ‘int main(int, char**)’:

/home/slamgpu/ORB_SLAM2_CUDA/Examples/RGB-D/rgbd_tum.cc:81:74: error:

‘CV_LOAD_IMAGE_UNCHANGED’ was not declared in this scope

imRGB =

```
cv::imread(string(argv[3])+"/*+vstrImageFilenamesRGB[ni],CV_LOAD_IMAGE_UNCHANGED);
```



```
slamgpu@slamgpu-desktop:~/ORB_SLAM2_CUDA
slamgpu@slamgpu-desktop:~/ORB_SLAM2_CUDA
/home/slamgpu/ORB_SLAM2_CUDA/include/LocalMapping.h:26:
From /home/slamgpu/ORB_SLAM2_CUDA/include/Tracking.h:31,
  From /home/slamgpu/ORB_SLAM2_CUDA/include/System.h:32,
    From /home/slamgpu/ORB_SLAM2_CUDA/examples/RGB-D/rgbd_tum.cc:30:
ap>:
/home/slamgpu/ORB_SLAM2_CUDA/thirdparty/g2o/g2o/types/types_seven_dof_expmap.h:30:36: warning: 'Eigen::AlignedBit' is deprecated [-Wdeprecated-declarations]
  typedef Eigen::Map<Matrix<double, 2, 1>, g2o::VertexSBAPointXYZ, g2o::VertexSIn3Exp
In file included from /usr/local/include/eigen3/Eigen/Core:1490,
  From /home/slamgpu/pangolin/include/pangolin/display/openGL_render_state.h:40,
    From /home/slamgpu/pangolin/include/pangolin/display/viewport.h:31,
      From /home/slamgpu/pangolin/include/pangolin/gl/gl.h:30
        From /home/slamgpu/pangolin/include/pangolin/gl/gl.h:30
          From /home/slamgpu/ORB_SLAM2_CUDA/include/MapDrawer.h:27,
            From /home/slamgpu/ORB_SLAM2_CUDA/include/Viewer.h:26,
              From /home/slamgpu/ORB_SLAM2_CUDA/include/Tracking.h:28,
                From /home/slamgpu/ORB_SLAM2_CUDA/include/LocalMapping.h:26,
                  From /home/slamgpu/ORB_SLAM2_CUDA/examples/RGB-D/rgbd_tum.cc:30:
/usr/local/include/eigen3/Eigen/src/Core/util/Constants.h:166:37: note: declared here
EIGEN_DEPRECATED const unsigned int AlignedBit = 0x80;
~~~~~
In file included from /home/slamgpu/ORB_SLAM2_CUDA/thirdparty/g2o/g2o/types/types_seven_dof_expmap.h:35:0,
  From /home/slamgpu/ORB_SLAM2_CUDA/include/LocalMapping.h:26,
    From /home/slamgpu/ORB_SLAM2_CUDA/include/Tracking.h:31,
      From /home/slamgpu/ORB_SLAM2_CUDA/include/System.h:32,
        From /home/slamgpu/ORB_SLAM2_CUDA/examples/RGB-D/rgbd_tum.cc:30:
/home/slamgpu/ORB_SLAM2_CUDA/thirdparty/g2o/g2o/types/core/base_binary_edge.h:68:80: warning: 'Eigen::AlignedBit' is deprecated [-Wdeprecated-declarations]
  typedef Eigen::Map<Matrix<double, 2, 1>, g2o::VertexSBAPointXYZ, g2o::VertexSIn3Exp
In file included from /usr/local/include/eigen3/Eigen/Core:1490,
  From /home/slamgpu/pangolin/include/pangolin/display/openGL_render_state.h:40,
    From /home/slamgpu/pangolin/include/pangolin/display/viewport.h:31,
      From /home/slamgpu/pangolin/include/pangolin/gl/gl.h:30
        From /home/slamgpu/pangolin/include/pangolin/gl/gl.h:30
          From /home/slamgpu/ORB_SLAM2_CUDA/include/MapDrawer.h:27,
            From /home/slamgpu/ORB_SLAM2_CUDA/include/Viewer.h:26,
              From /home/slamgpu/ORB_SLAM2_CUDA/include/Tracking.h:28,
                From /home/slamgpu/ORB_SLAM2_CUDA/include/System.h:32,
                  From /home/slamgpu/ORB_SLAM2_CUDA/examples/RGB-D/rgbd_tum.cc:30:
/usr/local/include/eigen3/Eigen/src/Core/util/Constants.h:166:37: note: declared here
EIGEN_DEPRECATED const unsigned int AlignedBit = 0x80;
~~~~~
/home/slamgpu/ORB_SLAM2_CUDA/examples/RGB-D/rgbd_tum.cc: In function ‘int main(int, char**)’:
/home/slamgpu/ORB_SLAM2_CUDA/examples/RGB-D/rgbd_tum.cc:81:74: error: ‘CV_LOAD_IMAGE_UNCHANGED’ was not declared in this scope
  imRGB = cv::imread(string(argv[3])+"/*+vstrImageFilenamesRGB[ni],CV_LOAD_IMAGE_UNCHANGED);

CMakeFiles/rgbd_tum.dir/build.make:62: recipe for target 'CMakeFiles/rgbd_tum.dir/Examples/RGB-D/rgbd_tum.cc.o' failed
make[2]: *** [CMakeFiles/rgbd_tum.dir/Examples/RGB-D/rgbd_tum.cc.o] Error 1
CMakeFiles/rgbd_tum.dir/build.make:67: recipe for target 'CMakeFiles/rgbd_tum.dir/all' failed
make[1]: *** [CMakeFiles/rgbd_tum.dir/all] Error 2
Makefile:33: recipe for target 'all' failed
make: *** [all] Error 2
slamgpu@slamgpu-desktop:~/ORB_SLAM2_CUDA$ vi /home/slamgpu/ORB_SLAM2_CUDA/Examples/RGB-D/rgbd_tum.cc
slamgpu@slamgpu-desktop:~/ORB_SLAM2_CUDA$
```

Solution:

Add to example cc file the following define

```
#define CV_LOAD_IMAGE_UNCHANGED -1 //new added one
```

https://github.com/raulmur/ORB_SLAM2/issues/451

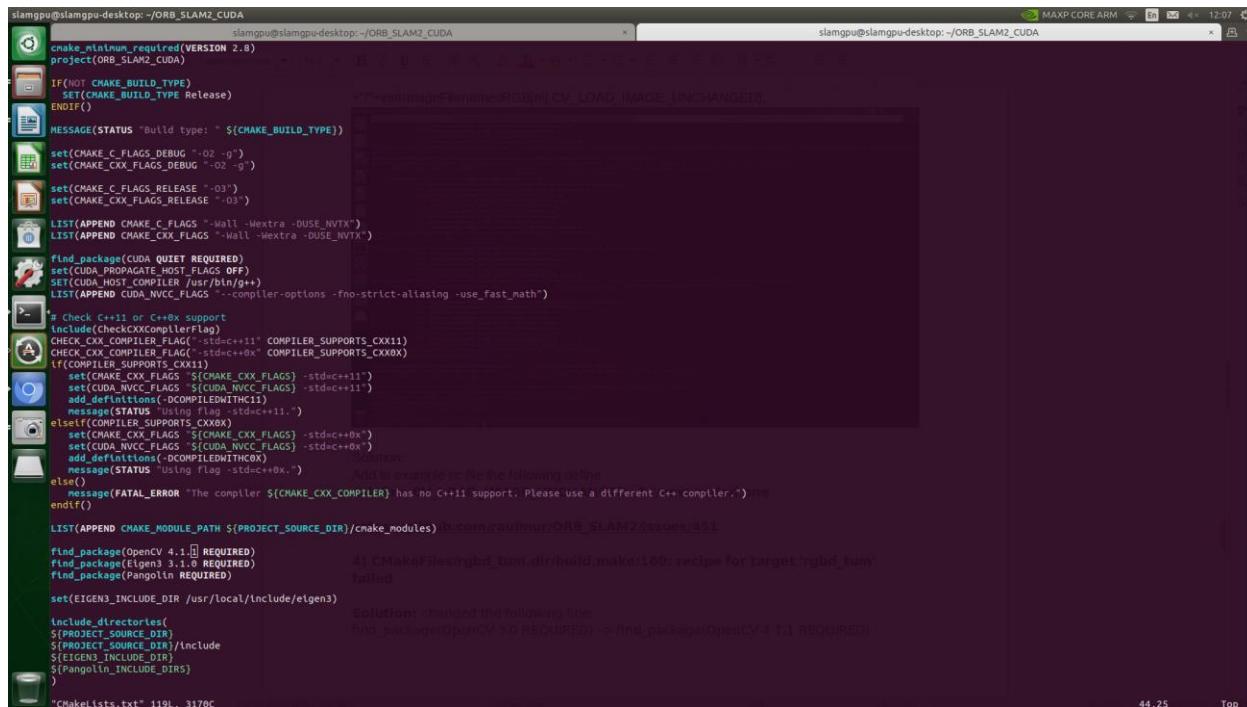
4) Issue - CMakeFiles/rgbd_tum.dir/build.make:160: recipe for target 'rgbd_tum' failed

Solution: changed the following line:

find_package(OpenCV 3.0 REQUIRED) -> find_package(OpenCV 4.1.1 REQUIRED)

in the file CMakeLists.txt

https://github.com/raulmur/ORB_SLAM2/issues/209



```
slamgpu@slamgpu-desktop: ~/ORB_SLAM2_CUDA
slamgpu@slamgpu-desktop: ~/ORB_SLAM2_CUDA
-- Configuring incomplete, errors occurred!
-- 
-- Solution: changed the following line:
-- find_package(OpenCV 3.0 REQUIRED) -> find_package(OpenCV 4.1.1 REQUIRED)
-- 
-- CMakeLists.txt 119L, 3170C
-- 
-- 41 CMakeFiles/rgbd_tum.dir/build.make:160: recipe for target 'rgbd_tum' failed
```

5) /home/slamgpu/ORB_SLAM2_CUDA/Examples/Stereo/stereo_kitti.cc: In function ‘int main(int, char**)’: /home/slamgpu/ORB_SLAM2_CUDA/Examples/Stereo/stereo_kitti.cc:70:47: error:
‘CV_LOAD_IMAGE_UNCHANGED’ was not declared in this scope
`imLeft = cv::imread(vstrImageLeft[ni],CV_LOAD_IMAGE_UNCHANGED);`

Solution: For the next examples should also add the macro to fix the error:
`#define CV_LOAD_IMAGE_UNCHANGED -1 //new added one`



```
EIGEN_DEPRECATED const unsigned int AlignedBit = 0x80;
/home/slamgpu/ORB_SLAM2_CUDA/Examples/Stereo/stereo_kitti.cc: In function ‘int main(int, char**)’:
/home/slamgpu/ORB_SLAM2_CUDA/Examples/Stereo/stereo_kitti.cc:70:47: error: ‘CV_LOAD_IMAGE_UNCHANGED’ was not declared in this scope
imLeft = cv::imread(vstrImageLeft[ni],CV_LOAD_IMAGE_UNCHANGED);

CMakeFiles/stereo_kitti.dir/build.make:62: recipe for target 'CMakeFiles/stereo_kitti.dir/Examples/Stereo/stereo_kitti.cc.o' failed
make[2]: *** [CMakeFiles/stereo_kitti.dir/Examples/Stereo/stereo_kitti.cc.o] Error 1
CMakeFiles/Makefile2:141: recipe for target 'CMakeFiles/stereo_kitti.dir/all' failed
make[1]: *** [CMakeFiles/stereo_kitti.dir/all] Error 2
Makefile:83: recipe for target 'all' failed
make: *** [all] Error 2
stampgus@slampu-desktop:/~/ORB_SLAM2_CUDA/v1 ~/home/slamgpu/ORB_SLAM2_CUDA/Examples/Stereo/stereo_kitti.cc]
```



```
from /home/slamgpu/Pangolin/include/pangolin/display/viewport.h:31,
from /home/slamgpu/Pangolin/include/pangolin/gl.h:30,
from /home/slamgpu/Pangolin/include/pangolin/pangolin.h:33,
from /home/slamgpu/ORB_SLAM2_CUDA/include/Viewer.h:27,
from /home/slamgpu/ORB_SLAM2_CUDA/include/Tracking.h:26,
from /home/slamgpu/ORB_SLAM2_CUDA/include/System.h:32,
from /home/slamgpu/ORB_SLAM2_CUDA/include/pangolin.h:30,
from /home/slamgpu/ORB_SLAM2_CUDA/Examples/Monocular/mono_tum.cc:30:
/usr/local/include/eigen3/Eigen/src/Core/util/Constants.h:166:37: note: declared here
EIGEN_DEPRECATED const unsigned int AlignedBit = 0x80;

/home/slamgpu/ORB_SLAM2_CUDA/Examples/Monocular/mono_tum.cc: In function ‘int main(int, char**)’:
/home/slamgpu/ORB_SLAM2_CUDA/Examples/Monocular/mono_tum.cc:74:48: error: ‘CV_LOAD_IMAGE_UNCHANGED’ was not declared in this scope
im = cv::imread(string(argv[3]),CV_LOAD_IMAGE_UNCHANGED);

CMakeFiles/mono_tum.dir/build.make:62: recipe for target 'CMakeFiles/mono_tum.dir/Examples/Monocular/mono_tum.cc.o' failed
make[2]: *** [CMakeFiles/mono_tum.dir/Examples/Monocular/mono_tum.cc.o] Error 1
CMakeFiles/Makefile2:178: recipe for target 'CMakeFiles/mono_tum.dir/all' failed
make[1]: *** [CMakeFiles/mono_tum.dir/all] Error 2
Makefile:83: recipe for target 'all' failed
make: *** [all] Error 2
stampgus@slampu-desktop:/~/ORB_SLAM2_CUDA/
```



```
from /home/slamgpu/Pangolin/include/pangolin/display/viewport.h:31,
from /home/slamgpu/Pangolin/include/pangolin/pangolin.h:33,
from /home/slamgpu/ORB_SLAM2_CUDA/include/MapDrawer.h:27,
from /home/slamgpu/ORB_SLAM2_CUDA/include/Viewer.h:26,
from /home/slamgpu/ORB_SLAM2_CUDA/include/Tracking.h:26,
from /home/slamgpu/ORB_SLAM2_CUDA/include/System.h:32,
from /home/slamgpu/ORB_SLAM2_CUDA/Examples/Monocular/mono_kitti.cc:31:
/usr/local/include/eigen3/Eigen/src/Core/util/Constants.h:166:37: note: declared here
EIGEN_DEPRECATED const unsigned int AlignedBit = 0x80;

/home/slamgpu/ORB_SLAM2_CUDA/Examples/Monocular/mono_kitti.cc: In function ‘int main(int, char**)’:
/home/slamgpu/ORB_SLAM2_CUDA/Examples/Monocular/mono_kitti.cc:73:48: error: ‘CV_LOAD_IMAGE_UNCHANGED’ was not declared in this scope
im = cv::imread(vstrImageFilenames[ni],CV_LOAD_IMAGE_UNCHANGED);

In file included from /home/slamgpu/ORB_SLAM2_CUDA/Examples/Monocular/mono_kitti.cc:32:0:
/home/slamgpu/ORB_SLAM2_CUDA/include/Utils.hpp:26:45: warning: missing initializer for member ‘nvtxEventAttributes_v2::size’ [-Wmissing-field-initializers]
    nvtxEventAttributes_t eventAttrib = {0}; \
^
/home/slamgpu/ORB_SLAM2_CUDA/Examples/Monocular/mono_kitti.cc:88:9: note: in expansion of macro ‘PUSH_RANGE’
    PUSH_RANGE("Track Image", 4);
^
/home/slamgpu/ORB_SLAM2_CUDA/include/Utils.hpp:26:45: warning: missing initializer for member ‘nvtxEventAttributes_v2::category’ [-Wmissing-field-initializers]

```

After fix all the issues we should see the follow:



```
/home/slamgpu/ORB_SLAM2_CUDA/gpu/ttx1.cpp:26:45: warning: missing initializer for member 'nvtxEventAttributes_v2::message_type' [-Wmissing-field-initializers]
   nvtxEventAttributes_t eventAttrib = {0};
                           ^
/home/slamgpu/ORB_SLAM2_CUDA/gpu/ttx1.cpp:67: note: in expansion of macro 'PUSH_RANGE'
 #define PUSH_RANGE(_trackImage_, _id) \
  _trackImage_/_id
/home/slamgpu/ORB_SLAM2_CUDA/include/Utils.h:26:45: warning: missing initializer for member 'nvtxEventAttributes_v2::message' [-Wmissing-field-initializers]
   nvtxEventAttributes_t eventAttrib = {0};
                           ^
/home/slamgpu/ORB_SLAM2_CUDA/gpu/ttx1.cpp:67: note: in expansion of macro 'PUSH_RANGE'
 #define PUSH_RANGE(_trackImage_, _id) \
  _trackImage_/_id
[100%] Linking CXX executable tx1
[100%] Built target tx1
slamgpu@slamgpu-desktop:/home/slamgpu/ORB_SLAM2_CUDA$ slamgpu@slamgpu-desktop:/home/slamgpu/ORB_SLAM2_CUDA$ slamgpu@slamgpu-desktop:/home/slamgpu/ORB_SLAM2_CUDA$ slamgpu@slamgpu-desktop:/home/slamgpu/ORB_SLAM2_CUDA$
```

Remember to run `build.sh` before building ROS because the `lib/libORB_SLAM2_CUDA.so` needs to be created first. To build ROS node:

Next step to build ROS by following commands:

```
export ROS_PACKAGE_PATH=${ROS_PACKAGE_PATH}:~/home/slamgpu/ORB_SLAM2_CUDA/Examples/ROS
chmod +x build_ros.sh
./build_ros.sh
```

Pay attention to the path for ROS

In case of issues run the following commands:

```
rosdep init
rosdep update
```

When building finished successfully we will show the follow:



```
[ 100%] Linking CXX executable .._Mono
/usr/bin/ld: warning: libopencv_core.so.3.2, needed by /opt/ros/melodic/lib/libcv_bridge.so, may conflict with libopencv_core.so.4.1
[100%] Built target Mono
slamgpu@slamgpu-desktop:/home/slamgpu/ORB_SLAM2_CUDA$ slamgpu@slamgpu-desktop:/home/slamgpu/ORB_SLAM2_CUDA$ slamgpu@slamgpu-desktop:/home/slamgpu/ORB_SLAM2_CUDA$ slamgpu@slamgpu-desktop:/home/slamgpu/ORB_SLAM2_CUDA$ slamgpu@slamgpu-desktop:/home/slamgpu/ORB_SLAM2_CUDA$ slamgpu@slamgpu-desktop:/home/slamgpu/ORB_SLAM2_CUDA$
```

Done.

If you have problem with the build, sometimes it helps to remove previous build folders:

```
# in ORB_SLAM2_CUDA folder
sudo rm -rf build
sudo rm -rf Thirdparty/DBow2/build
sudo rm -rf Thirdparty/g2o/build
./build.sh
sudo rm -rf Examples/ROS/ORB_SLAM2_CUDA/build
./build_ros.sh
```

Link to web tutorial: https://github.com/thien94/ORB_SLAM2_CUDA

Performance Tip commands

These commands help to increase the Jetson performance.
You should see the fan be active.

- Set power mode to MAXN: (once, saved after reboot)
`sudo nvpmode -m 0`
- Set clocks to high: (every reboot)
`sudo jetson_clocks`

Run SLAM application on the target

Do the following steps:

- 1) Add to end of file `~/.bashrc` the following lines:

```
export PATH=/usr/local/cuda-10.0/bin:$PATH  
export LD_LIBRARY_PATH=/usr/local/cuda-10.0/lib64:$LD_LIBRARY_PATH  
source /opt/ros/melodic/setup.bash  
export  
ROS_PACKAGE_PATH=${ROS_PACKAGE_PATH}:~/home/slamgpu/ORB_SLAM2_CUDA/Examples/ROS
```

- 2) Put the following highlight content to file `ros_mono.launch` in the path:

```
~/ORB_SLAM2_CUDA/Examples/ROS/ORB_SLAM2_CUDA/launch/ros_mono.launch
```

```
<?xml version="1.0"?>  
<launch>  
  
    <arg name="vocabulary_path" default="$(find ORB_SLAM2_CUDA)/../../../../Vocabulary/ORBvoc.txt" />  
    <arg name="camera_setting_path" default="$(find ORB_SLAM2_CUDA)/conf/euroc.yaml" />  
    <arg name="bUseViewer" default="true" />  
    <arg name="bEnablePublishROSTopic" default="false" />  
  
    <node name="ORB_SLAM2_CUDA" pkg="ORB_SLAM2_CUDA" type="Mono" output="screen"  
          args="$(arg vocabulary_path) $(arg camera_setting_path) $(arg bUseViewer) $(arg  
bEnablePublishROSTopic)">  
        <remap from="/camera/image_raw" to="/cam0/image_raw"/>  
    </node>  
  
</launch>
```

3) Add the file euroc.yaml to conf directory as

~/ORB_SLAM2_CUDA/Examples/ROS/ORB_SLAM2_CUDA/conf/euroc.yaml

You can see the content below

```
%YAML:1.0
#-----
# Camera Parameters. Adjust them!
#-----

# Camera calibration and distortion parameters (OpenCV)
Camera.fx: 435.2046959714599
Camera.fy: 435.2046959714599
Camera.cx: 367.4517211914062
Camera.cy: 252.2008514404297

Camera.k1: -0.28340811
Camera.k2: 0.07395907
Camera.p1: 0.00019359
Camera.p2: 1.76187114e-05

Camera.width: 752
Camera.height: 480

# Camera frames per second
Camera.fps: 20.0

# stereo baseline times fx
Camera.bf: 47.90639384423901

# Color order of the images (0: BGR, 1: RGB. It is ignored if images are grayscale)
Camera.RGB: 1

# Close/Far threshold. Baseline times.
ThDepth: 35

#-----
# Stereo Rectification. Only if you need to pre-rectify the images.
# Camera.fx, .fy, etc must be the same as in LEFT.P
#-----

LEFT.height: 480
LEFT.width: 752
LEFT.D: !opencv-matrix
rows: 1
cols: 5
dt: d
data:[-0.28340811, 0.07395907, 0.00019359, 1.76187114e-05, 0.0]
LEFT.K: !opencv-matrix
rows: 3
cols: 3
dt: d
data: [458.654, 0.0, 367.215, 0.0, 457.296, 248.375, 0.0, 0.0, 1.0]
LEFT.R: !opencv-matrix
rows: 3
cols: 3
dt: d
data: [0.999966347530033, -0.001422739138722922, 0.008079580483432283, 0.001365741834644127, 0.9999741760894847,
0.007055629199258132, -0.008089410156878961, -0.007044357138835809, 0.9999424675829176]
LEFT.P: !opencv-matrix
rows: 3
cols: 4
dt: d
data: [435.2046959714599, 0, 367.4517211914062, 0, 0, 435.2046959714599, 252.2008514404297, 0, 0, 0, 1, 0]

RIGHT.height: 480
```

```

RIGHT.width: 752
RIGHT.D: !!opencv-matrix
  rows: 1
  cols: 5
  dt: d
  data:[-0.28368365, 0.07451284, -0.00010473, -3.555907e-05, 0.0]
RIGHT.K: !!opencv-matrix
  rows: 3
  cols: 3
  dt: d
  data: [457.587, 0.0, 379.999, 0.0, 456.134, 255.238, 0.0, 0.0, 1]
RIGHT.R: !!opencv-matrix
  rows: 3
  cols: 3
  dt: d
  data: [0.9999633526194376, -0.003625811871560086, 0.007755443660172947, 0.003680398547259526, 0.9999684752771629, -
0.007035845251224894, -0.007729688520722713, 0.007064130529506649, 0.999945173484644]
RIGHT.P: !!opencv-matrix
  rows: 3
  cols: 4
  dt: d
  data: [435.2046959714599, 0, 367.4517211914062, -47.90639384423901, 0, 435.2046959714599, 252.2008514404297, 0, 0, 1, 0]

#-----
# ORB Parameters
#-----

# ORB Extractor: Number of features per image
ORBextractor.nFeatures: 2000

# ORB Extractor: Scale factor between levels in the scale pyramid
ORBextractor.scaleFactor: 1.2

# ORB Extractor: Number of levels in the scale pyramid
ORBextractor.nLevels: 8

# ORB Extractor: Fast threshold
# Image is divided in a grid. At each cell FAST are extracted imposing a minimum response.
# Firstly we impose iniThFAST. If no corners are detected we impose a lower value minThFAST
# You can lower these values if your images have low contrast
ORBextractor.iniThFAST: 20
ORBextractor.minThFAST: 7

#-----
# Viewer Parameters
#-----

Viewer.KeyFrameSize: 0.05
Viewer.KeyFrameLineWidth: 1
Viewer.GraphLineWidth: 0.9
Viewer.PointSize: 2
Viewer.CameraSize: 0.08
Viewer.CameraLineWidth: 3
Viewer.ViewpointX: 0
Viewer.ViewpointY: -0.7
Viewer.ViewpointZ: -1.8
Viewer.ViewpointF: 500

```

Run ORB SLAM with ROS

1) Set clocks to high: (every reboot) - Optional

```
sudo jetson_clocks
```

2) Run SLAM application:

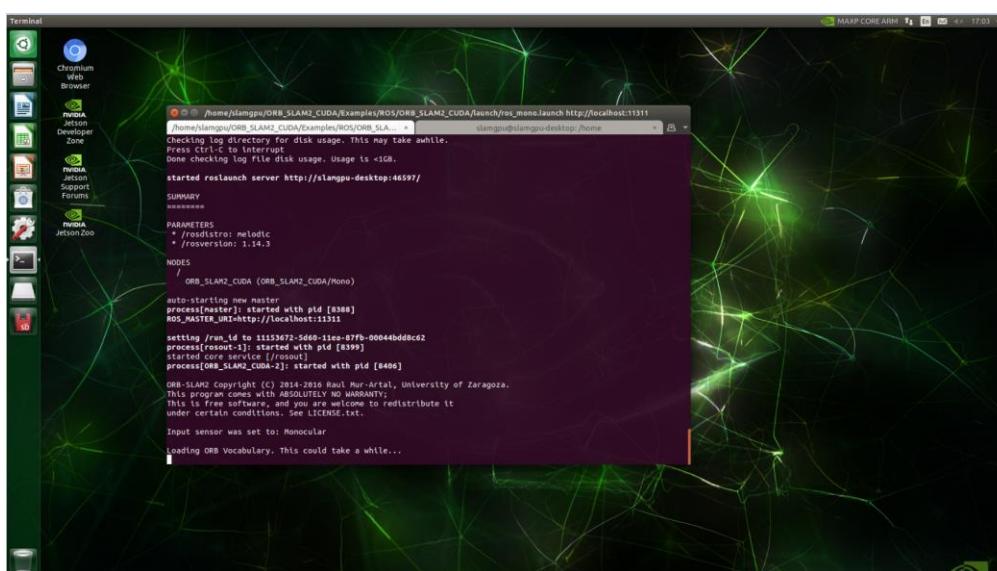
```
roslaunch /home/slampu/ORB_SLAM2_CUDA/Examples/ROS/ORB_SLAM2_CUDA/launch/ros_mono.launch
```

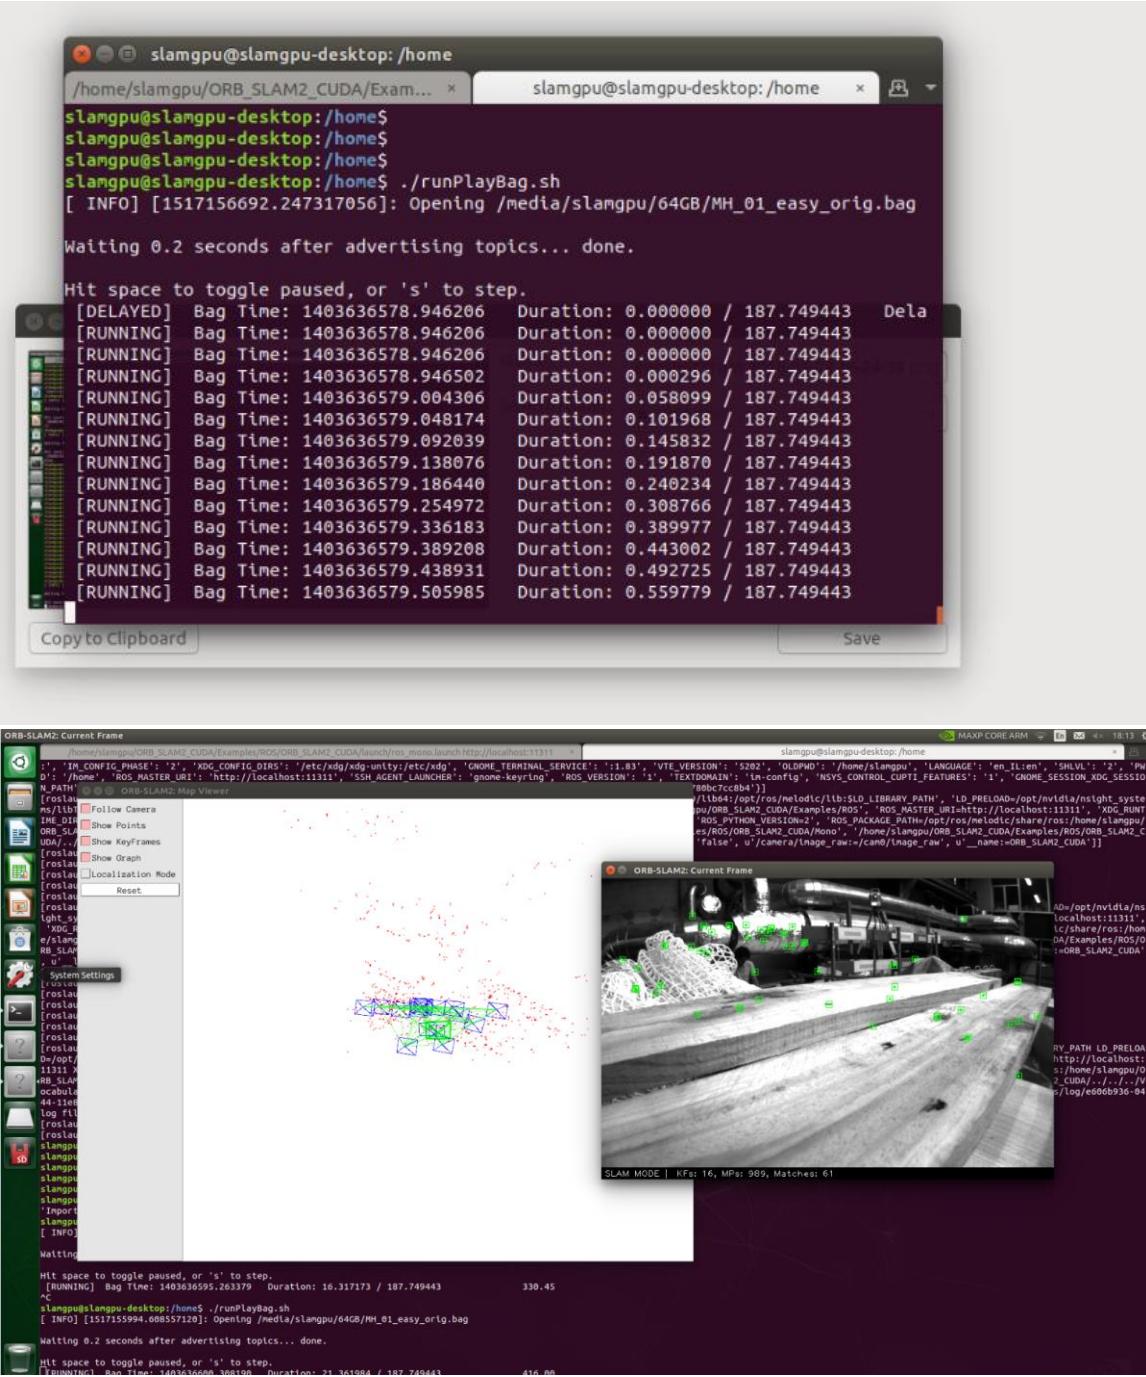
3) Play bag file to SLAM application (bag file located in SDcard)

```
rosbag play /media/slampu/64GB/MH_01_easy_orig.bag
```

You can run it from device\host.

For running in remote (host) you should change the display id.





Run ORB SLAM without ROS

1) Set clocks to high: (every reboot)

```
sudo jetson_clocks
```

2) Run SLAM application including parameters:

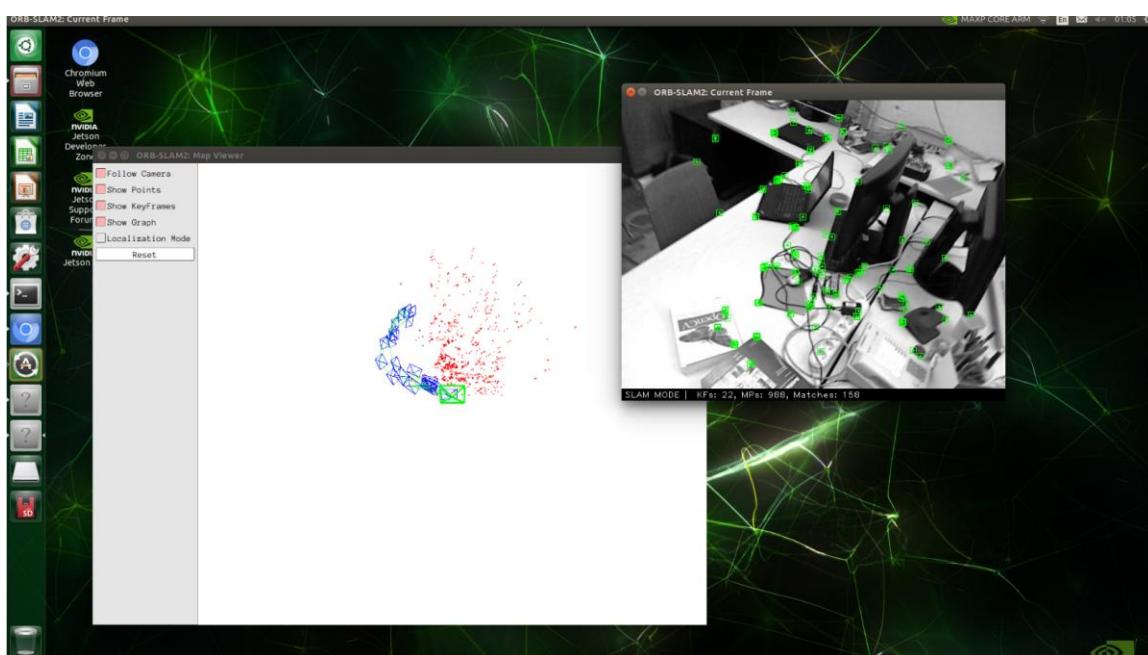
```
cd /home/slamgpu/ORB_SLAM2_CUDA
```

```
./build/mono_tum Vocabulary/ORBvoc.txt Examples/Monocular/TUM1.yaml
```

```
/media/slamgpu/64GB/rgbd_dataset_freiburg1_desk true
```

You can run it from device\host.

For running in remote (host) you should change the display id.





Configuration issue during SLAM running

In case of error as below you should configure the environment and modify the configuration files

- 1)~/ORB_SLAM2_CUDA/Examples/ROS/ORB_SLAM2_CUDA/conf/euroc.yaml
- 2)~/ORB_SLAM2_CUDA/Examples/ROS/ORB_SLAM2_CUDA/launch/ros_mono.launch

```
/home/slangpu/ORB_SLAM2_CUDA/Examples/ROS/ORB_SLAM2_CUDA/launch/ros_mono.launch
NODES
/
ORB_SLAM2_CUDA (ORB_SLAM2_CUDA/Mono)

auto-starting new master
process[master]: started with pid [11741]
ROS_MASTER_URI=http://localhost:11311

setting /run_id to 2c102856-1daf-11ea-96a8-920e2eb0fd39
process[rosout-1]: started with pid [11752]
started core service [/rosout]
process[ORB_SLAM2_CUDA-2]: started with pid [11759]

Usage: rosrun ORB_SLAM2 Mono path_to_vocabulary path_to_settings bUseViewer bEnablePublishROSTopic
[ORB_SLAM2_CUDA-2] process has died [pid 11759, exit code 1, cmd /home/slangpu/ORB_SLAM2_CUDA/Examples/ROS/ORB_SLAM2_CUDA/Mono /home/slangpu/ORB_SLAM2_CUDA/Examples/ROS/ORB_SLAM2_CUDA/.../.../.../Vocabulary/ORBvoc.txt /home/slangpu/ORB_SLAM2_CUDA/Examples/ROS/ORB_SLAM2_CUDA/conf/euroc.yaml true false false input_stream:=cam0/image_raw __name:=ORB_SLAM2_CUDA __log:=/home/slangpu/.ros/log/2c102856-1daf-11ea-96a8-920e2eb0fd39/ORB_SLAM2_CUDA-2*.log]
log file: /home/slangpu/.ros/log/2c102856-1daf-11ea-96a8-920e2eb0fd39/ORB_SLAM2_CUDA-2*.log
```

ROS Useful commands

You can read about the ROS commands in the following links:

- <https://github.com/bosch-io/bcx18-openADx-examples/wiki/Useful-ROS-Commands>
- <http://wiki.ros.org/ROS/CommandLineTools>
- https://subscription.packtpub.com/book/hardware_and_creative/9781782175193/1/ch011_vl1sec15/ros-commands-summary

roscore command

- roscore -p <port>: start a roscore process directly on given port

roslaunch command

- roslaunch <launchfile>: run a launch file (which starts a roscore if none is running)

rosnode command

- rosnode list: show all running nodes
- rosnode info <nodename>: show infos for node
- rosnode ping <nodename>: try to reach a running node

rosbag command

- rosbag info <bagfile>: show all topics contained in a bagfile
- rosbag play <bagfile>: play all messages recorded in a bagfile
- rosbag record -o <bagfile> -a: record all topics into a bagfile

rostopic command

- rostopic echo <topic_name>: listen to topic
- rostopic list: show all published and subscribed topics
- rostopic pub <topic_name> <message_type> 'data': manually publish a message to a topic, e.g. rostopic pub /flip_image std_msgs/Bool '{data:true}' or rostopic pub /mavros/rc/override mavros_msgs/OverrideRCIn '[yaw, 0, throttle, 0, 0, 0, 0]'
- rostopic info <topic_name>: print infos about topic
- rostopic type <topic_name>: the type of the topic

rvizrviz tool

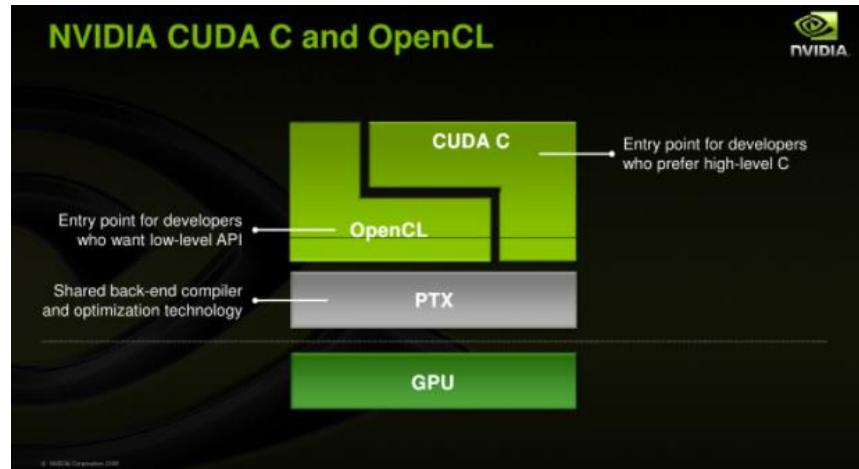
Tool for visualizing sensor data

rqt_console

GUI for displaying and filtering ROS messages

CUDA

CUDA® is a parallel computing platform and programming model invented by NVIDIA. It enables dramatic increases in computing performance by harnessing the power of the graphics processing unit (GPU).

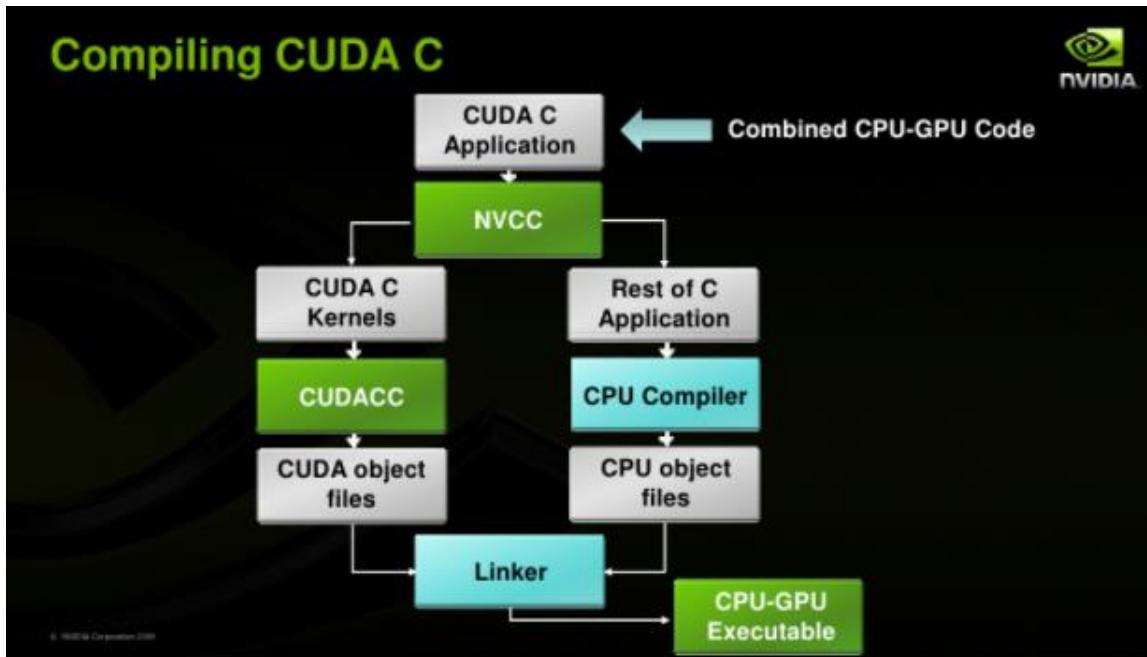


CUDA was developed with several design goals in mind:

- Provide a small set of extensions to standard programming languages, like C, that enable a straightforward implementation of parallel algorithms. With CUDA C/C++, programmers can focus on the task of parallelization of the algorithms rather than spending time on their implementation.
- Support heterogeneous computation where applications use both the CPU and GPU. Serial portions of applications are run on the CPU, and parallel portions are offloaded to the GPU. As such, CUDA can be incrementally applied to existing applications. The CPU and GPU are treated as separate devices that have their own memory spaces. This configuration also allows simultaneous computation on the CPU and GPU without contention for memory resources.

CUDA-capable GPUs have hundreds of cores that can collectively run thousands of computing threads. These cores have shared resources including a register file and a shared memory. The on-chip shared memory allows parallel tasks running on these cores to share data without sending it over the system memory bus.

Cuda Compilation Flow



Cuda application is compiled with NVCC compiler which create the object files for CPU and GPU. These object files are linked to execution file which ready for running on CPU-GPU enviroment.

Cuda Entry Point and Reference

NVIDIA Site:

Cuda toolkit link: <https://docs.nvidia.com/cuda/index.html>

Installation Guide: <https://docs.nvidia.com/cuda/index.html#installation-guides>

Programming Guide: <https://docs.nvidia.com/cuda/index.html#programming-guides>

Cuda API: <https://docs.nvidia.com/cuda/index.html#cuda-api-references>

Cuda Samples: <https://docs.nvidia.com/cuda/cuda-samples/index.html>

Debuging API: <https://docs.nvidia.com/cuda/debugger-api/index.html>

From Internet:

Cuda tutorial : <https://www.slideshare.net/maheshkha/cuda-tutorial>

Communication between Host and Device

The communication between the Host (CPU) and Device (GPU) is occurred by transfer memory buffers and threads synchronization (barriers) for indication all the threads are finished their work.

These buffers are used to transfer data to GPU for perform parallel calculations and then the buffer results is returned to the CPU.

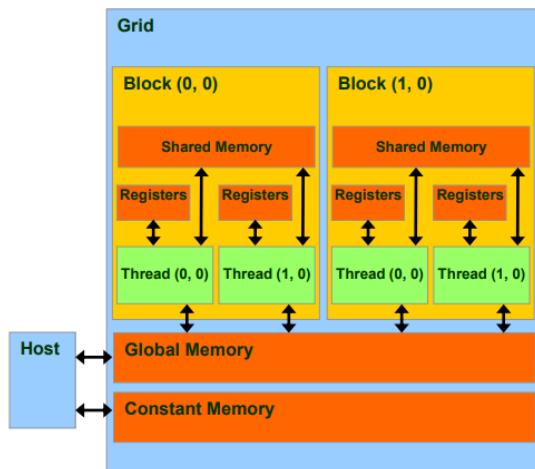
When the one thread achieved the desirable result we can update other threads to stop their calculation and jump to the barrier for waiting all the other threads will finish their work and when all the threads ready to complete the task and return the results to the host.

Memories Hierarchy

1. Registers ~1 cycle
2. Shared memory ~5 cycles
3. Global memory ~500 cycles

There is one more memory for constants ~5 cycles with caching.

Host can copy data (RW) to the global memory.



The bus between CPU and GPU is PCIe.

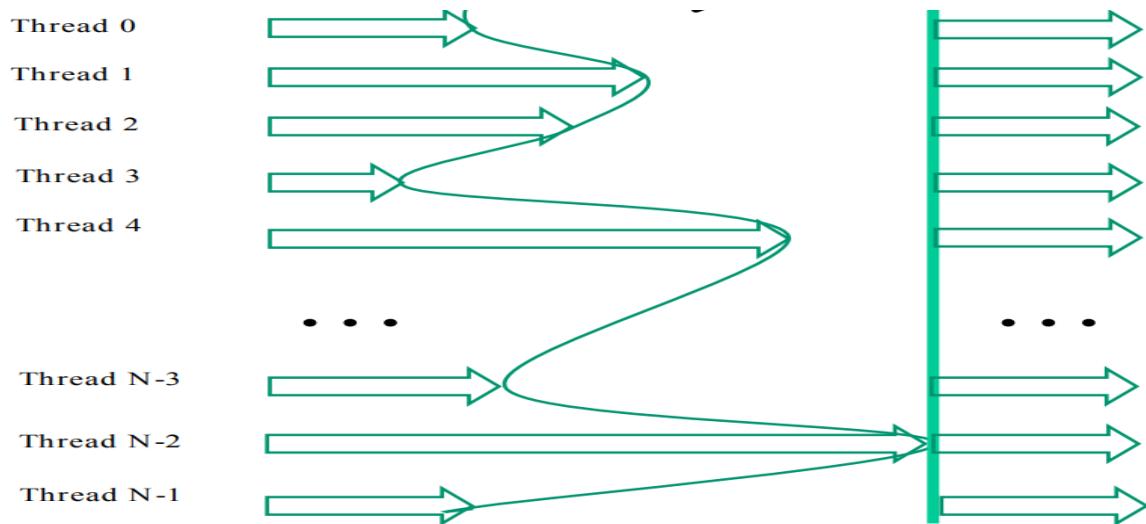
By using shared memory we can increase the memory bandwidth as we can use each memory data several times and improve the bus efficiency.

Cuda API for transfer data

Fucntion	Usage	Note
cudaMalloc()	Allocates buffer in the device memory	
cudaFree()	Free allocated device buffer	
cudaMemcpy()	memory data transfer (Copy data from\to Host to\from Device)	Two flags: cudaMemcpyHostToDevice cudaMemcpyDeviceToHost

Cuda API for synchronization

Function	Usage	Note
<code>__syncthreads()</code>	All threads in the same block must reach the <code>__syncthreads()</code> before any can move on	



ORB SLAM GPU Cuda Code

Code Structure



src\cuda\Fast_gpu.cu – Fast feature detection implementation using GPU.

src\cuda\Orb_gpu.cu – ORB compute descriptors using GPU.

src\cuda\Allocator_gpu.cu – allocation/free buffer in GPU.

src\ORBExtractor.cc – CPU code entry point for CPU->GPU calls.

Jetson code path (Target - Tetson):

/home/slamgpu/ORB_SLAM2_CUDA

Main Kernel Cuda functions

Tasks from CPU to GPU to improve performance are:

- calcOrb_kernel()
- tileCalcKeypoints_kernel()
- IC_Angle_kernel()
- addBorder_kernel()

The following files contains these function:



I will explain the first two mains kernel functions:

tileCalcKeypoints_kernel()

```
__global__ void tileCalcKeypoints_kernel(const PtrStepSz<img>, short2 * kpLoc, float * kpScore, const unsigned int maxKeypoints, const int highThreshold, const int lowThreshold, PtrStepI scoreMat,
                                         const int counter_ptr) {
    const int i = threadIdx.x + blockIdx.x * blockDim.x + 3;
    const int j = threadIdx.y + blockIdx.y * blockDim.y + 4 + 3;
    const int tid = threadIdx.y * blockDim.x + threadIdx.x;
    __shared__ bool haskp;
    if (tid == 0) {
        haskp = false;
    }
    bool isKp[4] = {0};
    for (int t = 0; t < 4; ++t) {
        if ((i+t < img.rows - 3 && j < img.cols - 3)) {
            isKp[t] = isKeyPoint2(img, i+t, j, highThreshold, scoreMat);
        }
    }
    // barrier
    __syncthreads();
    for (int t = 0; t < 4; ++t) {
        if (isKp[t]) {
            isKp[t] = false;
            short2 loc = make_short2(j, i+t);
            if (isMax(loc, scoreMat)) {
                int score = scoreMat(loc.y, loc.x);
                haskp = true;
                const unsigned int ind = atomicInc(counter_ptr, (unsigned int)(-1));
                if (ind < maxKeypoints) {
                    kpLoc[ind] = loc;
                    kpScore[ind] = static_cast<float>(score);
                }
            }
        }
    }
    // barrier
    __syncthreads();
    if (haskp) return;
    // lower the threshold and try again
    for (int t = 0; t < 4; ++t) {
        if ((i+t < img.rows - 3 && j < img.cols - 3)) {
            isKp[t] = isKeyPoint2(img, i+t, j, lowThreshold, scoreMat);
        }
    }
    // barrier
    __syncthreads();
    for (int t = 0; t < 4; ++t) {
        if (isKp[t]) {
            short2 loc = make_short2(j, i+t);
            if (isMax(loc, scoreMat)) {
                int score = scoreMat(loc.y, loc.x);
                const unsigned int ind = atomicInc(counter_ptr, (unsigned int)(-1));
                if (ind < maxKeypoints) {
                    kpLoc[ind] = loc;
                    kpScore[ind] = static_cast<float>(score);
                }
            }
        }
    }
}
```

This function calculates key points in parallel computation in GPU.

For input image:

```

void GpuFast::detectAsync(InputArray _image) {
    const cv::cuda::GpuMat image = _image.getGpuMat();
    if (scoreMat.empty()) {
        // If it is not empty, then it's already allocated by previous iteration
        // and I ASSUME THE DIMENSIONS ARE CONSISTENT ACROSS EVERY ITERATION
        // else THIS WILL BREAK
        scoreMat = GpuMat(image.size(), CV_32SC1);
    }
    scoreMat.setTo(Scalar::all(0), cvStream);
    checkCudaErrors(cudaMemsetAsync(counter_ptr, 0, sizeof(unsigned int), stream));
    dim3 dimBlock(32, 8);
    dim3 dimGrid(divUp(image.cols, dimBlock.x), divUp(image.rows, dimBlock.y * 4));
    tileCalcKeypoints_kernel<<>(dimGrid, dimBlock, 0, stream>>>(image, kpLoc, kpScore, maxKeypoints, highThreshold,
    lowThreshold, scoreMat, counter_ptr);
    checkCudaErrors(cudaGetLastError());
}

```

We calculate the Grid size (number of blocks) and block size (Threads) per image size.

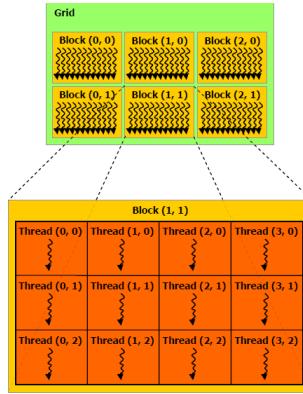


Image source : <https://www.3dgep.com/cuda-thread-execution-model/>

In our case threads per block are 32*8

By dividing image size cordination by coresponding thread count we get the blocks per X or Y.

For dim3 dimGrid(divUp(image.cols, dimBlock.x), divUp(image.rows, dimBlock.y * 4));

Number blocks in X dimention = divUp(image.cols, dimBlock.x)

Number blocks in Y dimention = divUp(image.rows, dimBlock.y * 4))

For Y dimention we create 4 calutlations per thread.

The key point calculation per thread is divided to several steps:

- 1) Indexs calculation per thread
- 2) thread 0 in the block set shared memory flag to 0 which indicate if key point is found by threads
- 3) Check 4 potentialion key points per thread
- 4) Syncronize all the threads in block
- 5) Verify the potentialion key points results
- 6) Syncronize all the threads
- 7) if any key point is found all the threads will finish caluclation
- 8) if not found key point retry steps 3-7 with lower threshold

The target of all threads is to collect several key points.

When key point is found and is verified then atomic counter is incremented and the key point is saved with his score.

The counter is shared between all the threads in all the blocks.

The array for all key points is shared between all the threads in all the blocks.

calcOrb_kernel()

Calculation per key point

```
void GpuOrb::launch_async(InputArray _image, const KeyPoint * _keypoints, const int npoints) {
    if (npoints == 0) {
        POP_RANGE;
        return ;
    }
    const GpuMat image = _image.getGpuMat();
    checkCudaErrors( cudaMemcpyAsync(_keypoints, _keypoints, sizeof(KeyPoint) * npoints, cudaMemcpyHostToDevice, stream) );
    desc = descriptors.rowRange(0, npoints);
    desc.setTo(scalar::all(0), cvStream);

    dim3 dimBlock(32);
    dim3 dimGrid(npoints);
    calcOrb_kernel<<<dimGrid, dimBlock, 0, stream>>>(image, keypoints, npoints, desc);
    checkCudaErrors( cudaGetLastError() );
}
```

Block size 32 threads

Grid size is per number of key points

```
__global__ void calcOrb_kernel(const PtrStepb image, KeyPoint * keypoints, const int npoints,
    PtrStepb descriptors) {
    int id = blockIdx.x;
    int tid = threadIdx.x;
    if (id >= npoints) return;

    const KeyPoint &kpt = keypoints[id];
    short2 loc = make_short2(kpt.pt.x, kpt.pt.y);
    const Point * pattern = ((Point *)c_pattern) + 16 * tid;

    uchar * desc = descriptors.ptr(id);
    const float factorPI = (float)(CV_PI/180.f);
    float angle = (float)kpt.angle * factorPI;
    float a = (float)cosf(angle), b = (float)sinf(angle);

    int t0, t1, val;
    t0 = GET_VALUE(0); t1 = GET_VALUE(1);
    val = t0 < t1;
    t0 = GET_VALUE(2); t1 = GET_VALUE(3);
    val |= (t0 < t1) << 1;
    t0 = GET_VALUE(4); t1 = GET_VALUE(5);
    val |= (t0 < t1) << 2;
    t0 = GET_VALUE(6); t1 = GET_VALUE(7);
    val |= (t0 < t1) << 3;
    t0 = GET_VALUE(8); t1 = GET_VALUE(9);
    val |= (t0 < t1) << 4;
    t0 = GET_VALUE(10); t1 = GET_VALUE(11);
    val |= (t0 < t1) << 5;
    t0 = GET_VALUE(12); t1 = GET_VALUE(13);
    val |= (t0 < t1) << 6;
    t0 = GET_VALUE(14); t1 = GET_VALUE(15);
    val |= (t0 < t1) << 7;

    desc[tid] = (uchar)val;
}
```

Per thread in block there is check for existing key point.

Each thread do calculation on the key point.

Profiling Tools

The Visual Profiler is a graphical profiling tool that displays a timeline of your application's CPU and GPU activity, and that includes an automated analysis engine to identify optimization opportunities.

Old tools:

- [Nvprof profiling](#) tool enables you to collect and view profiling data from the command-line.
- [Visual Profiler](#) allow to view the data in graphical view.

These old tools will be depressed in the future.

New tools: [Nvidia tools overview](#)

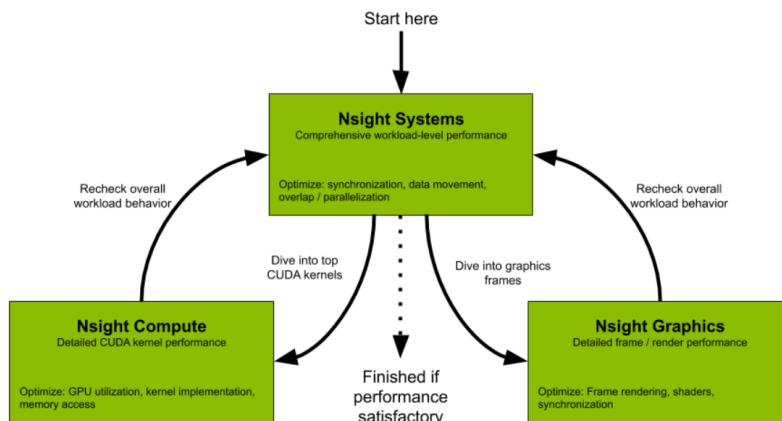


Figure 1. Flowchart describing working with new NVIDIA Nsight tools for performance optimization

[NVIDIA Nsight Compute](#) is an interactive kernel profiler for CUDA applications. It provides detailed performance metrics and API debugging via a user interface and command line tool. In addition, its baseline feature allows users to compare results within the tool. Nsight Compute provides a customizable and data-driven user interface and metric collection and can be extended with analysis scripts for post-processing results.

[NVIDIA Nsight Systems](#) is a system-wide performance analysis tool designed to visualize an application's algorithms, help you identify the largest opportunities to optimize, and tune to scale efficiently across any quantity or size of CPUs and GPUs; from large server to our smallest SoC.

[Nsight Graphics](#) is a standalone developer tool that enables you to debug, profile, and export frames built with Direct3D (11, 12, DXR), Vulkan (1.2, NV Vulkan Ray Tracing Extension), OpenGL, OpenVR, and the Oculus SDK.

Comparisons Table

Tool	System View (OS)	Cuda View	Cuda Debuger	Debug on Remote Host	Debug on Device (Jetson)	For which tasks the tool is optimized
nvprof	X	Y	X	Y	Y	Capture Cuda trace for analysis.
Visual Profiler	X	Y	X	Y	Y	Graphical view
Nsight Compute	X	Y	Y	Y	N	Optimized for Cuda We can use this tool as debugger to follow the code.
Nsight System	Y	High level Cuda information	X	Y	N	Optimized for System view You can see events and timelines for threads in CPU and GPU. In each time which libraries were used. Understand the Cuda calls. Percentage for Computation VS Memory. Context switches in GPU.
Nsight Graphics	?	?	?	Y	N	Allow real-time analysis

How to Install JetPack

Depending on your Jetson device, there are one or two ways to install JetPack.

[NVIDIA SDK Manager](#) supports JetPack installation on these Jetson products:

- NVIDIA Jetson Nano modules on a Jetson Nano Developer Kit carrier board
- NVIDIA Jetson TX1 on a Jetson TX2 Developer Kit carrier board
- NVIDIA Jetson TX2 series modules on a Jetson TX2 Developer Kit carrier board
- NVIDIA Jetson AGX Xavier on a Jetson AGX Xavier Developer Kit carrier board

A Linux host computer running Ubuntu Linux x64 version 18.04 or 16.04 is required to run SDK Manager.

Detailed instructions can be found here: <https://docs.nvidia.com/sdk-manager/index.html>

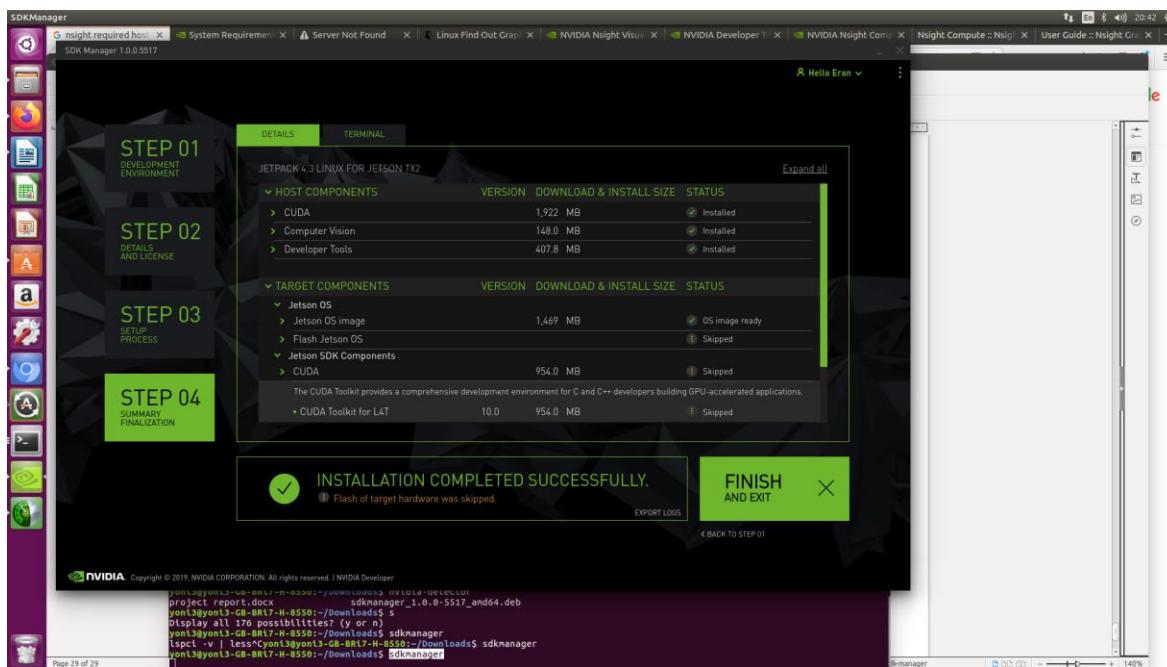
You can install in other options see following link:

<https://docs.nvidia.com/jetson/jetpack/install-jetpack/index.html>

After sdkmanager installed, you can run the GUI by the command:

`sdkmanager`

Then do the step1 – step4 to install Jetpack with all packages as you can see below



When done you will see INSTALLATION COMPLETED SUCCESSFULLY.

SLAM APPLICATION RUNNING MODES

There are two modes I used during the research:

- 1) SLAM Application with ROS.

- a) First command launch ROS core:

```
roslaunch  
/home/slamgpu/ORB_SLAM2_CUDA/Examples/ROS/ORB_SLAM2_CUDA/launch/ros_mono  
.launch
```

- b) Second command to provide the bag file:

```
rosbag play /media/slamgpu/64GB/MH_01_easy_orig.bag
```

- 2) SLAM Application in standalone mode.

This mode uses mono_tum application with some parameters to load the dataset

Application:

```
./build/mono_tum
```

Working Directory:

```
/home/slamgpu/ORB_SLAM2_CUDA
```

Application parameters:

```
Vocabulary/ORBvoc.txt Examples/Monocular/TUM1.yaml  
/media/slamgpu/64GB/rgbd_dataset_freiburg1_desk true
```

environment

```
DISPLAY=:1
```

In this mode I used with the tools as we must connect in root mode.

I tried also in ROS mode however ROS isn't supported in default to run in root mode so it can be new task for the next project.

Nvprof (Nvidia profiling tool)

Configuration nvprof to capture and profile Cuda applications

- 1) Move to root

```
sudo su
```

- 2) Configure envirment variables permanent for terminal windows

Add to end of file ~/.bashrc (root bashrc) the following lines:

```
export PATH=/usr/local/cuda-10.0/bin:$PATH
export LD_LIBRARY_PATH=/usr/local/cuda-10.0/lib64:$LD_LIBRARY_PATH
source /opt/ros/melodic/setup.bash
export
ROS_PACKAGE_PATH=${ROS_PACKAGE_PATH}:~/home/slamgpu/ORB_SLAM2_CUDA/Examples/ROS
```

Or you can run manually on each terminal window the command

source ~/home/tx2/.bashrc to prepare the environment

- 3) Verify you have nvprof installed by run command nvprof

In case the tool is missing use sudo apt get to install the tool

```
sudo apt-get install nvprof
```

- 4) Run the nvprof profiler command with ORB SLAM application

With Sudo (Root mode - Preferred)

```
sudo nvprof --profile-child-processes --export-profile /media/slamgpu/64GB/output_%p.prof
roslaunch
/home/slamgpu/ORB_SLAM2_CUDA/Examples/ROS/ORB_SLAM2_CUDA/launch/ros_mono.launch
```

Without Sudo

```
nvprof --profile-child-processes --export-profile /media/slamgpu/64GB/output_%p.prof roslaunch
/home/slamgpu/ORB_SLAM2_CUDA/Examples/ROS/ORB_SLAM2_CUDA/launch/ros_mono.launch
```

- 5) Run the bag file (dataset) with the captured data on different termainl window

```
rosbag play /media/slamgpu/64GB/MH_01_easy_orig.bag
```

- 6) The output file will be saved in the sdcard (optional) as

```
/media/slamgpu/64GB/output_%p.prof
```

For example /media/slampu/64GB/output_12212.prof
%p is for the process id

Analysys the output file can be done with nvprof or visual profiler:

nvprof command for analyze:

nvprof --import-profile file_name

Note: The profilers use SQLite as the format of the export profiles. Writing files in such format may require more disk operations than writing a plain file. Thus, exporting profiles to slower devices such as a network drive may slow down the execution of the application.

Example:

nvprof --import-profile /media/slampu/64GB/output_12212.prof

For analysis:

Transfer the output .prof files to the host PC.

run /usr/local/cuda-10.0/bin/nvvp (nvidia visual profiler).

Use file->import to open the file (nvprof type).

```
root@slampu-desktop:/home
root@slampu-desktop:/home# ls -l /media/slampu/64GB/output_12212.prof
219 -rw-r--r-- 1 slampu slampu 62M Dec 17 18:35 /media/slampu/64GB/output_12212.prof
root@slampu-desktop:/home# nvprof -c
root@slampu-desktop:/home# nvprof --import-profile /media/slampu/64GB/output_12212.prof
===== Profiling results:
GPU activities: 55.95% 12.3478s 28760 594.79us 76.039us 3.1778ns
:PrStep<int>, unsigned int*)
10.66% 2.35389s 28760 113.35us 50.679us 345.79us ORB_SLAM2::cuda::tileCalcKeypoints_kernel<(cv::cuda::PtrStepSz<unsigned char>, short2*, float*, unsigned int, int, int, cv::cuda
:PrStep<int>, unsigned int*)
9.44% 2.08334s 28760 100.35us 26.239us 246.24us void cv::cuda::grid_copy_detail<copy>:cv::cuda::RemapPtr<cv::cuda::BrdBase<cv::cuda::BrdReflect101, cv::cuda::GlobPtr<uns
signed char>, _GLOBPTR<cv::cuda::BrdBase<cv::cuda::BrdReflect101, cv::cuda::GlobPtr<unsigned char>, cv::cuda::BrdReflect101, cv::cuda::BrdReflect101, cv::cuda::GlobP
d char>, cv::cuda::GlobPtr<unsigned char>, int, int>
7.67% 1.09288s 28760 81.54us 16.479us 1.2864ns void column_filter::linearColumnFilter<int>, float, unsigned char, cv::cuda::device::BrdColReflect101<float>,(cv::cuda::PtrSt
psz<float>, cv::cuda::PtrStep<unsigned char>, float const *, int, float)
5.37% 1.18450s 28760 57.056us 13.120us 373.92us void row_filter::linearRowFilter<int>, unsigned char, float, cv::cuda::device::BrdRowReflect101<unsigned char>,(cv::cuda::PtrSt
tepSz<unsigned char>, cv::cuda::PtrStep<float>, float const *, int, unsigned char)
4.76% 1.09214s 28760 50.848us 17.750us 10.980us ORB_SLAM2::cuda::ICAngle kernel<(cv::cuda::PtrStep<unsigned char>, cv::KeyPoint*, int, int)
4.26% 948.910us 18165 51.010us 14.239us 23.810us void cv::cuda::device::resize_linear<unsigned char>,(cv::cuda::PtrStepSz<unsigned char>, unsigned char, float, float)
1.04% 230.17ms 44123 5.2160us 320ns 87.199us [CUDA memcpy HtoD]
0.29% 63.830us 28760 3.0740us 1.2800ns 8.1600ns ORB_SLAM2::cuda::addBorder_kernel<(cv::KeyPoint*, int, int, int, int)
0.26% 57.164ms 62280 917ns 160ns 6.2400ns [CUDA memcpy DtoH]
0.08% 18.060ms 62280 289ns 159ns 2.6880ns [CUDA memset]
41.32% 20.0913s 15302 122.90us 30.71us 1.22771s
23.71% 11.7723s 83040 141.830s 25.471us cudaStreamCreate
10.76% 5.32850s 64875 82.1340s 10.490us 16.567ns cudaStreamSynchronize
9.32% 4.01511s 28760 222.31us 56.543us 14.301ns cudaMemcpy2DAsync
6.36% 3.15855s 41520 75.880us 35.391us 2.7383ns cudaMemset2DAsync
2.22% 1.10098s 2595 424.27us 65.824us 11.959ns cudaFree
1.84% 913.590us 28760 44.060us 22.367us 1.0790ns cudaMemset3DAsync or with the visual Profiler.
1.07% 299.200us 2599 14.752us 7.570us 1.5890ns cudaMemset3D
1.53% 760.070us 14 54.290ns 10.399us 440.77ns cudaStreamCreate
0.65% 322.440us 2661 123.97us 27.880ns 1.8783ns cudaMallocPitch
0.36% 146.090ms 163485 893ns 480ns 641.94us cudaGetLastFrob
0.26% 98.4300s 28760 4.7410us 2.8150us 636.47us cudaGetDevice
0.06% 11.2070s 28760 2.970us 1.180us 673.00us cudaGetDeviceName
0.06% 728.230us 10 72.024us 19.136us 1.060us cudaGetDevice
0.06% 691.190us 468 1.4780us 704ns 50.271us cuDeviceGetAttribute
0.06% 226.170us 4 56.543us 39.264us 74.719us cudaMemcpyToSymbol
0.06% 67.984us 1 67.984us 67.984us 67.984us cuDeviceGetProperty
0.06% 67.742us 6 11.290us 8.1000us 15.900ns cuDeviceDestroy
0.06% 6.6000us 5 13.000us 8.0000us 15.900ns cuDeviceGetMem
0.06% 28.542us 4 1.150us 1.5510us 10.840ns cuDeviceAttachMemAsync
0.06% 12.512us 7 1.7870us 1.1200us 4.0000ns cuDeviceGetMemCount
0.06% 12.384us 4 3.0960us 2.8480us 3.7440us cuInt
0.06% 10.240us 4 2.5600us 1.7600us 3.3920us cuDriverGetVersion Ubuntu 18.10
0.06% 10.112us 5 2.0220us 1.8880us 2.1760us cuDeviceGetName
0.06% 7.7120us 6 1.2800us 896ns 1.5600us cuDeviceGetUuid
0.06% 4.7870us 5 540ns 896ns 1.02400us cuDeviceGetUuid
0.06% 544ns 1 544ns 544ns 544ns cuDeviceGetDeviceCount
Note: The -c option is only required when jRE is not included in CUDA Toolkit package and jRE 1.8 is not in the default path.
```

To run Visual Profiler on Ubuntu 18.04 or Ubuntu 18.10:

Make sure that you invoke Visual Profiler with the command-line option included as shown below:

```
nvvp -vm /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/java
```

Note: The `-vm` option is only required when JRE is not included in CUDA Toolkit package and JRE 1.8 is not in the default path. On Ubuntu 18.10, if you get error "no `swt-pi-gtk` in `java.library.path`" when running Visual Profiler, then you need to install GTK2. Type the below command to install the required GTK2.

```
apt-get install libgtk2.0-0
```

Remote profiling is the process of collecting profile data from a remote system that is different than the host system at which that profile data will be viewed and analyzed. There are two ways to perform remote profiling. You can profile your remote application directly from `nvidia-nvprof` or the Visual Profiler. Or you can use `nvprof` to collect the profile data on the remote system (target) and then use `nvvp` on the host system to view and analyze the data.

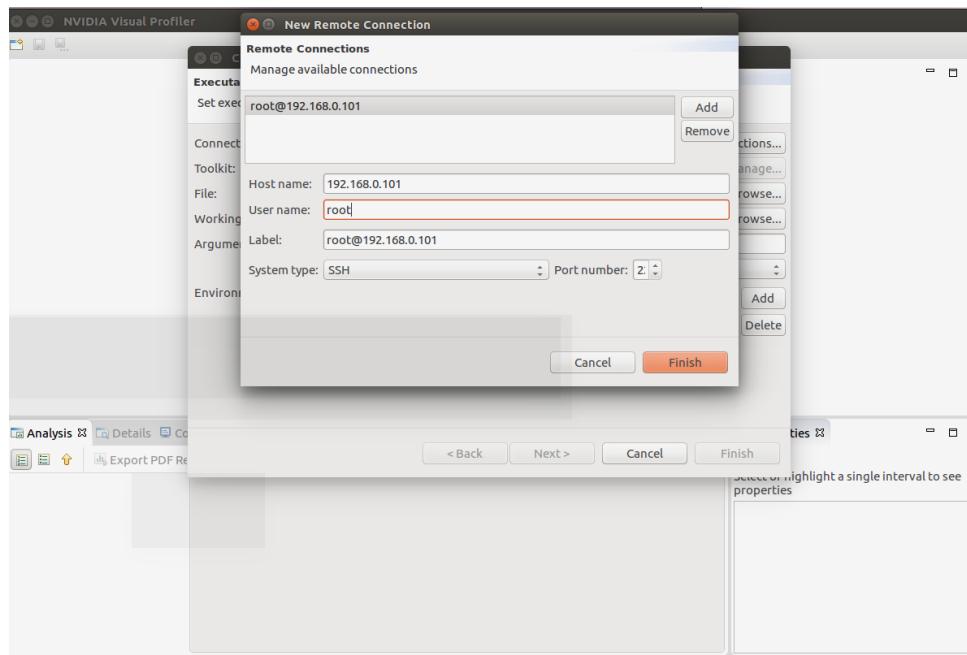
<https://docs.nvidia.com/cuda/profiler-users-guide/index.html>

Visual profiler

Make sure you have visual profiler install on the target by running the following command

```
nvvp -vm /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/java
```

Configure the connection



Configure application details:

Application:

./build/mono_tum

Working Directory:

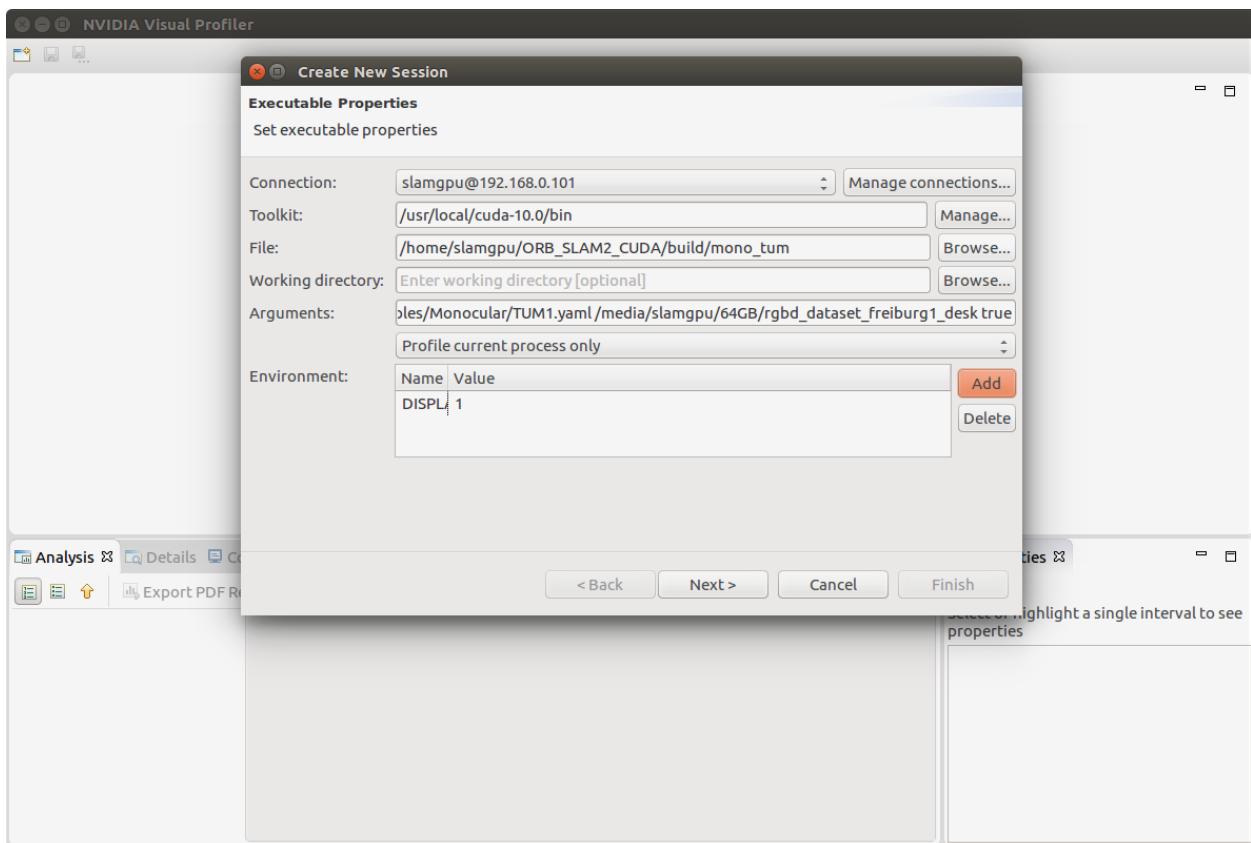
/home/slamgpu/ORB_SLAM2_CUDA

Application parameters:

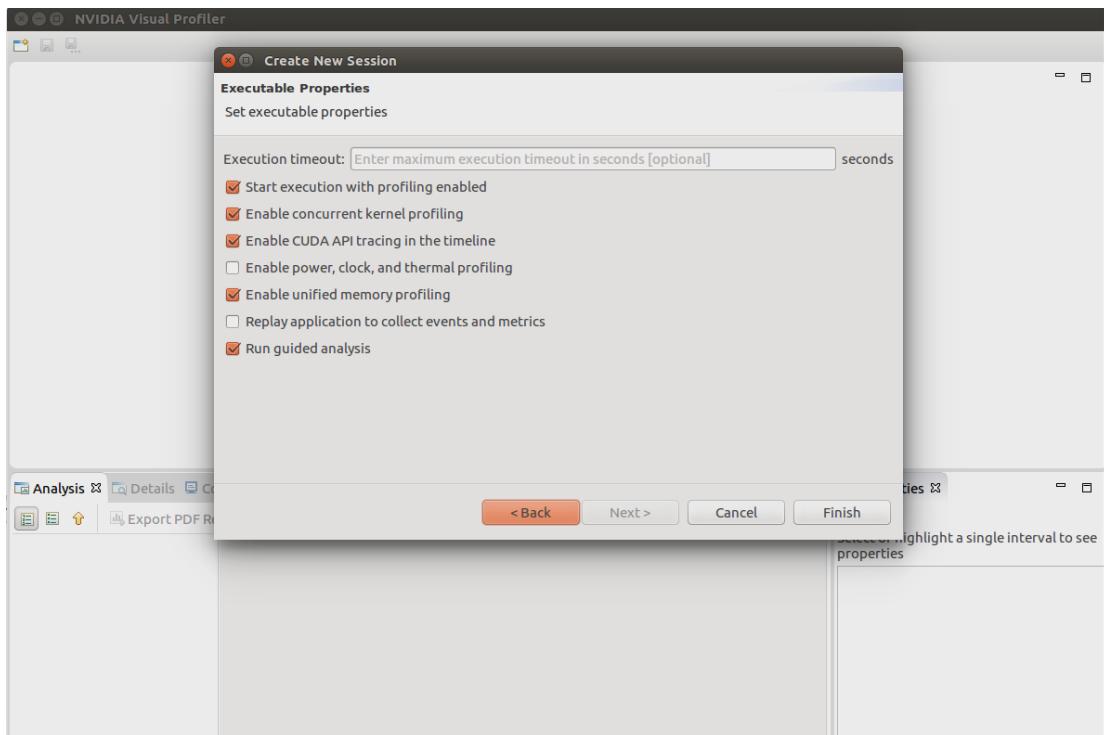
Vocabulary/ORBvoc.txt Examples/Monocular/TUM1.yaml /media/slamgpu/64GB/rgbd_dataset_freiburg1_desk true

environment

DISPLAY=:1



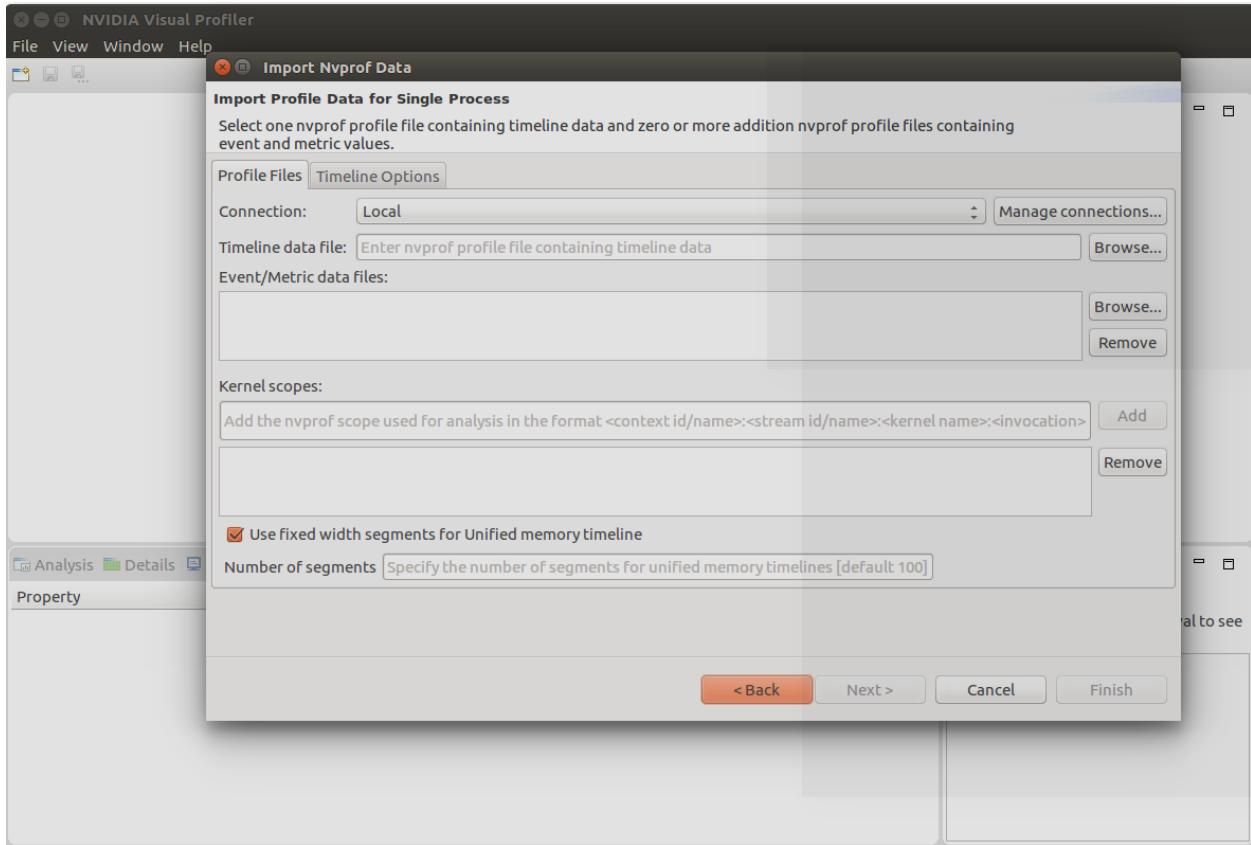
Click Next for the next window



Click finish.

You can start the application on the host side by also with the command
/usr/local/cuda-10.0/bin/nvvp

This command will allow you to load saved nvprof file.

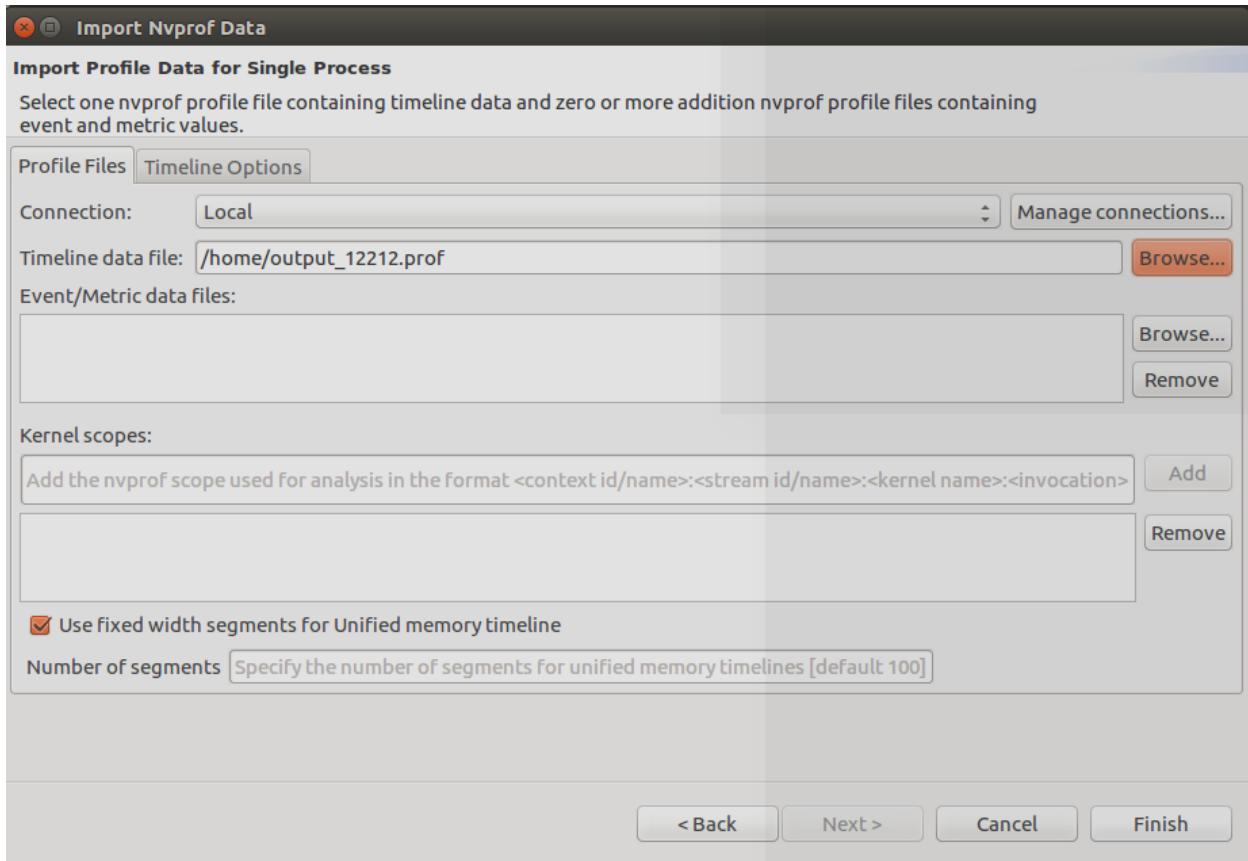


You can copy files from device by SCP

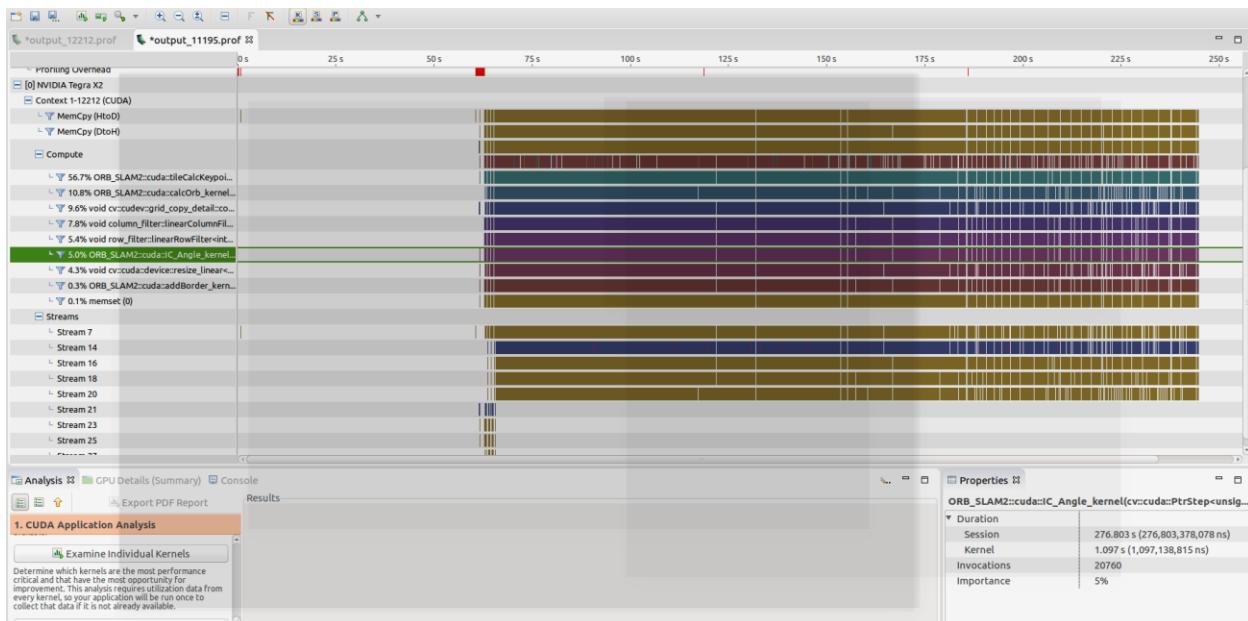
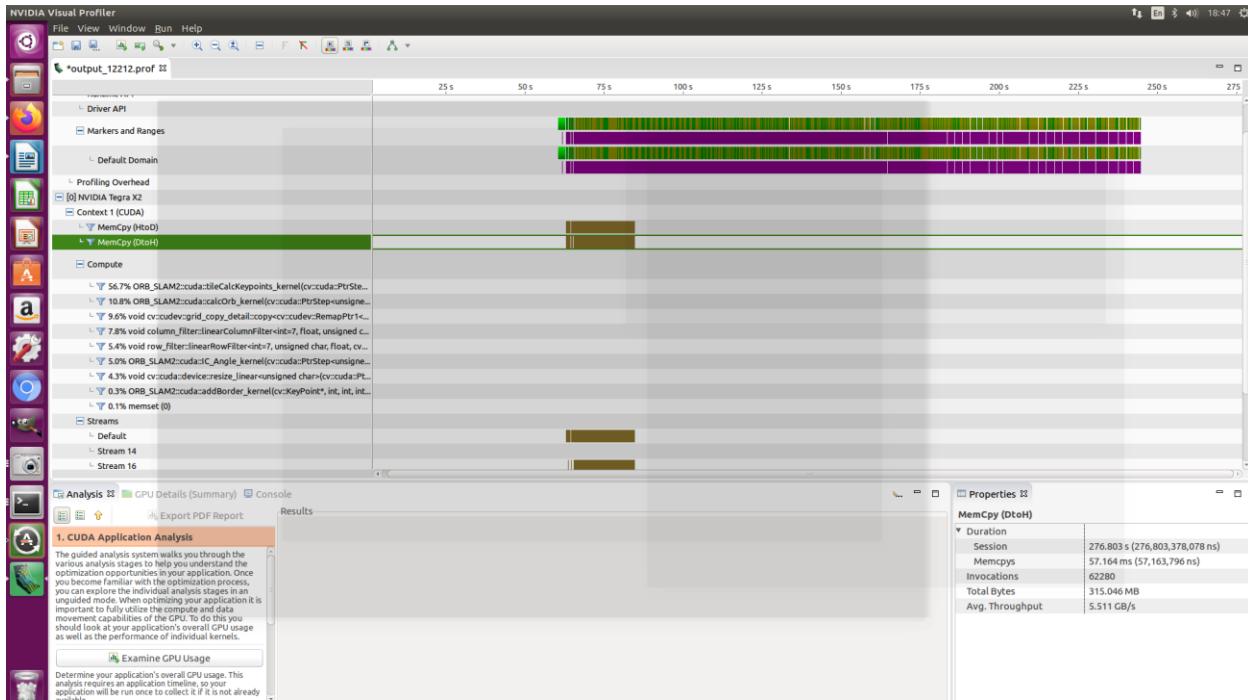
Example:

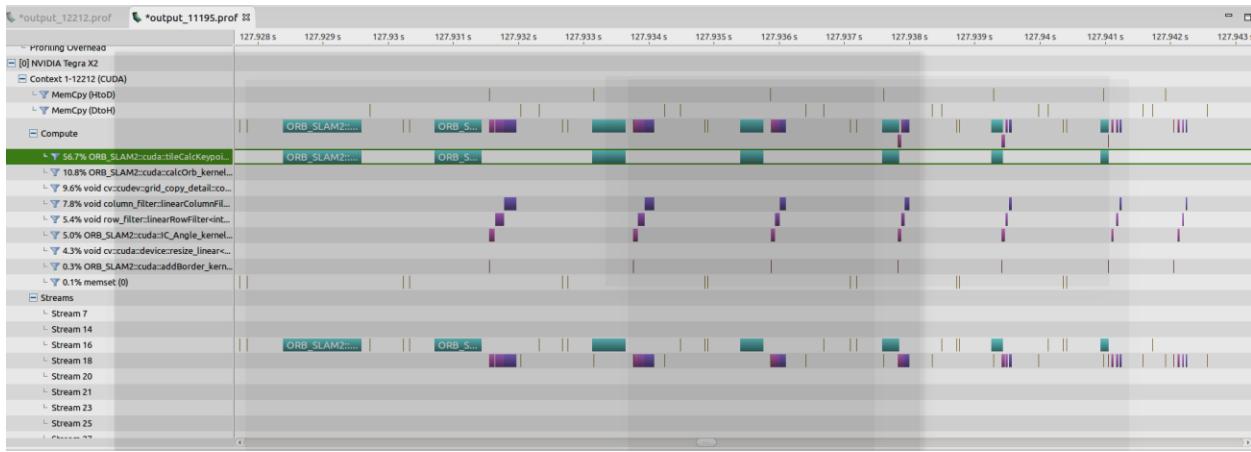
```
sudo scp root@192.168.0.101:/media/slampu/64GB/output_12212.prof  
/home/output_12212.prof
```

Fill the path to the file and then click Finish.

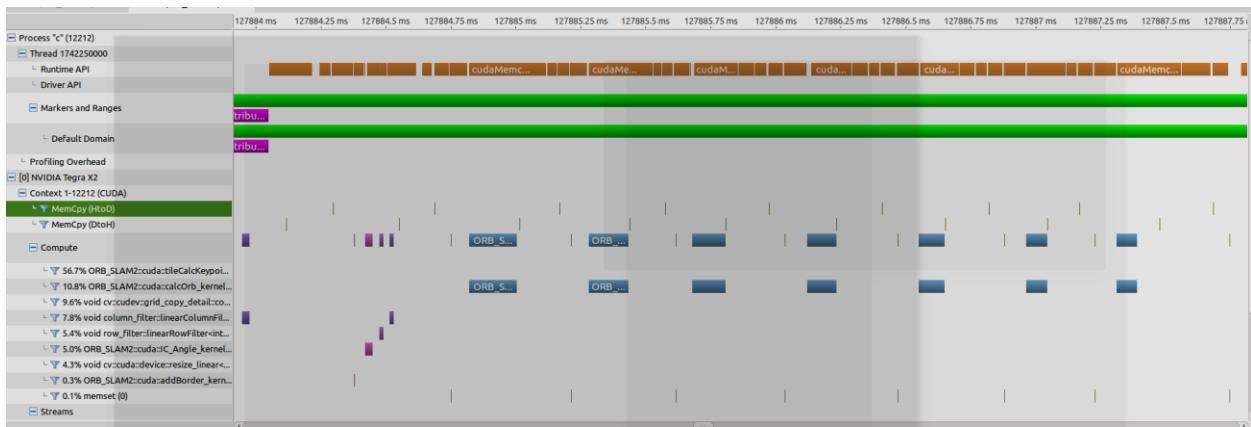


The tool will import and will show the timelines.





You can see the times for transfer data between host to device and vice versa.



Nsight eclipse edition

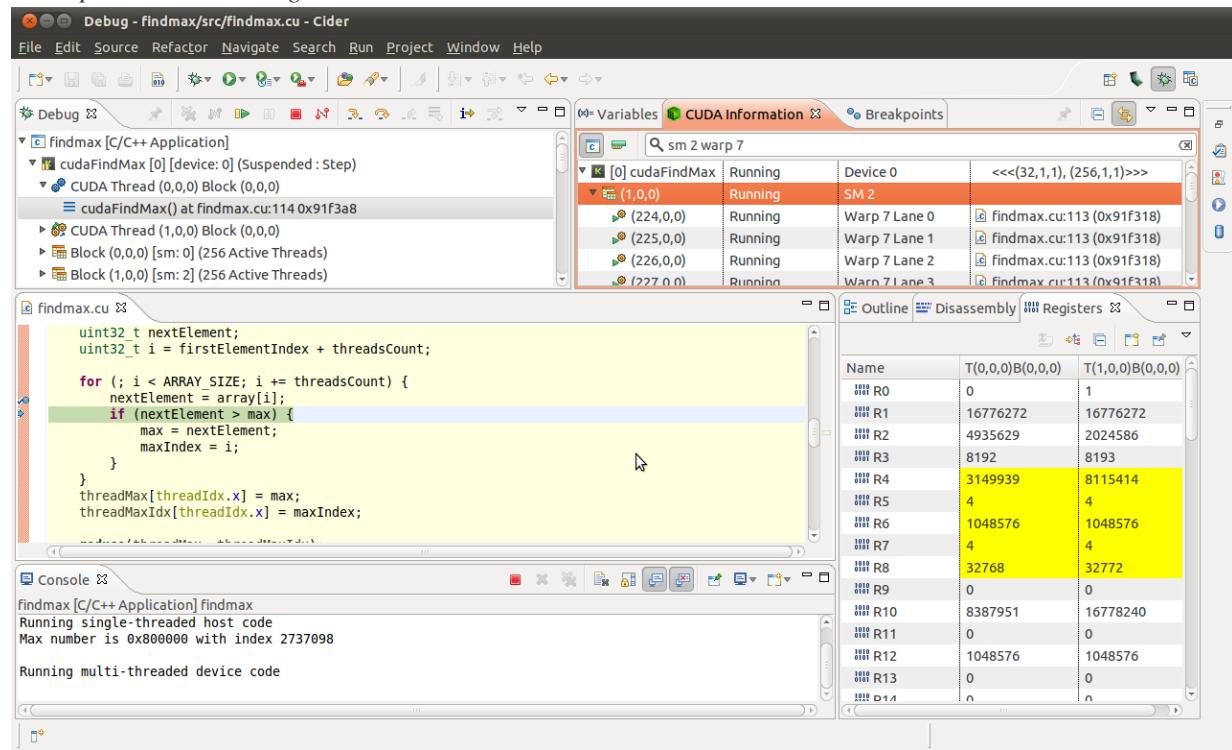
NVIDIA® Nsight™ Eclipse Edition is a unified CPU plus GPU integrated development environment (IDE) for developing CUDA® applications on Linux and Mac OS X for the x86, POWER and ARM platforms.

You can see information about the tool in the following path:

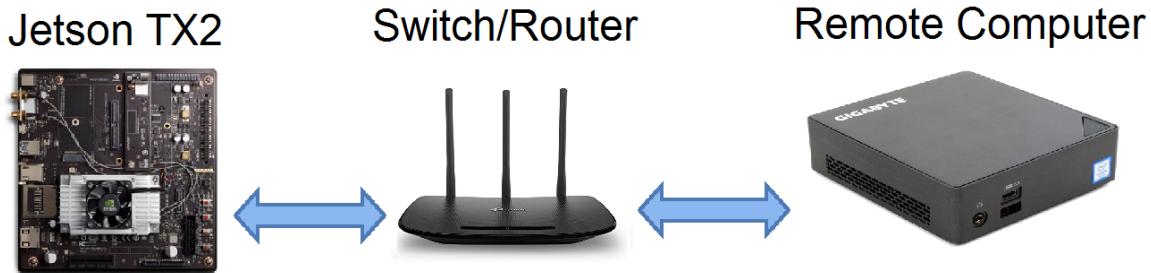
<https://docs.nvidia.com/cuda/nsight-eclipse-edition-getting-started-guide/index.html>

For manully tool installtion run the command:

sudo apt install nvidia-nsight



Network Configuration for Remote Profiling



I used TPLink Switch\Router.

- 1) Check device IP (Jetson) by command: ifconfig

```
root@slamgpu-desktop:~#
root@slamgpu-desktop:~# ifconfig
eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
        inet 192.168.0.101 netmask 255.255.255.0 broadcast 192.168.0.255
              inet6 fe80::ae0a:61fa:1c0b:23da prefixlen 64 scopeid 0x20<link>
                ether 00:04:4b:dd:8c:62 txqueuelen 1000 (Ethernet)
                  RX packets 6089 bytes 475646 (475.6 KB)
                  RX errors 0 dropped 0 overruns 0 frame 0
                  TX packets 11669 bytes 975562 (975.5 KB)
                  TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
                device interrupt 41

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
        inet 127.0.0.1 netmask 255.0.0.0
        inet6 ::1 prefixlen 128 scopeid 0x10<host>
          loop txqueuelen 1 (Local Loopback)
            RX packets 1729 bytes 128065 (128.0 KB)
            RX errors 0 dropped 0 overruns 0 frame 0
```

Internal interface for local connection eth0 with IP address 192.168.0.101

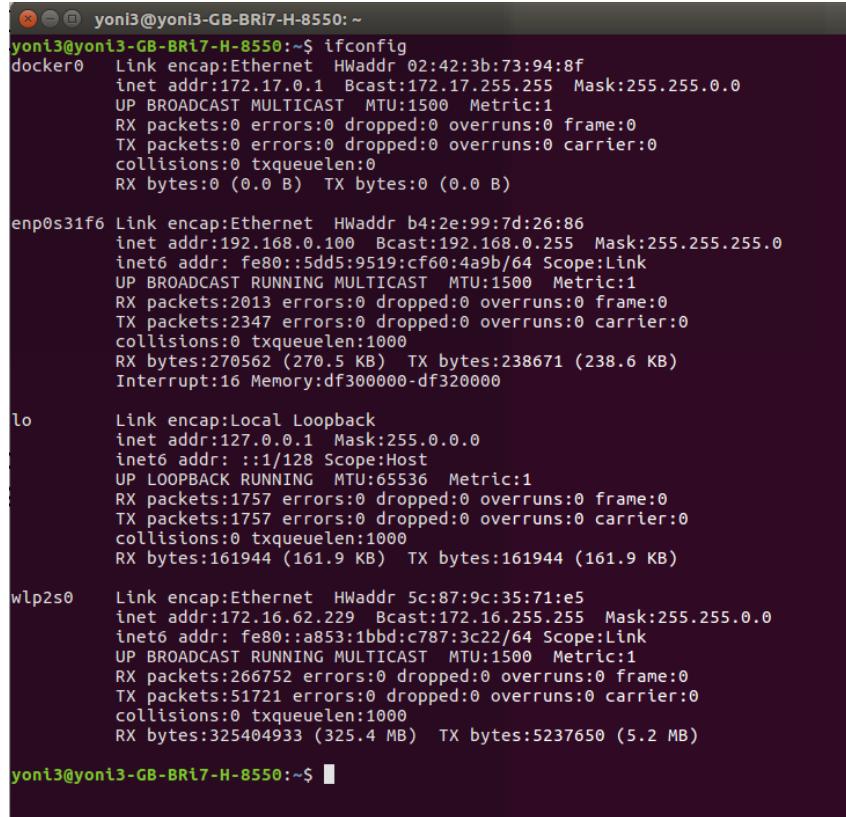
External interface for Wifi:

```
wlan0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
        inet 172.16.205.6 netmask 255.255.0.0 broadcast 172.16.255.255
              inet6 fe80::d7d1:43ce:6780:b374 prefixlen 64 scopeid 0x20<link>
                ether 00:04:4b:dd:8c:60 txqueuelen 1000 (Ethernet)
                  RX packets 51176 bytes 32003734 (32.0 MB)
                  RX errors 0 dropped 0 overruns 0 frame 0
                  TX packets 14495 bytes 1672565 (1.6 MB)
                  TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

2) Host side Configuration (Remote computer)

Make sure the internal interface netmask is the same as in device side.

In the current configuration internal interface IP address is 192.168.0.100



```
yoni3@yoni3-GB-BRi7-H-8550:~$ ifconfig
docker0  Link encap:Ethernet HWaddr 02:42:3b:73:94:8f
          inet addr:172.17.0.1 Bcast:172.17.255.255 Mask:255.255.0.0
              UP BROADCAST MULTICAST MTU:1500 Metric:1
              RX packets:0 errors:0 dropped:0 overruns:0 frame:0
              TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
              collisions:0 txqueuelen:0
              RX bytes:0 (0.0 B) TX bytes:0 (0.0 B)

enp0s31f6 Link encap:Ethernet HWaddr b4:2e:99:7d:26:86
          inet addr:192.168.0.100 Bcast:192.168.0.255 Mask:255.255.255.0
              inet6 addr: fe80::5dd5:9519:cf60:4a9b/64 Scope:Link
                  UP BROADCAST RUNNING MULTICAST MTU:1500 Metric:1
                  RX packets:2013 errors:0 dropped:0 overruns:0 frame:0
                  TX packets:2347 errors:0 dropped:0 overruns:0 carrier:0
                  collisions:0 txqueuelen:1000
                  RX bytes:270562 (270.5 KB) TX bytes:238671 (238.6 KB)
                  Interrupt:16 Memory:df300000-df320000

lo      Link encap:Local Loopback
          inet addr:127.0.0.1 Mask:255.0.0.0
              inet6 addr: ::1/128 Scope:Host
                  UP LOOPBACK RUNNING MTU:65536 Metric:1
                  RX packets:1757 errors:0 dropped:0 overruns:0 frame:0
                  TX packets:1757 errors:0 dropped:0 overruns:0 carrier:0
                  collisions:0 txqueuelen:1000
                  RX bytes:161944 (161.9 KB) TX bytes:161944 (161.9 KB)

wlp2s0  Link encap:Ethernet HWaddr 5c:87:9c:35:71:e5
          inet addr:172.16.62.229 Bcast:172.16.255.255 Mask:255.255.0.0
              inet6 addr: fe80::a853:1bbd:c787:3c22/64 Scope:Link
                  UP BROADCAST RUNNING MULTICAST MTU:1500 Metric:1
                  RX packets:266752 errors:0 dropped:0 overruns:0 frame:0
                  TX packets:51721 errors:0 dropped:0 overruns:0 carrier:0
                  collisions:0 txqueuelen:1000
                  RX bytes:325404933 (325.4 MB) TX bytes:5237650 (5.2 MB)

yoni3@yoni3-GB-BRi7-H-8550:~$
```

The default GW is 192.168.0.1

In our setup all the three on the same subnet mask 255.255.255.0

3) Device side configuration

The device is configured with the same network configuration as the host and configured with unique internal IP 192.168.0.101

CLONE JETSON TX2 INTO NEW SD CARD

- connect jetson via USB to host PC with jetpack (same version that used for installation)
- download the script described in this post:
<https://devtalk.nvidia.com/default/topic/1000105/jetson-tx2/tx2-cloning/post/5111893/#5111893>
to flash.sh file to the host PC.
- run with the following command on the host PC:
`sudo ./flash.sh -r -k APP -G my_backup.img jetson-tx2 mmcblk0p1`
- (not sure about this. Try to skip first) squash the new image on host with:
`sudo mksparse -v --fillpattern=0 system.raw system.img`
- rename the resulting file to be system.img and place it under the jetpack root filesystem that is used for the jetson installation. It should be under **bootloader**/ directory. Note that the old system.img file should be erased/renamed/moved.
- connect a new jetson to the host pc. (you might need to force it into emergency mode if not identified as emmc device in the host PC). run the following command to flash the **new** device:
`sudo ./flash.sh -r -k APP jetson-tx2 mmcblk0p1`

GPU PCIe Configuration Check

To can check which GPU you have on host by the Linux util commands:

<http://manpages.ubuntu.com/manpages/xenial/man8/lspci.8.html>

Detailed Device Information

`lspci -v`

Dump PCI Info in Different Format

`lspci -m`

Link for tools

<https://developer.nvidia.com/tools-overview>

Nsight compute

<https://docs.nvidia.com/nsight-compute/NsightCompute/index.html>

https://developer.nvidia.com/nsight-compute-2019_5

<https://docs.nvidia.com/nsight-graphics/UserGuide/index.html#getting started how to launch and connect>

- /usr/local/cuda-10.0/NsightCompute-1.0/nv-nsight-cu
- nv-nsight-gfx-for-l4t
- nsight-sys

How to run device SLAM application via SSH connection

Note : This application mode is ROS mode.

Open SSH connection to Device from the host

Check the DISPLAY id

```
slamgpu@slamgpu-desktop:~$  
slamgpu@slamgpu-desktop:~$  
slamgpu@slamgpu-desktop:~$  
slamgpu@slamgpu-desktop:~$  
slamgpu@slamgpu-desktop:~$  
slamgpu@slamgpu-desktop:~$  
slamgpu@slamgpu-desktop:~$  
slamgpu@slamgpu-desktop:~$  
slamgpu@slamgpu-desktop:~$  
slamgpu@slamgpu-desktop:~$ who  
slamgpu :1 2019-12-24 09:24 (:1)  
slamgpu pts/2 2019-12-24 09:33 (192.168.0.100)  
slamgpu@slamgpu-desktop:~$ █
```

In our case :1 is the Display id to run on the device

Change the the global varible to open the display in the device by

```
export DISPLAY=:1
```

In case of error as below :

```
ORB Extractor Parameters:  
- Number of Features: 1000  
- Scale Levels: 8  
- Scale Factor: 1.2  
- Initial Fast Threshold: 20  
- Minimum Fast Threshold: 7  
  
-----  
Start processing sequence ...  
Images in the sequence: 613  
  
error: XDG_RUNTIME_DIR not set in the environment.  
No protocol specified  
terminate called after throwing an instance of 'std::runtime_error'  
what(): Pangolin X11: Failed to open X display  
Aborted  
root@slamgpu-desktop:/home/slampu/ORB_SLAM2_CUDA# █
```

```
export XDG_RUNTIME_DIR=/run/user/1000
```

run the ROS application in one SSH terminal

```
roslaunch  
/home/slamgpu/ORB_SLAM2_CUDA/Examples/ROS/ORB_SLAM2_CUDA/launch/ros_mono.launch
```

Then run in second SSH terminal (new terminal) the bag script

```
rosbag play /media/slamgpu/64GB/MH_01_easy_orig.bag
```

Solutions for issues during SLAM ROS application

1) Error during running ROS:

Failed to load module "canberra-gtk-module"

```
Loading ORB Vocabulary. This could take a while...
Vocabulary loaded!

Camera Parameters:
- fx: 435.205
- fy: 435.205
- cx: 367.452
- cy: 252.201
- k1: -0.283408
- k2: 0.0739591
- p1: 0.00019359
- p2: 1.76187e-05
- fps: 20
- color order: RGB (ignored if grayscale)

ORB Extractor Parameters:
- Number of Features: 2000
- Scale Levels: 8
- Scale Factor: 1.2
- Initial Fast Threshold: 20
- Minimum Fast Threshold: 7
Gtk-Message: 09:54:56.958: Failed to load module "canberra-gtk-module"
```

Fix the issue by:

```
sudo apt install libcanberra-gtk-module libcanberra-gtk3-module
```

- 2) Issue during running ROS application:
"unable to find package: image transport"

```

/home/slamgpu/ORB_SLAM2_CUDA/Examples/ROS/ORB_SLAM2_CUDA/launch/ros_mono.launch http://localhost:11311
/home/slamgpu/ORB_SLAM2_CUDA/Examples/ROS/ORB_SLAM2_CUDA/launch/ros_mono.launch http://localhost:11311

Camera Parameters:
- fx: 435.205
- fy: 435.205
- cx: 367.452
- cy: 252.201
- k1: -0.283408
- k2: 0.0739591
- p1: 0.00019359
- p2: 1.761876e-05
- fps: 20
- color order: RGB (ignored if grayscale)

ORB Extractor Parameters:
- Number of Features: 2000
- Scale Levels: 8
- Scale Factor: 1.2
- Initial Fast Threshold: 20
- Minimum Fast Threshold: 7
[rospack] Error: package 'image_transport' not found
[librospace]: error while executing command
terminate called after throwing an instance of 'pluginlib::ClassLoaderException'
what(): Unable to find package: image_transport

```

Solution:

Search for the relevant package :

```

slamgpu@slamgpu-desktop:/home$ sudo apt-get install ros-*~image-transport
Reading package lists... Done
Building dependency tree
Reading state information... Done
Note, selecting 'ros-melodic-compressed-depth-image-transport' for glob 'ros-*~image-transport'
Note, selecting 'ros-melodic-theora-image-transport' for glob 'ros-*~image-transport'
Note, selecting 'ros-melodic-imagezero-image-transport' for glob 'ros-*~image-transport'
Note, selecting 'ros-melodic-resized-image-transport' for glob 'ros-*~image-transport'
Note, selecting 'ros-melodic-codec-image-transport' for glob 'ros-*~image-transport'
Note, selecting 'ros-melodic-compressed-image-transport' for glob 'ros-*~image-transport'
Note, selecting 'ros-melodic-image-transport' for glob 'ros-*~image-transport'

```

By command:

```

sudo apt-get install ros-*~image-transport
slamgpu@slamgpu-desktop:/home$ sudo apt-get install ros-melodic-image-transport
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following packages will be upgraded:
  ros-melodic-image-transport
1 upgraded, 0 newly installed, 0 to remove and 587 not upgraded.
Need to get 210 kB of archives.
After this operation, 0 B of additional disk space will be used.
Get:1 http://packages.ros.org/ros/ubuntu bionic/main arm64 ros-melodic-image-transport arm64 1.11.13-0bionic.20200321.033016 [210 kB]
Fetched 210 kB in 8s (26.9 kB/s)
debconf: delaying package configuration, since apt-utils is not installed
(Reading database ... 187044 files and directories currently installed.)
Preparing to unpack .../ros-melodic-image-transport_1.11.13-0bionic.20200321.033016_arm64.deb ...
Unpacking ros-melodic-image-transport (1.11.13-0bionic.20200321.033016) over (1.11.13-0bionic.20191009.141759) ...
Setting up ros-melodic-image-transport (1.11.13-0bionic.20200321.033016) ...

```

In our case ROS version is melodic so I used the command:

sudo apt-get install ros-melodic-image-transport

Compilation Cuda Application with Debug Symbols

Compile Cuda example on Jetson TX2 with debug symbols:

```
sudo sh -c "dbg=1 make"
```

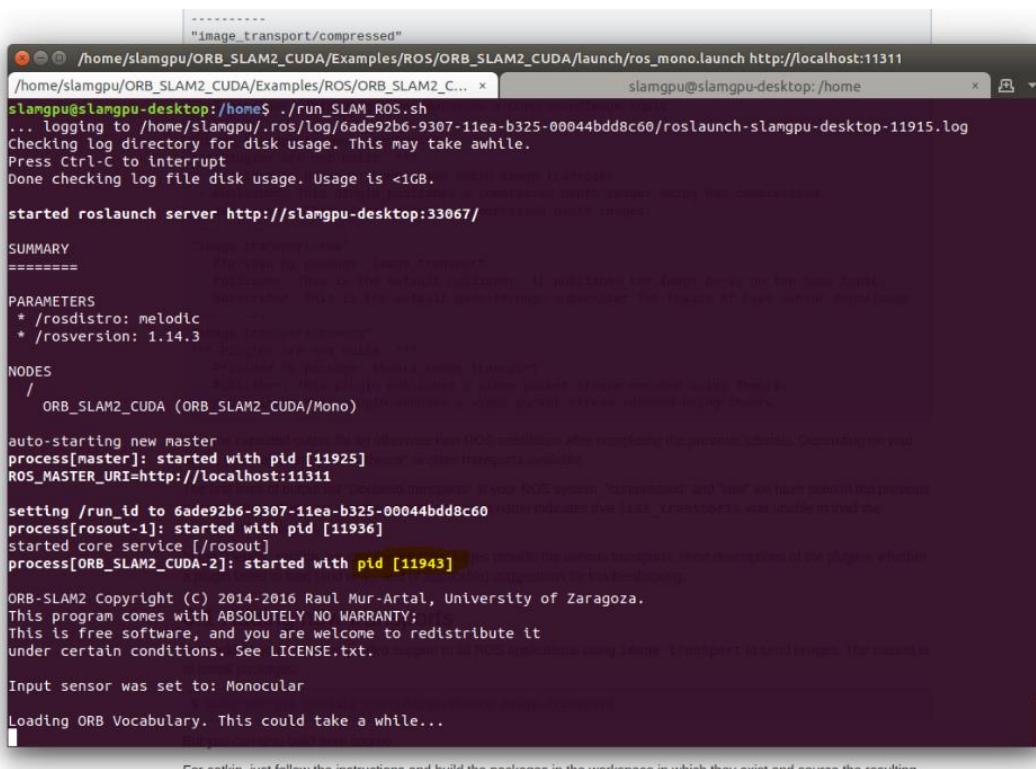
Example

```
slamgpu@slamgpu-desktop:/usr/local/cuda-10.0/samples/0_Simple/vectorAdd$ sudo sh -c "dbg=1
make"/usr/local/cuda-10.0/bin/nvcc -ccbin g++ -I../../common/inc -m64 -g -G -gencode
arch=compute_30,code=sm_30 -gencode arch=compute_32,code=sm_32 -gencode arch=compute_53,code=sm_53 -
gencode arch=compute_61,code=sm_61 -gencode arch=compute_62,code=sm_62 -gencode
arch=compute_70,code=sm_70 -gencode arch=compute_72,code=sm_72 -gencode arch=compute_75,code=sm_75 -
gencode arch=compute_75,code=compute_75 -o vectorAdd.o -c vectorAdd.cu
/usr/local/cuda-10.0/bin/nvcc g++ -m64 -g -G -gencode arch=compute_30,code=sm_30 -gencode
arch=compute_32,code=sm_32 -gencode arch=compute_53,code=sm_53 -gencode arch=compute_61,code=sm_61 -
gencode arch=compute_62,code=sm_62 -gencode arch=compute_70,code=sm_70 -gencode
arch=compute_72,code=sm_72 -gencode arch=compute_75,code=sm_75 -gencode
arch=compute_75,code=compute_75 -o vectorAdd vectorAdd.o
mkdir -p ../../bin/aarch64/linux/debug
cp vectorAdd ../../bin/aarch64/linux/debug
```

Nsight System

SLAM application in ROS mode

In the following tutorial I run the SLAM application in ROS mode and I connect the profiling tool (Nsight System) to process id (PID).



```
-----  
"image_transport/compressed"  
/home/slamgpu/ORB_SLAM2_CUDA/Examples/ROS/ORB_SLAM2_CUDA/launch/ros_mono.launch http://localhost:11311  
/home/slamgpu/ORB_SLAM2_CUDA/Examples/ROS/ORB_SLAM2_C... x slamgpu@slamgpu-desktop:/home x   
slamgpu@slamgpu-desktop:~/home$ ./run_SLAM_ROS.sh  
... logging to /home/slamgpu/.ros/log/6ade92b6-9307-11ea-b325-00044bdd8c60/roslaunch-slamgpu-desktop-11915.log  
Checking log directory for disk usage. This may take awhile.  
Press Ctrl-C to interrupt  
Done checking log file disk usage. Usage is <1GB.  
started roslaunch server http://slamgpu-desktop:33067/  
  
SUMMARY  
=====  
PARAMETERS  
* /rosdistro: melodic  
* /rosversion: 1.14.3  
  
NODES  
/  
    ORB_SLAM2_CUDA (ORB_SLAM2_CUDA/Mono)  
  
auto-starting new master  
process[master]: started with pid [11925]  
ROS_MASTER_URI=http://localhost:11311  
  
setting /run_id to 6ade92b6-9307-11ea-b325-00044bdd8c60  
process[rosout-1]: started with pid [11936]  
started core service [/rosout]  
process[ORB_SLAM2_CUDA-2]: started with pid [11943]  
  
ORB-SLAM2 Copyright (C) 2014-2016 Raul Mur-Artal, University of Zaragoza.  
This program comes with ABSOLUTELY NO WARRANTY;  
This is free software, and you are welcome to redistribute it  
under certain conditions. See LICENSE.txt.  
  
Input sensor was set to: Monocular  
Loading ORB Vocabulary. This could take a while...  
[ 0% ]  
  
For certain instructions and build the workspace in which they exist and source the resulting
```

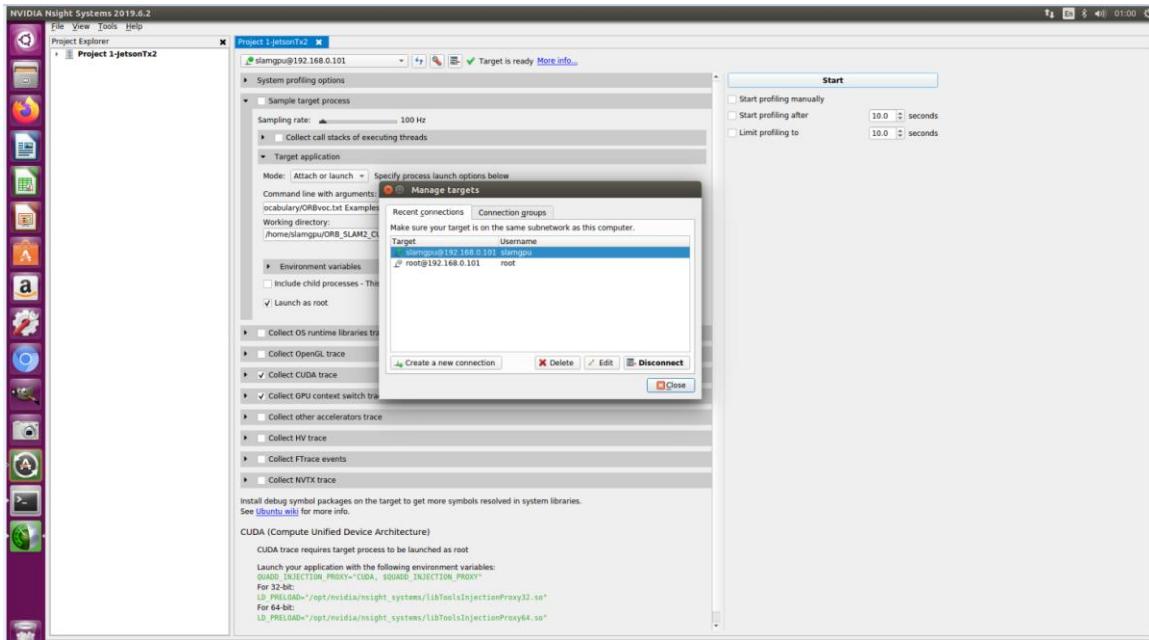
1) Open the tool on the host by the command :
`nsight-sys`

2) Open new project

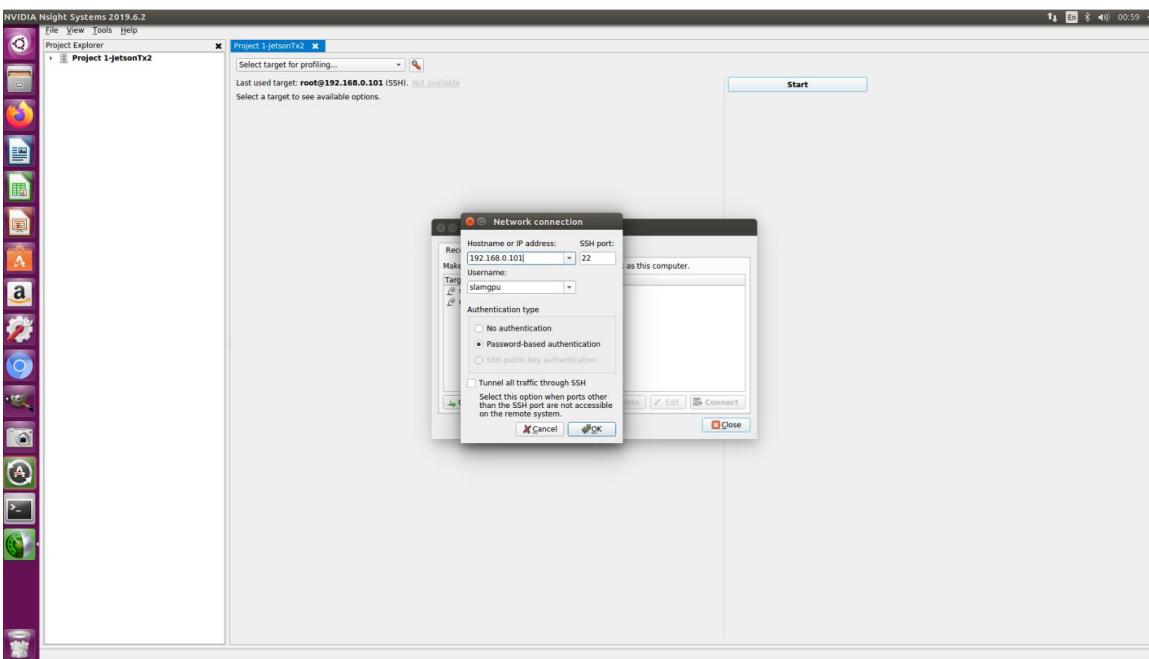
3) Configure SSH connection to device

Create new connection regular user or root.

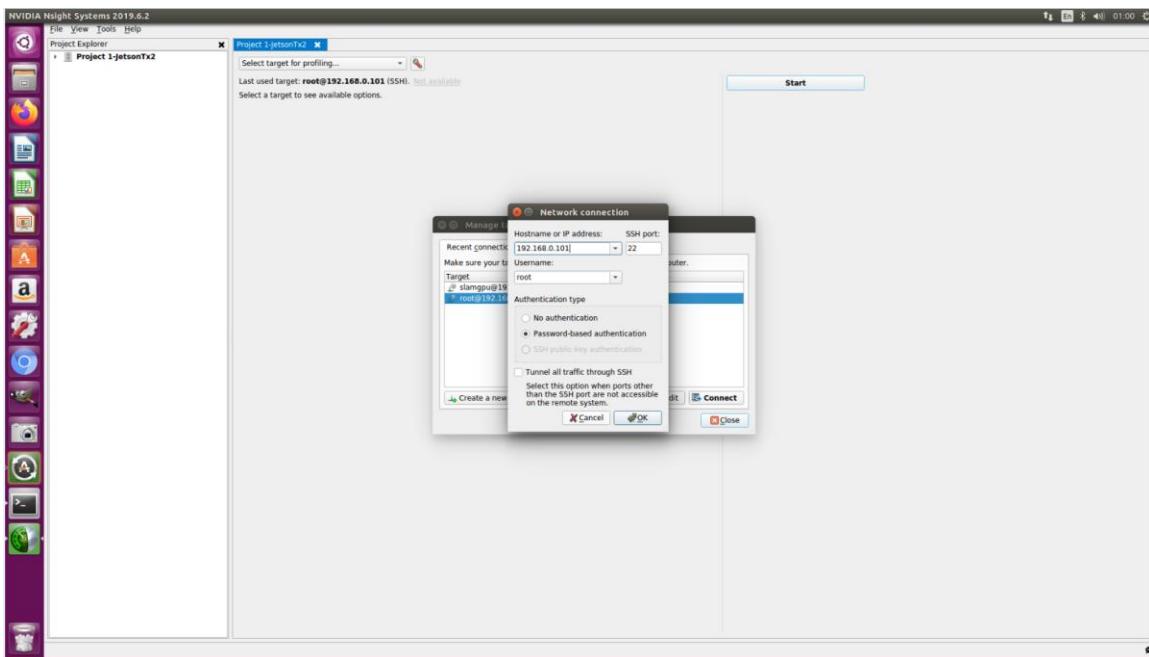
Fill IP address and user name for login to the device.



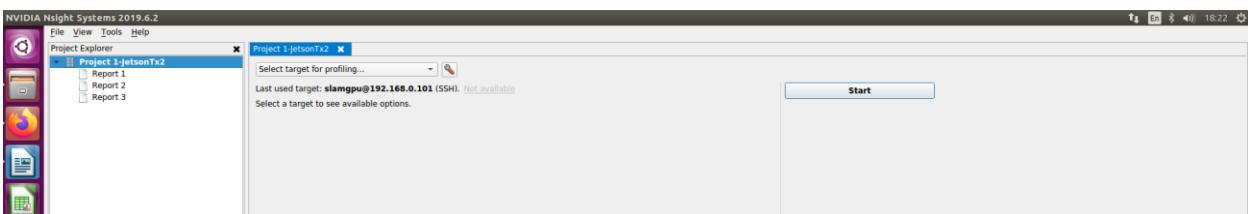
User: `slamgpu`



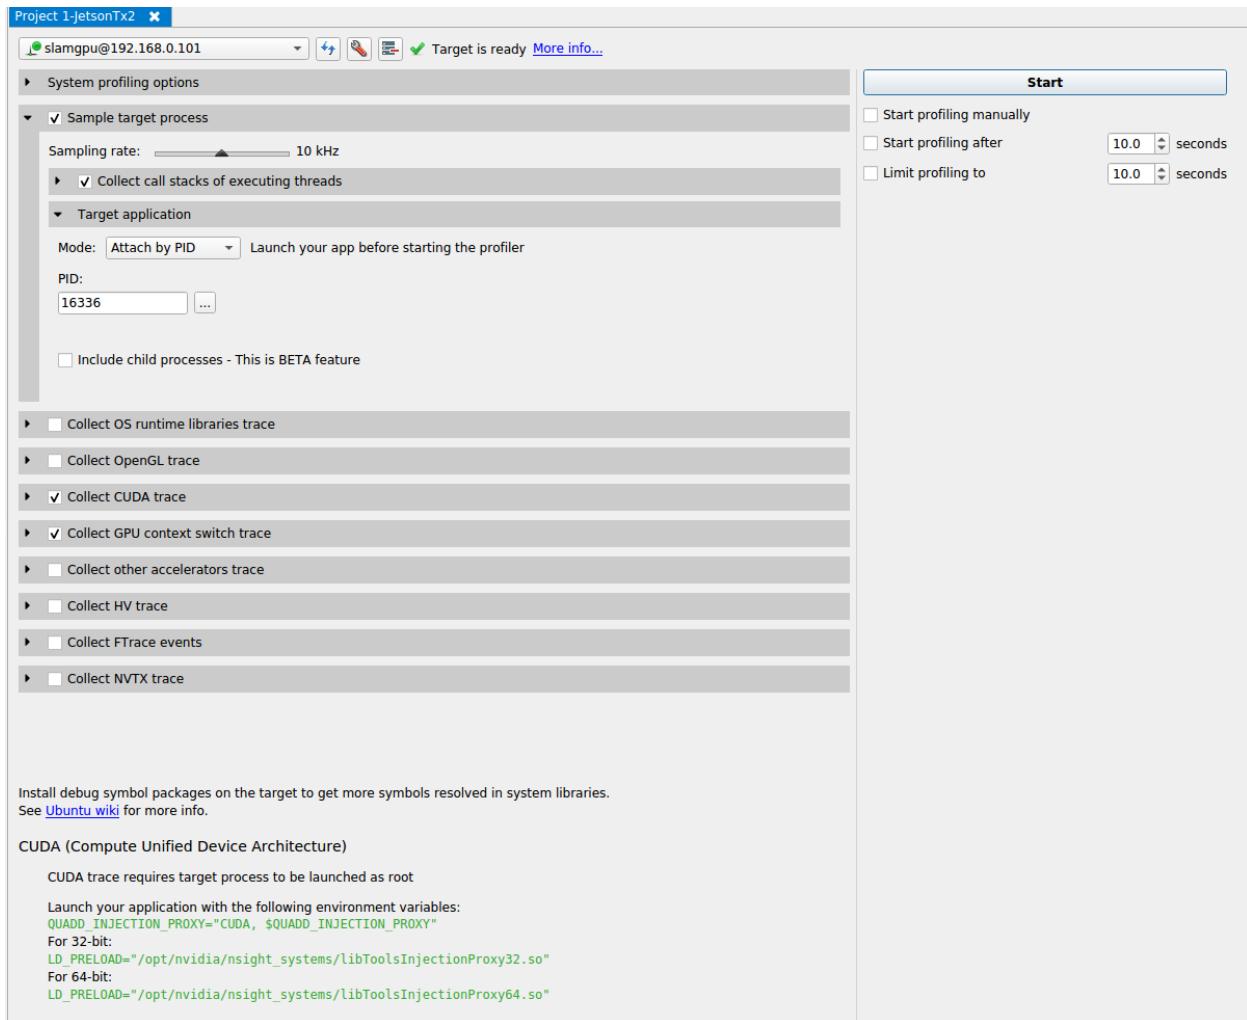
User: `Root`



- 4) Select the remote target (User or Root) and click on Start

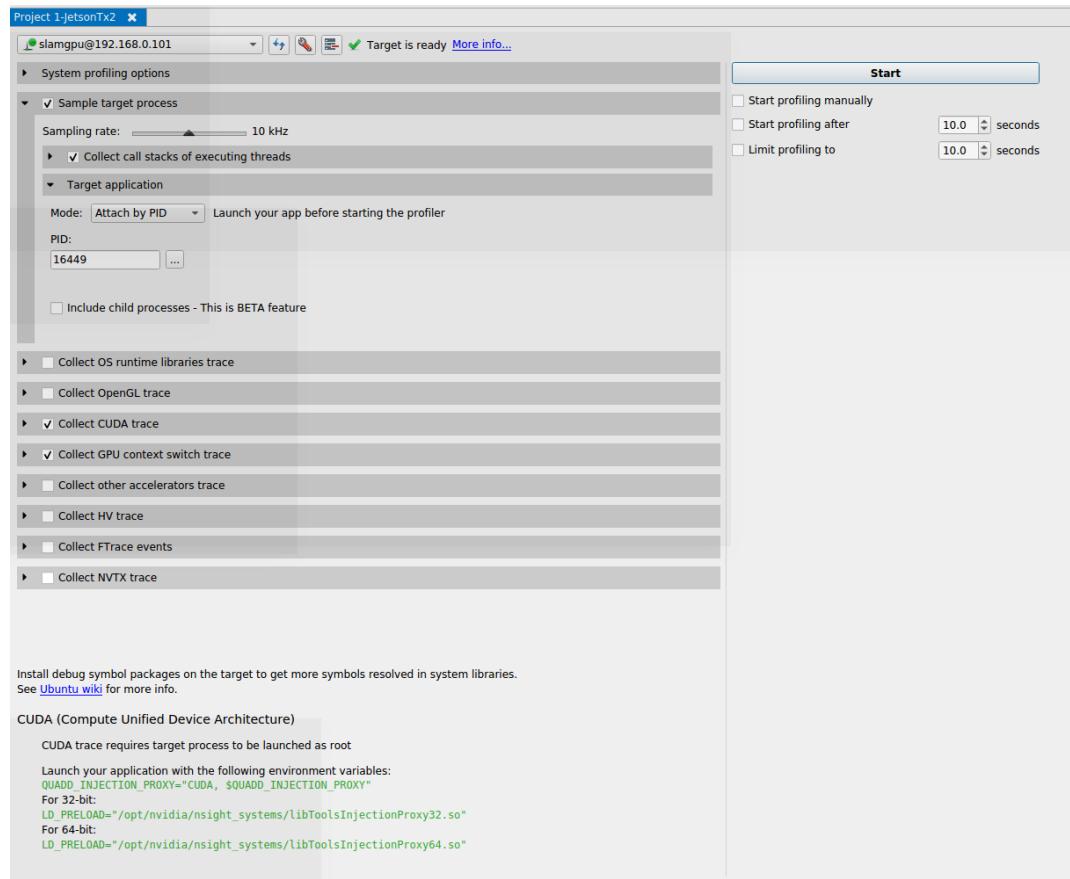


5) Configure the profiling options

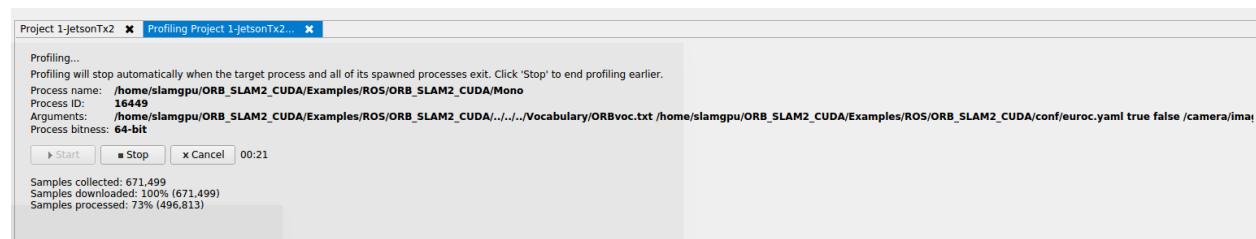


6) Run the application in the device side and configure the process id to the nsight system

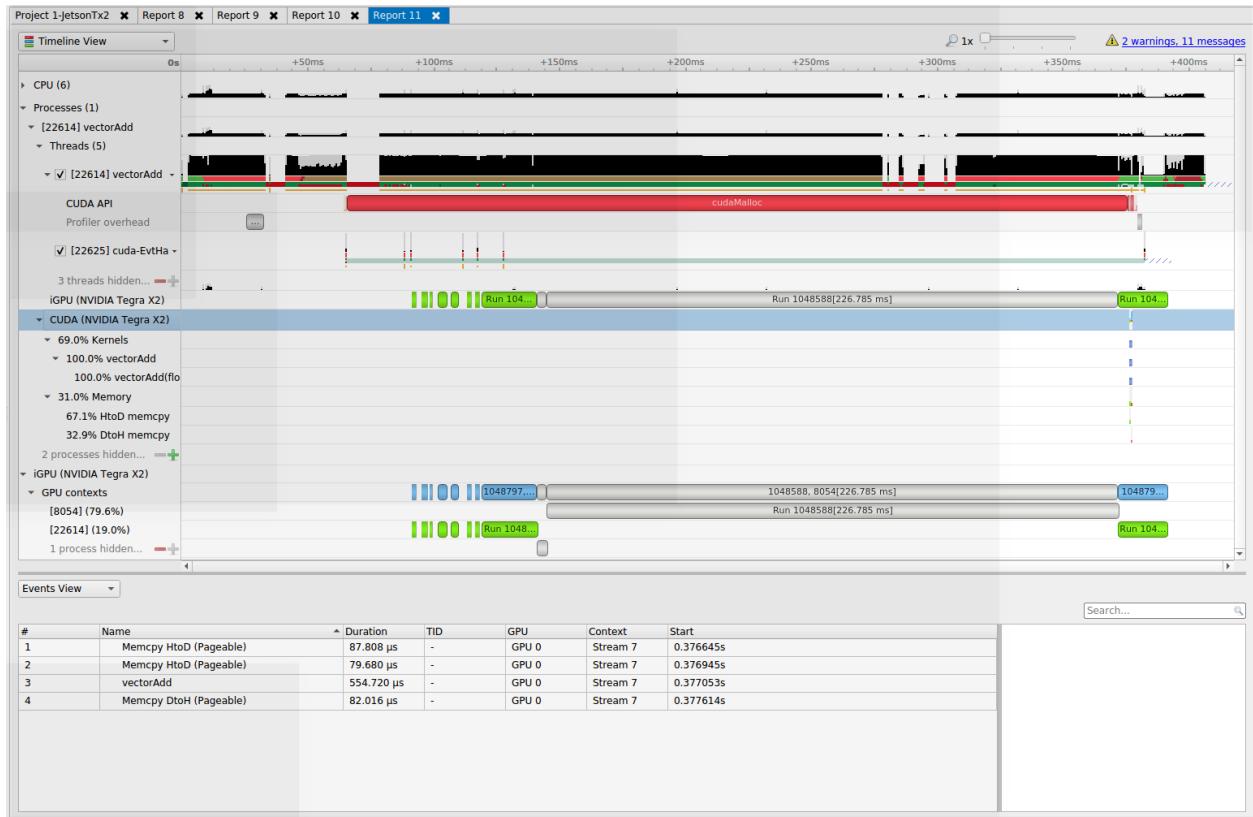
In the below image I used PID 16449



Then click on start



When finish you could and analyze the results.



For capture Cuda Information better use :

CUDA trace requires target process to be launched as root

Launch your application with the following environment variables:

```
QUADD_INJECTION_PROXY="CUDA, $QUADD_INJECTION_PROXY"
```

For 32-bit:

```
LD_PRELOAD="/opt/nvidia/nsight_systems/libToolsInjectionProxy32.so"
```

For 64-bit:

```
LD_PRELOAD="/opt/nvidia/nsight_systems/libToolsInjectionProxy64.so"
```

By adding the defines to bashrc you can reload by the command:

```
source ~/.bashrc
```

Nsight System

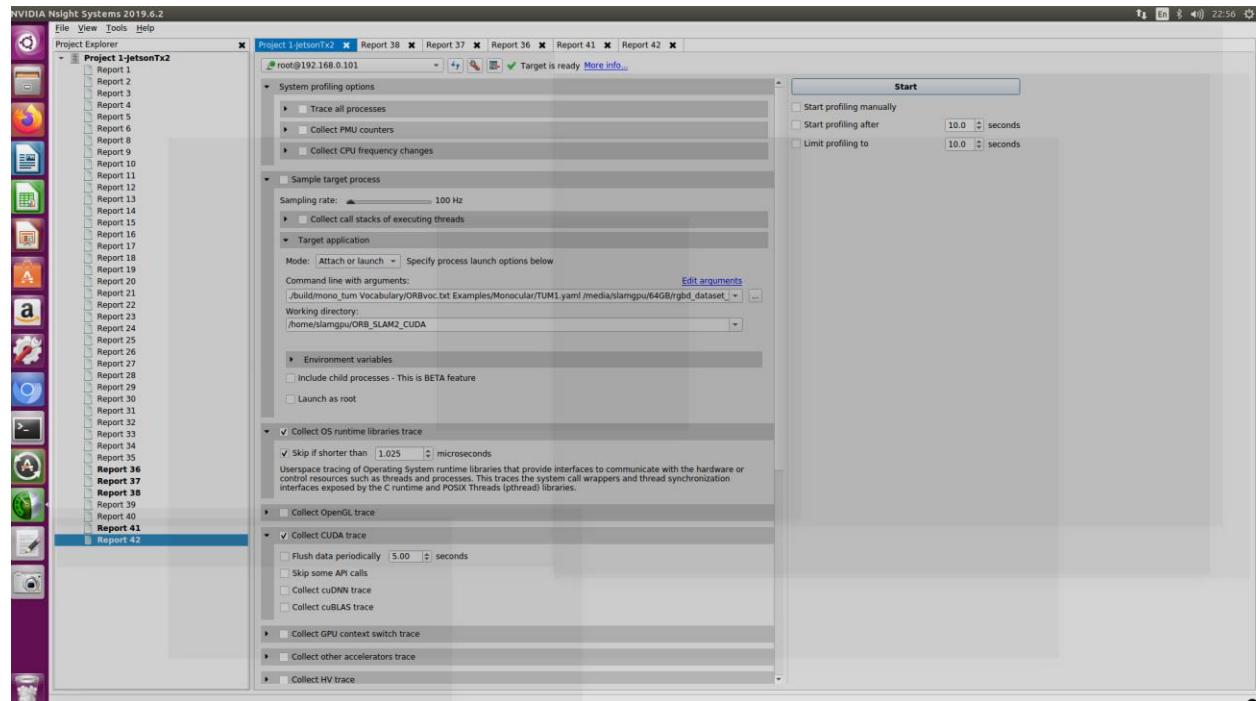
SLAM application in Standalone Mode

This mode is preferred in our research as the SLAM application is run in root mode.

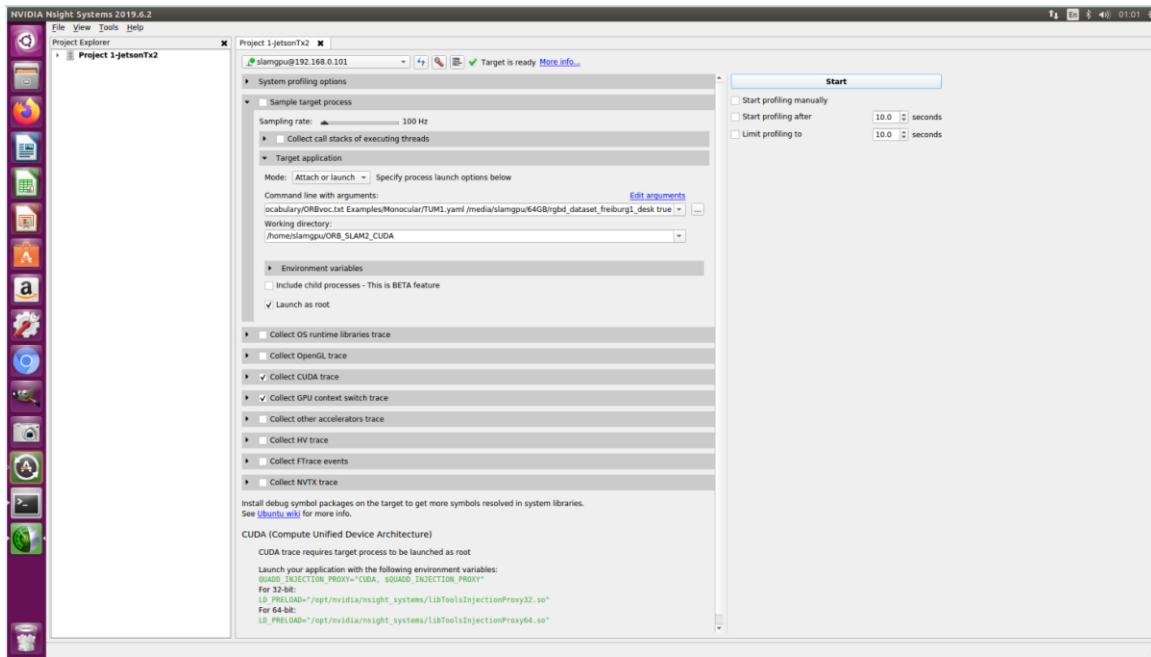
You can launch the tool by command:

```
nsight-sys
```

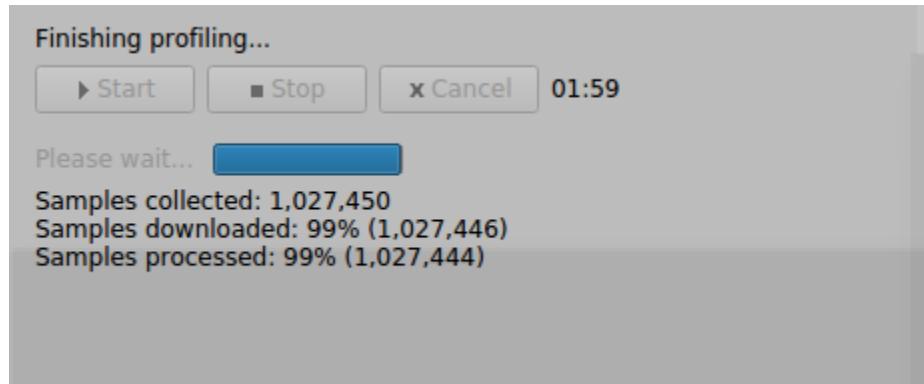
Configure connection and application parameters with root permissions:



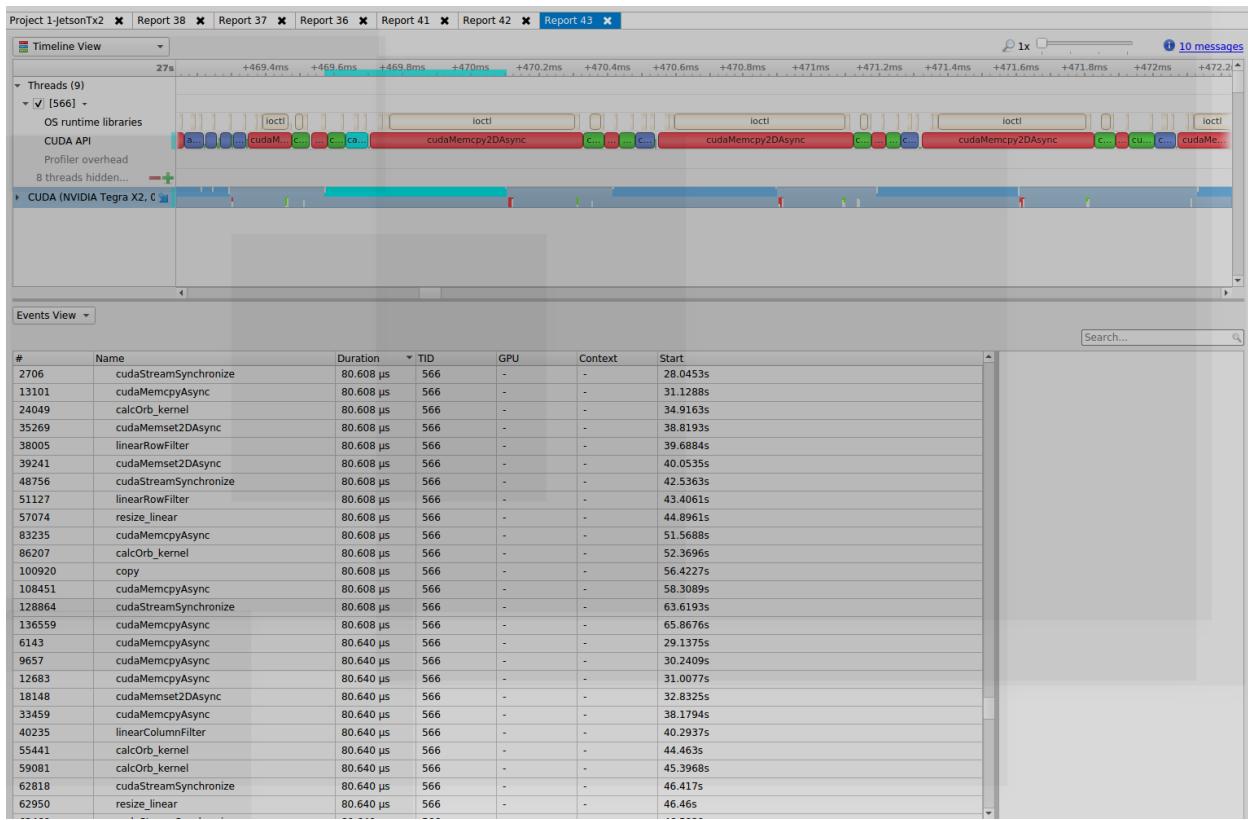
Preferred to click on Launch as root to collect more information:



Click start for launch the session then waiting for parsing



When processing is finished you can start analyze the results:



You can see the Cuda profiling done correctly by

Messages			
Source	Target, Process	Time	Description
Analysis		00:00:005	Profiling has started.
Daemon	slamgpu@192.168.0.101 (:1:0:14186)	00:00:006	Process was launched by the profiler, see /tmp/nvidia/insight_systems/streams/pid_14186_stdout.log and stderr.log for program output
Daemon	slamgpu@192.168.0.101 (:1:0:14186)	00:00:007	Profiler attached to the process.
Injection	slamgpu@192.168.0.101 (:1:0:14186)	00:00:096	Common injection library initialized successfully.
Injection	slamgpu@192.168.0.101 (:1:0:14186)	00:00:155	OS runtime libraries injection initialized successfully.
Injection	slamgpu@192.168.0.101 (:1:0:14186)	00:23:784	CUDA injection initialized successfully.
Injection	slamgpu@192.168.0.101 (:1:0:14186)	01:08:394	Number of CUPTI events produced: 232934, CUPTI buffers: 20.
Analysis		03:37:780	Profiling has stopped.
Analysis	slamgpu@192.168.0.101 (:1:0:14186)	05:42:041	Number of CUDA events collected: 217,585.
Analysis	slamgpu@192.168.0.101 (:1:0:14186)	05:42:041	Number of OS runtime libraries events collected: 4,872,211.
Analysis	slamgpu@192.168.0.101 (:1:0)	05:42:041	Number of GPU context switch events collected: 226,433.

You can see application output parameters:

The screenshot shows the NVIDIA Insight Systems 2019.6.2 software interface. On the left is the Project Explorer, which lists multiple reports from a project named "Project 1-jetsonTx2". The right side of the window is a log file viewer displaying the contents of "/tmp/nvidia/insight_systems/streams/pid_15843_stdout.log". The log file contains detailed configuration parameters for an ORB-SLAM2 system, including camera parameters (fx=537.206, fy=516.469, cx=355.443, cy=255.314, k1=0.262383, k2=0.112004, k3=1.16331, p1=-0.000058, p2=0.002428, -tp: 50), ORB Extractor parameters (Number of Features: 1000, Scale Levels: 8, Scale Factor: 1.2, Initial Fast Threshold: 20, Minimum Fast Threshold: 7), and the start of a processing sequence with 77 points. The log ends with a new map created with 100 points.

```
ORB-SLAM2 Copyright (C) 2014-2016 Raul Mur-Artal, University of Zaragoza.
This program comes with ABSOLUTELY NO WARRANTY;
This program is free software: you are free to redistribute it
under certain conditions. See LICENSE.txt.

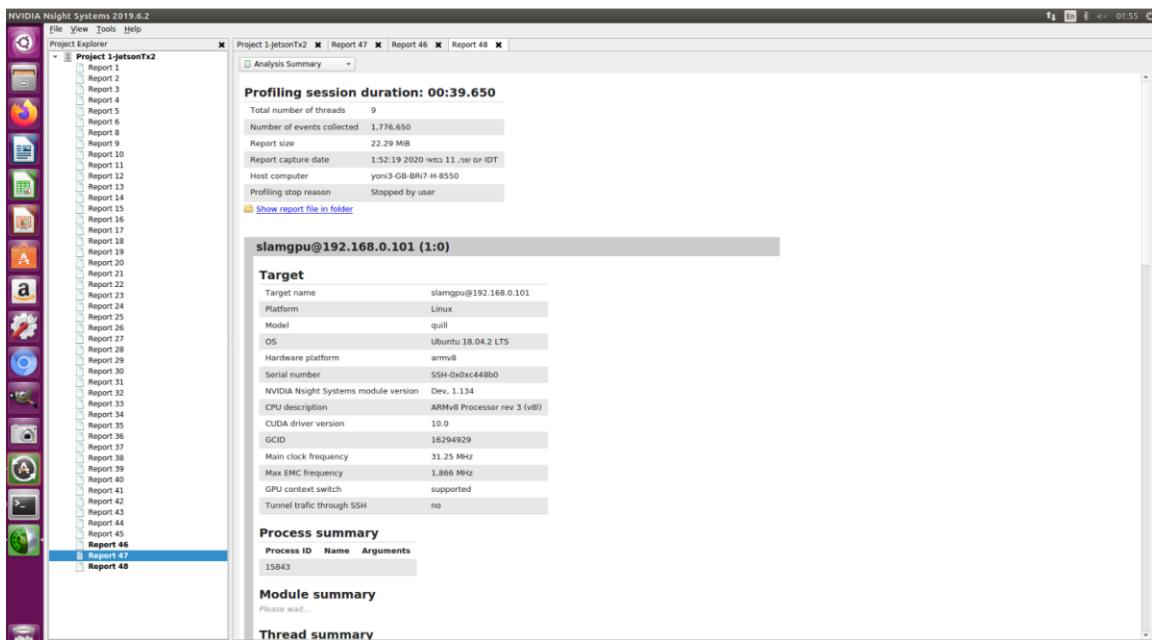
Input sensor was set to: Monocular
Loading ORB Vocabulary. This could take a while...
Vocabulary loaded!

Camera Parameters:
-fx: 537.206
-fy: 516.469
-cx: 355.443
-cy: 255.314
-k1: 0.262383
-k2: 0.112004
-k3: 1.16331
-p1: -0.000058
-p2: 0.002428
-tp: 50
-colorOrder: RGB (ignored if grayscale)

ORB Extractor Parameters:
- Number of Features: 1000
- Scale Levels: 8
- Scale Factor: 1.2
- Initial Fast Threshold: 20
- Minimum Fast Threshold: 7

Start processing sequence ...
Images in the sequence: 613
New Map created with 77 points
Wrong initialization, resetting...
System Resetting
Resetting Local Mapper... done
Resetting Loop Closing... done
Resetting Database... done
New Map created with 95 points
Wrong initialization, resetting...
System Resetting
Resetting Local Mapper... done
Resetting Loop Closing... done
Resetting Database... done
New Map created with 100 points
```

You can see high level information about the profiling session:



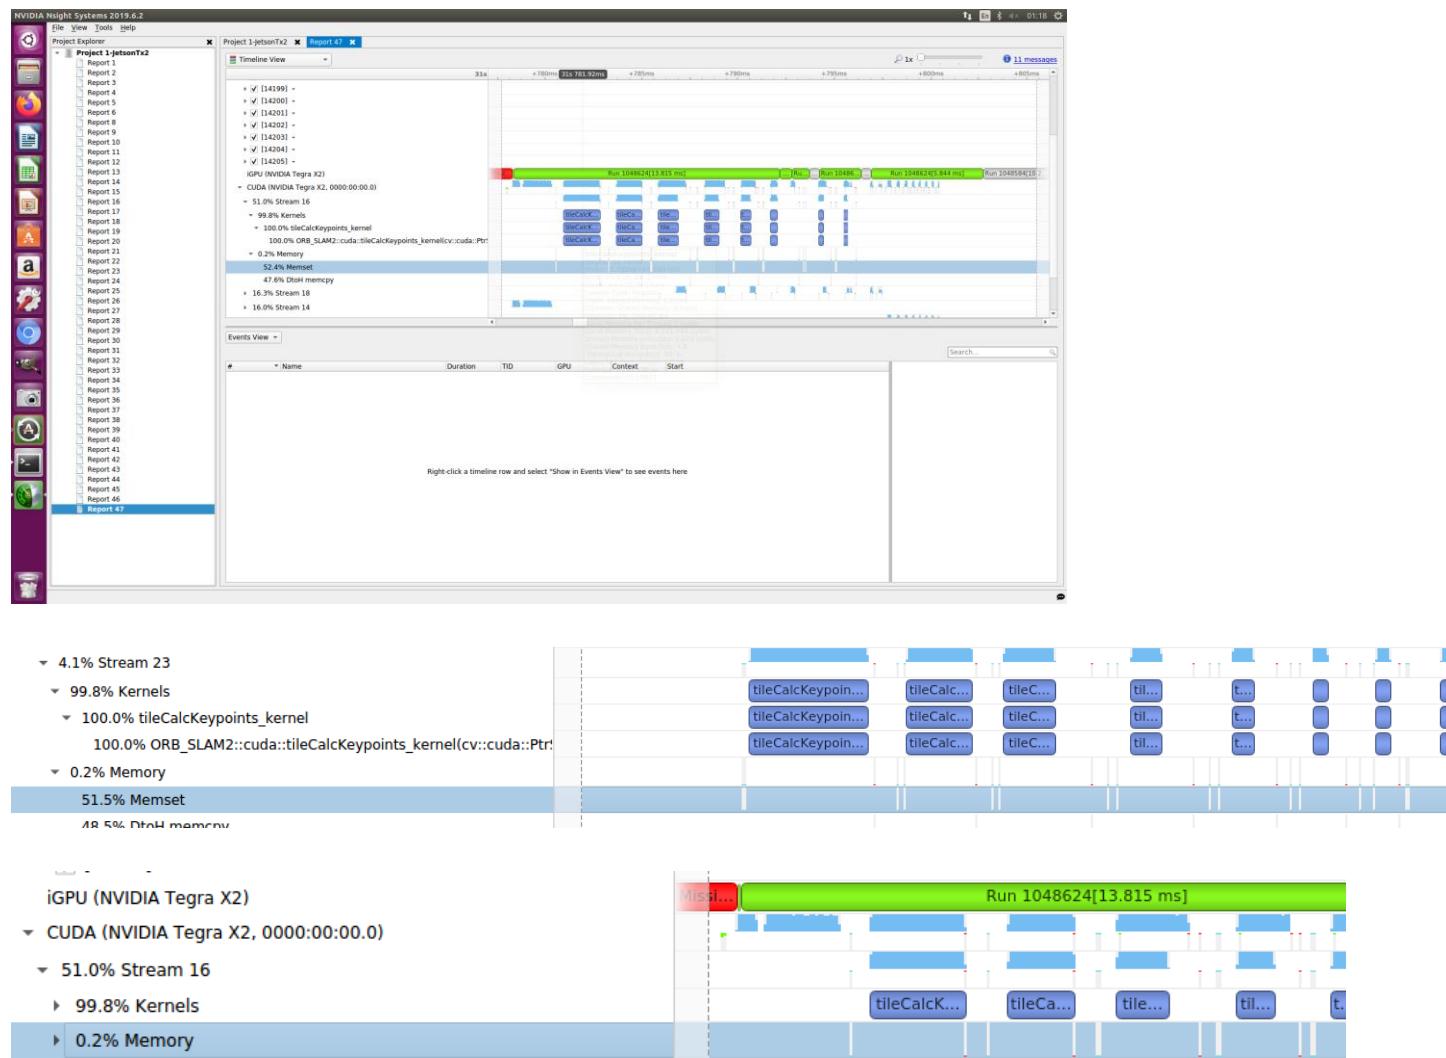
Analyze the SLAM CUDA application with Nsight system

Cuda Streams :

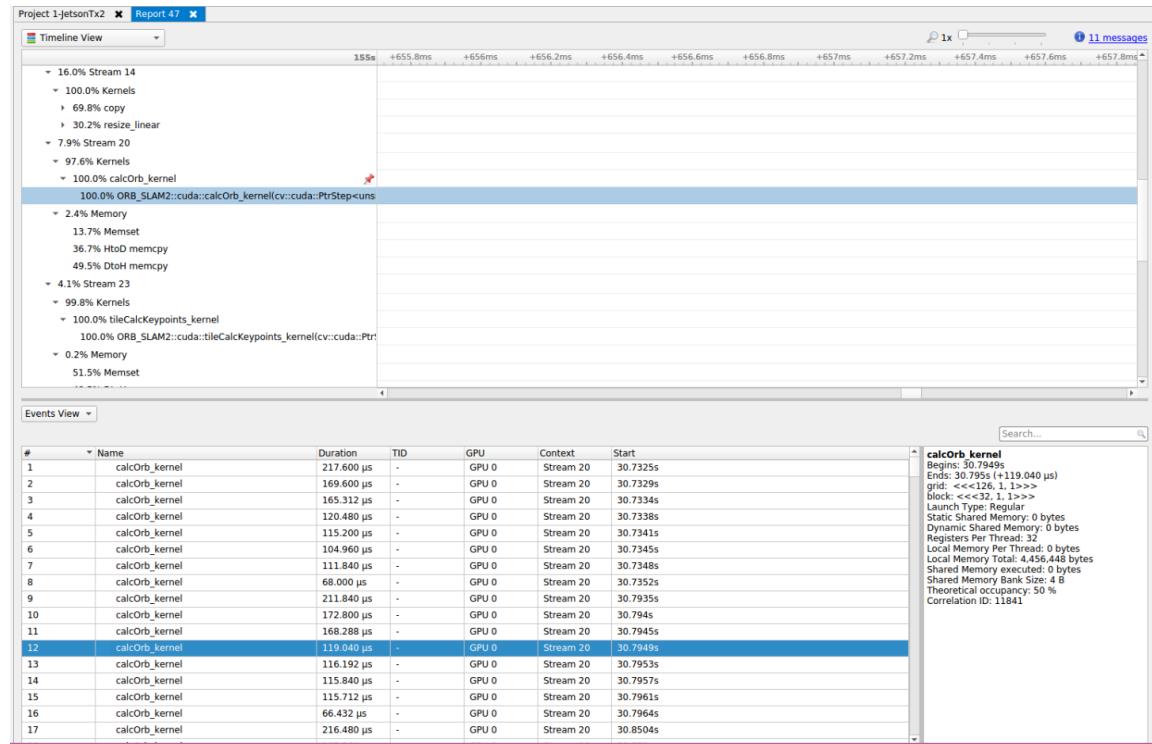
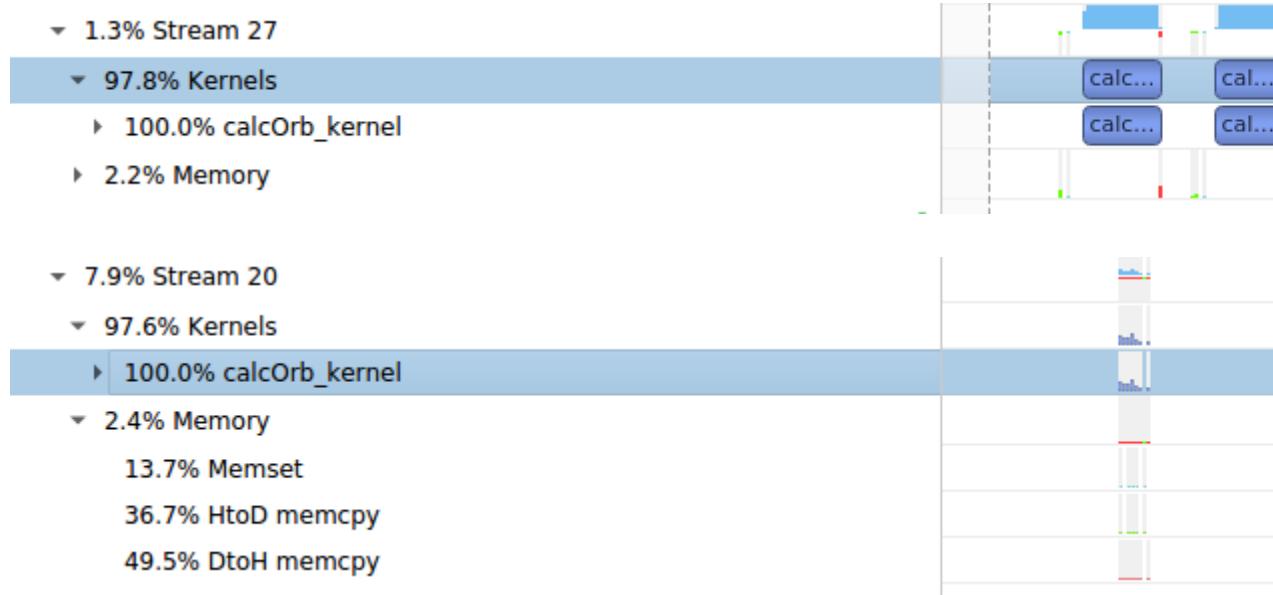
Most of the time run kernels:

Example:

`tileCalcKeypoints_kernel()`

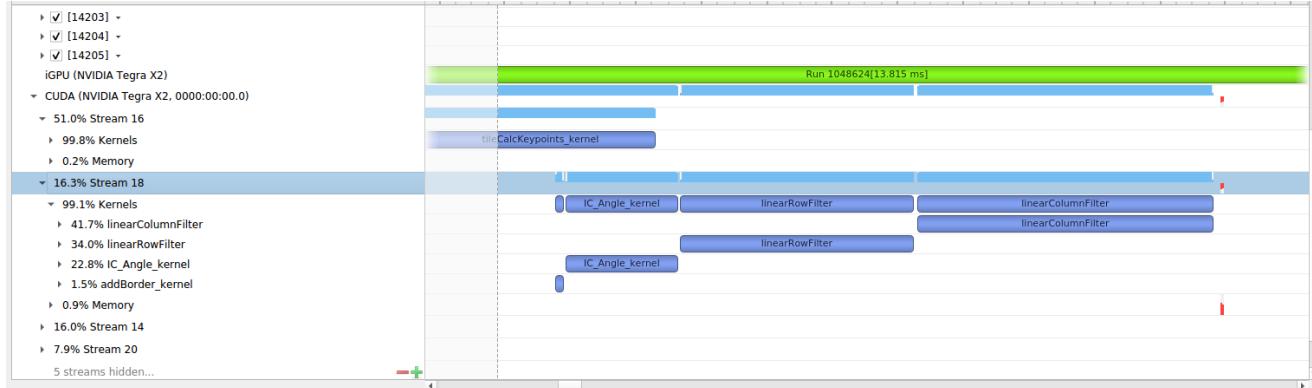


CalcOrb_kernel()

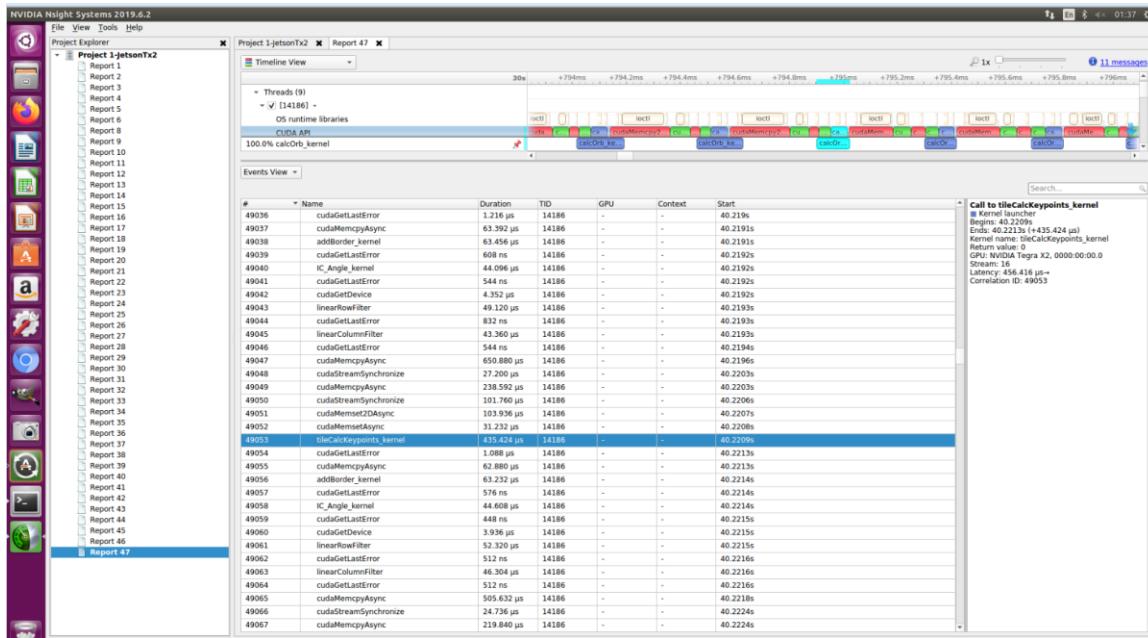


Another example:

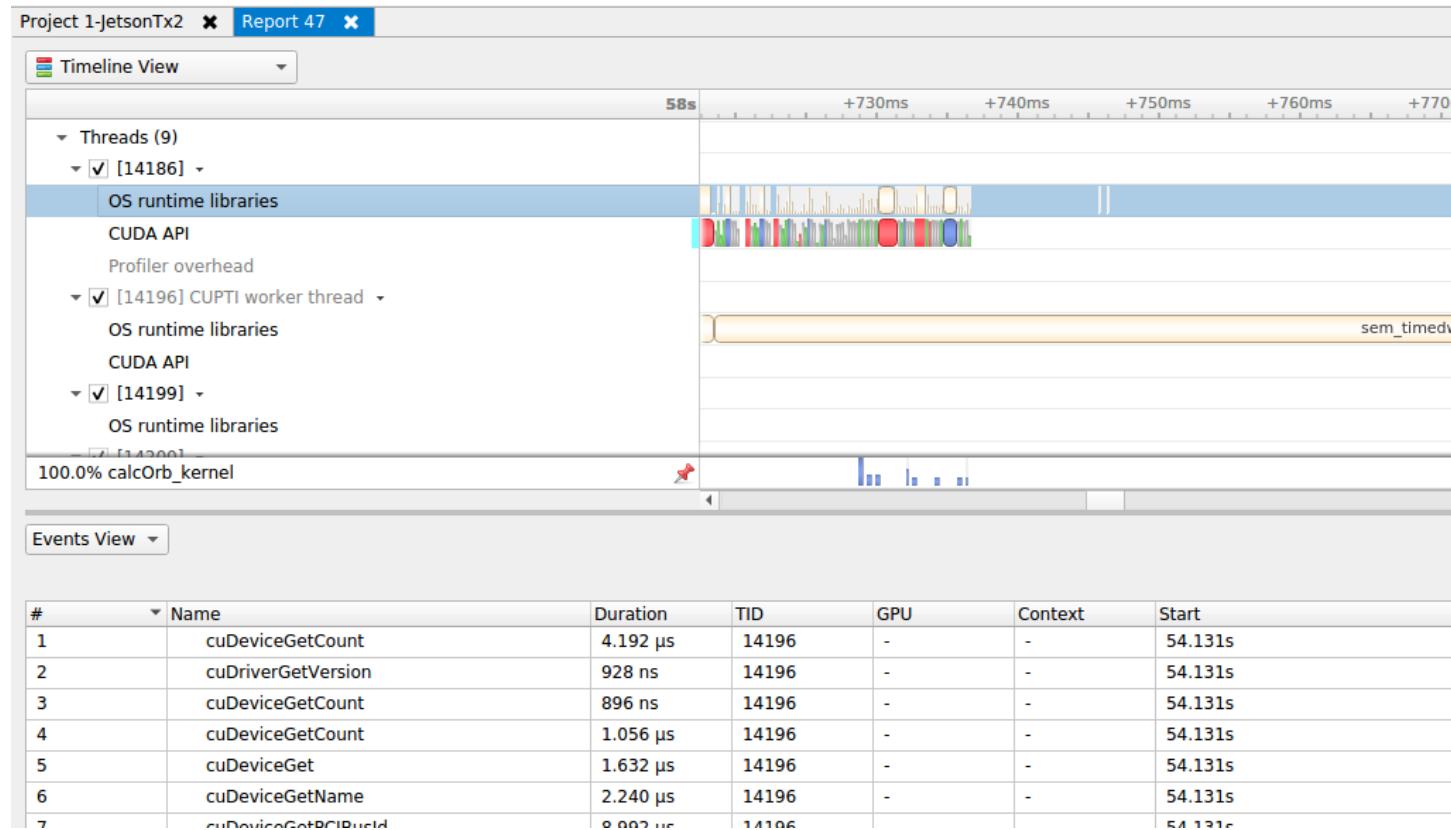
IC_Angle_kernel()



You can see timeline for the CUDA function and Cuda API:



You can see OS runtime libraries usage:



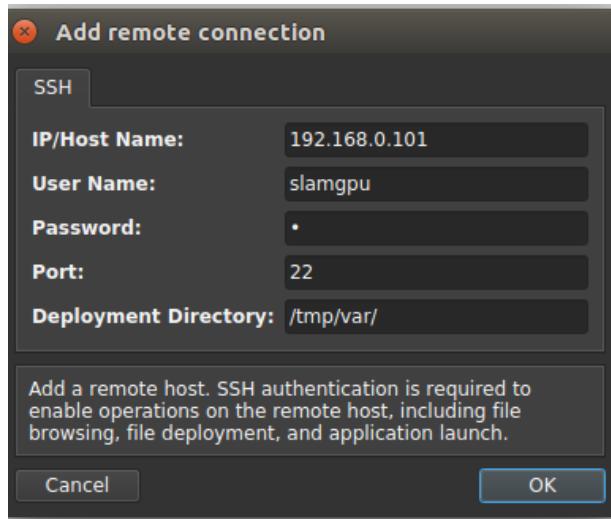
Nsight Compute

SLAM application in Standalone mode

Start the application with the following command

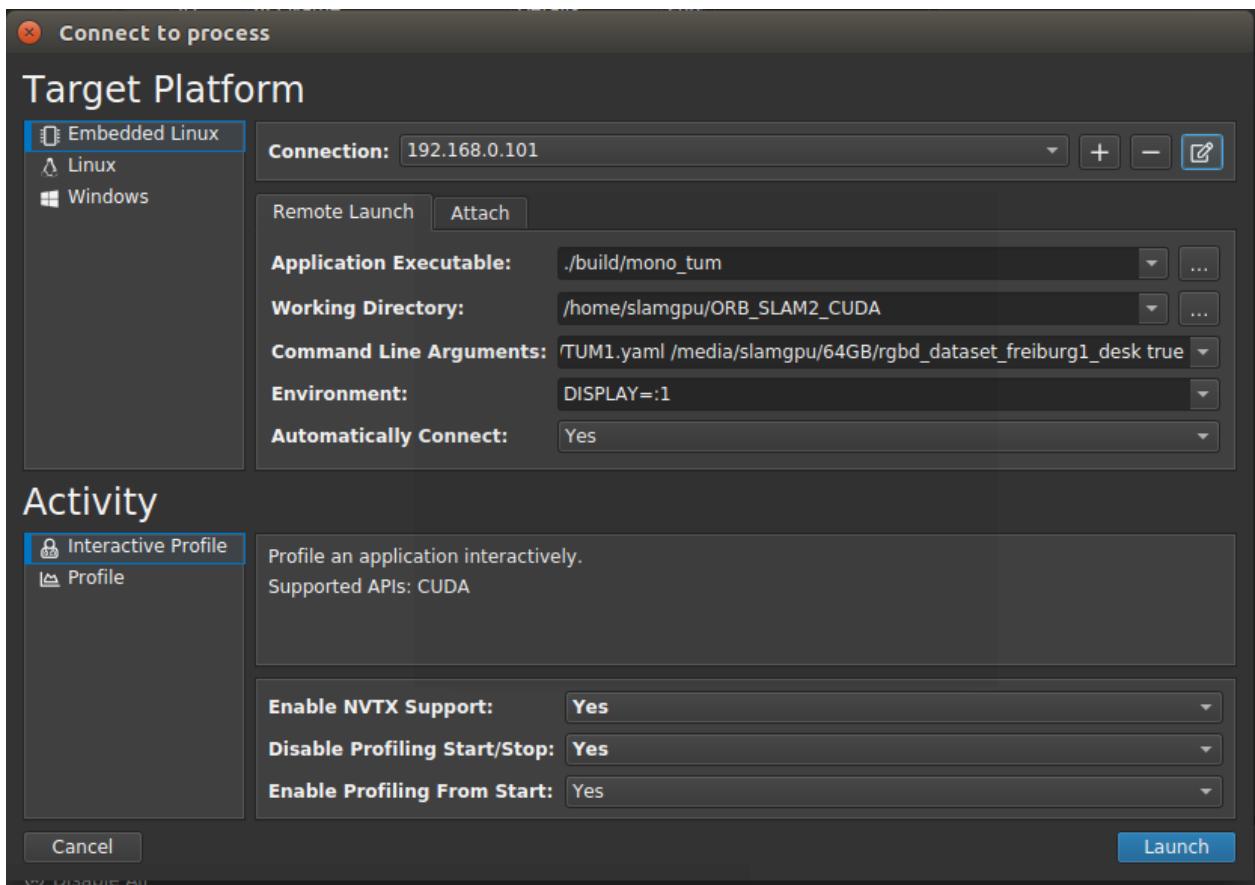
```
/usr/local/cuda-10.0/NsightCompute-1.0/nv-nsgith-cu
```

Configure connection to Jetson TX2 :



Recommendation use root user instead slamgpu user

Configure the application parameters:



Application:

./build/mono_tum

Working Directroy:

/home/slamgpu/ORB_SLAM2_CUDA

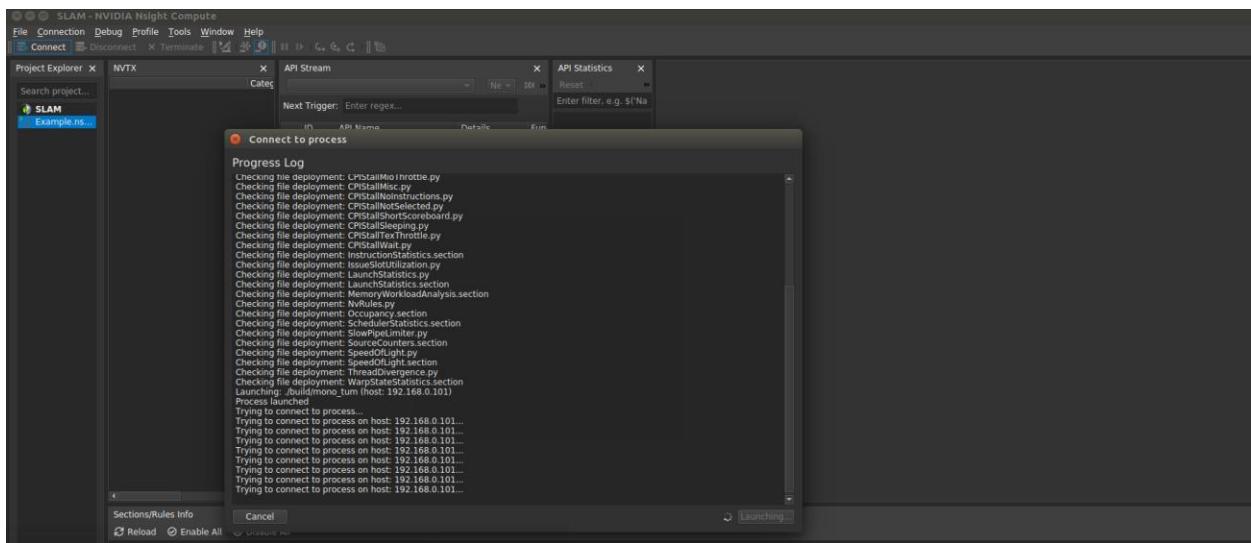
Application parameters:

Vocabulary/ORBvoc.txt Examples/Monocular/TUM1.yaml /media/slamgpu/64GB/rgbd_dataset_freiburg1_desk true

enviroment

DISPLAY=:1

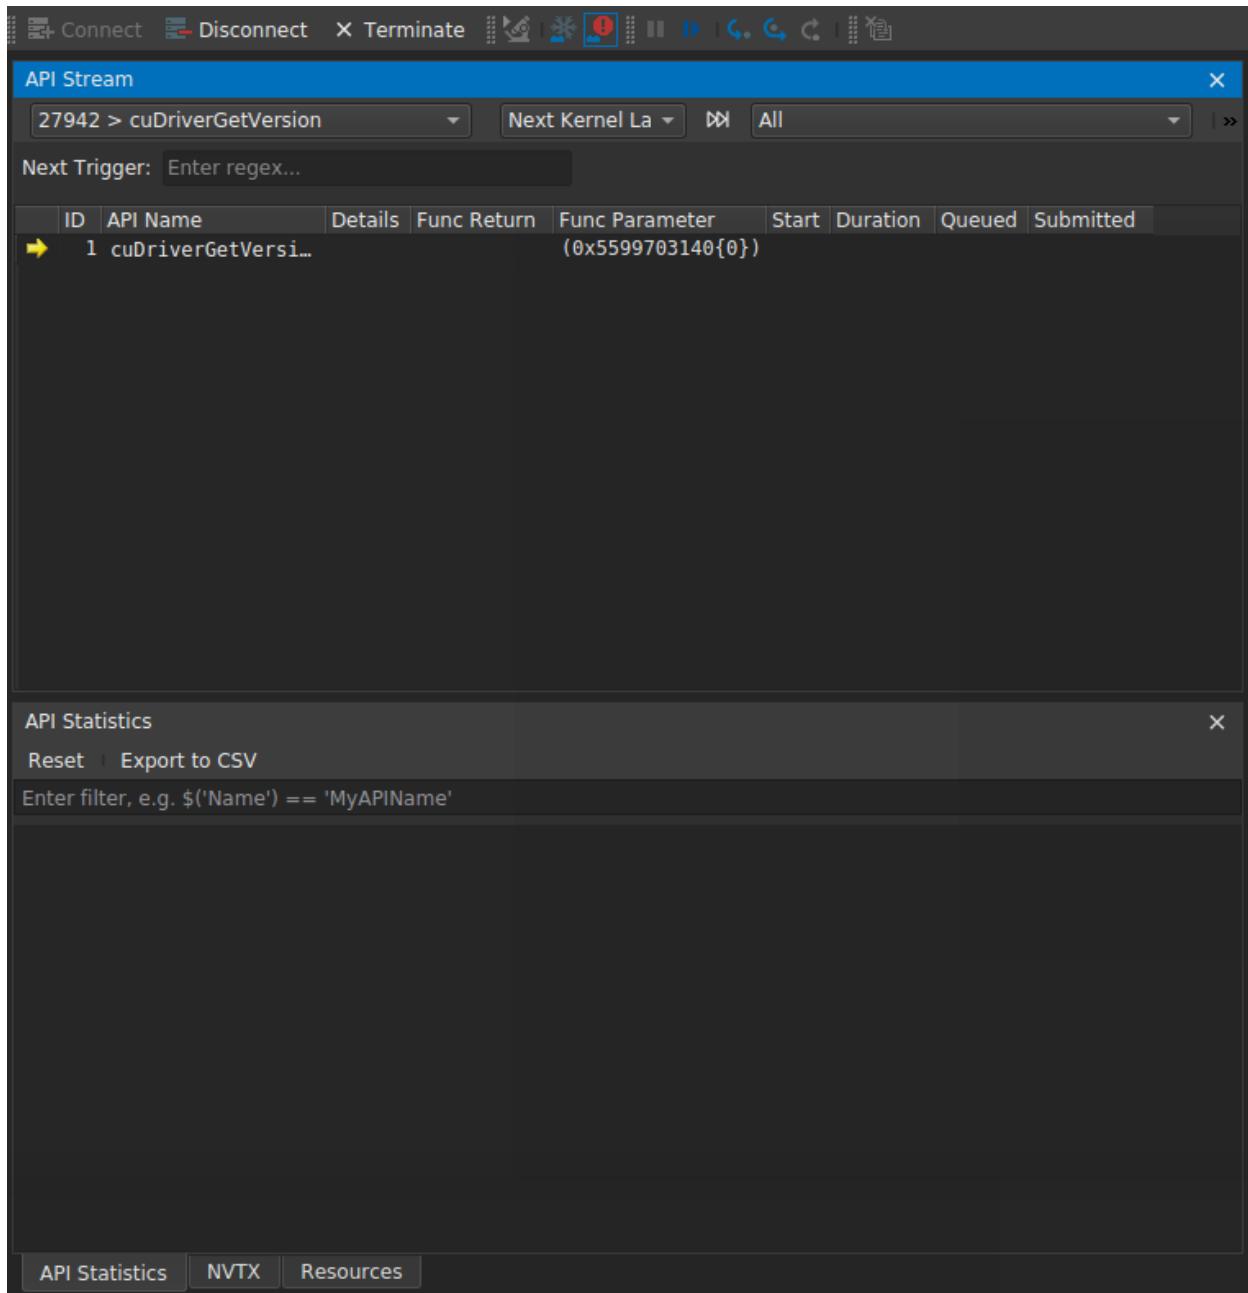
Use User root for interactive mode and click Launch to start



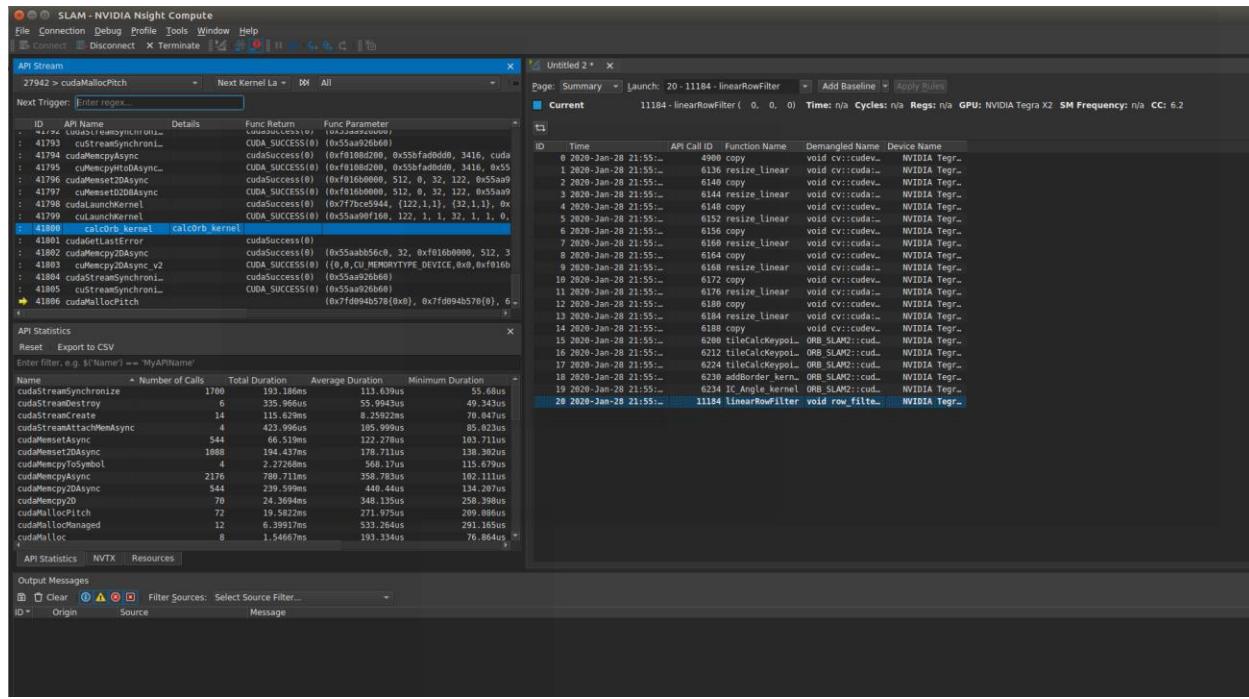
Waiting to application loading

When finished you start debug the application and jump between kernel and Cuda API.

In the following image you can see the first call (Entry point) to Cuda API call.



You can see the functions calls



You can debug source code

Untitled 2 * X

Page: Source Launch: 20 - 11184 - linearRowFilter Add Baseline Apply Rules

Current 11184 - linearRowFilter (0, 0, 0) Time: n/a Cycles: n/a Regs: n/a GPU: NVIDIA Tegra X2 SM Frequency: n/a CC: 6.2

View: SASS

_ZN10row_filter15linearRowFilterLi32EhfN2cv4cuda6device16BrdRowReflect101lhEEEEvNS2_9PtrStepSzIT0_EENS2_7PtrStepIT1_EEPKfIT2_

#	Address	Source	Live Registers
1		<u>_ZN10row_filter15linearRow</u>	
2	0207b7c8	MOV R1, c[0x0][0x20]	1
3	0207b7d0	S2R R12, SR_CTAID.Y	2
4	0207b7d8	S2R R19, SR_TID.Y	3
5	0207b7e8	ISCADD R12, R12, R19, 0x3	3
6	0207b7f0	ISETP.GE.AND P0, PT, R12, c	3
7	0207b7f8	NOP	3
8	0207b808	NOP	3
9	0207b810 @P0	EXIT	3
10	0207b818	SHR R13, R12, 0x1f	4
11	0207b828	S2R R3, SR_CTAID.X	5
12	0207b830	S2R R10, SR_TID.X	6
13	0207b838	ISETP.NE.AND P0, PT, R3, RZ	6
14	0207b848	ISCADD R0, R3, R10, 0x7	7
15	0207b850	IADD32I R4, R0, -0x20	8
16	0207b858 @P0	BRA 0x207b9b8	8
17	0207b868	MOV R2, c[0x0][0x174]	9
18	0207b870	I2I.S32.S32 R11, R4	10
19	0207b878	IADD32I R2, R2, 0x1	10
20	0207b888	I2F.F32.S32.RP R5, R2	11
21	0207b890	I2I.S32.S32 R9, R2	12
22	0207b898	ISETP.EQ.AND P1, PT, R2, RZ	12
23	0207b8a8	MUFU.RCP R5, R5	12
24	0207b8b0	IADD32I R6, R5, 0xffffffff	13
25	0207b8b8	F2I.FTZ.U32.F32.TRUNC R6, R6	12
26	0207b8c8	XMAD R7, R9, R6, RZ	13
27	0207b8d0	XMAD.MRG R8, R9, R6.H1, RZ	14
28	0207b8d8	XMAD.PSL.CBCC R7, R9.H1, R8.	14
29	0207b8e8	IADD R16, -R7, RZ	14
30	0207b8f0	XMAD R7, R6, R16, RZ	14
31	0207b8f8	XMAD R8, R6, R16.H1, RZ	15
32	0207b908	XMAD R5, R6.H1, R16.H1, R6	16

You can see information about the session

Untitled 2 * X

Page: Session Launch: 6200 - tileCalcKeypoints_kernel Add Baseline Apply Rules

Current 6200 - tileCalcKeypoints_kernel (0, 0, 0) Time: n/a Cycles: n/a Regs: n/a GPU: NVIDIA Tegra X2 SM Frequency: n/a CC: 6.2

Launch Settings

Executable	./build/mono_tum
Working Directory	/home/slampu/ORB_SLAM2_CUDA
Command Line Arguments	Vocabulary/ORBvoc.txt Examples/Monocular/TUM1.yaml /media/slampu/64GB/rgbd_dataset_freiburg1_desk true
Environment	DISPLAY=:1
Activity Type	Interactive Profile
Enable NVTX Support	Yes
Disable Profiling Start/Stop	Yes
Enable Profiling From Start	Yes

Session Info

Created	2020-Jan-28 21:55:28
Host Name	yoni3-GB-BRi7-H-8550
Host OS	Linux - #81~16.04.1-Ubuntu SMP Tue Nov 26 16:34:21 UTC 2019
Host Architecture	x86_64
Host Processor	Intel(R) Core(TM) i7-8550U CPU @ 1.80GHz
Target Name	slampu-desktop
Target OS	Linux - #1 SMP PREEMPT Mon Aug 12 21:29:52 PDT 2019
Target Architecture	aarch64
CUDA Version	10.0
NVIDIA Nsight Compute	1.0 (Build 25942742)

Device Attributes

Attribute	NVIDIA Tegra X2 (0)
architecture	304
async_engine_count	1
can_flush_remote_writes	0
can_map_host_memory	1
can_tex2d_gather	1
can_use_64_bit_stream_mem_ops	0
can_use_host_pointer_for_registered_mem	0
can_use_stream_mem_ops	0
can_use_stream_wait_value_nor	0
chip	315
clock_rate_mhz	1020000

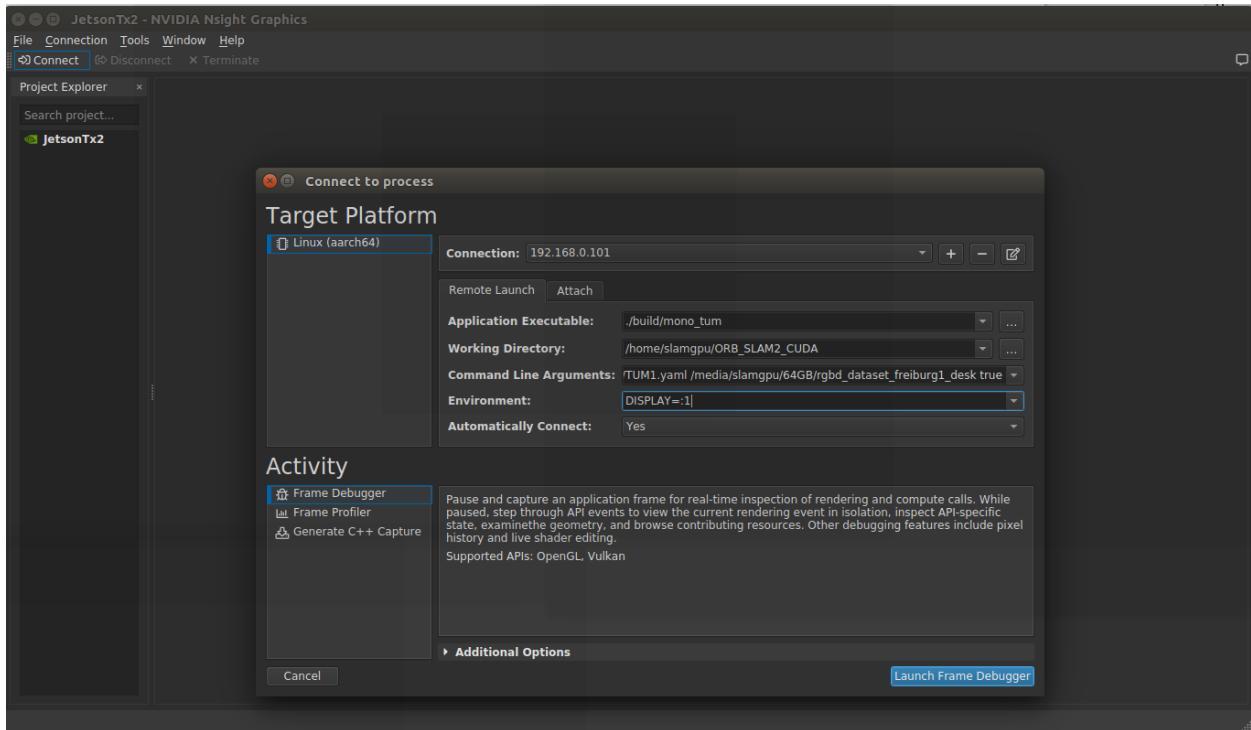
Nsight Graphics

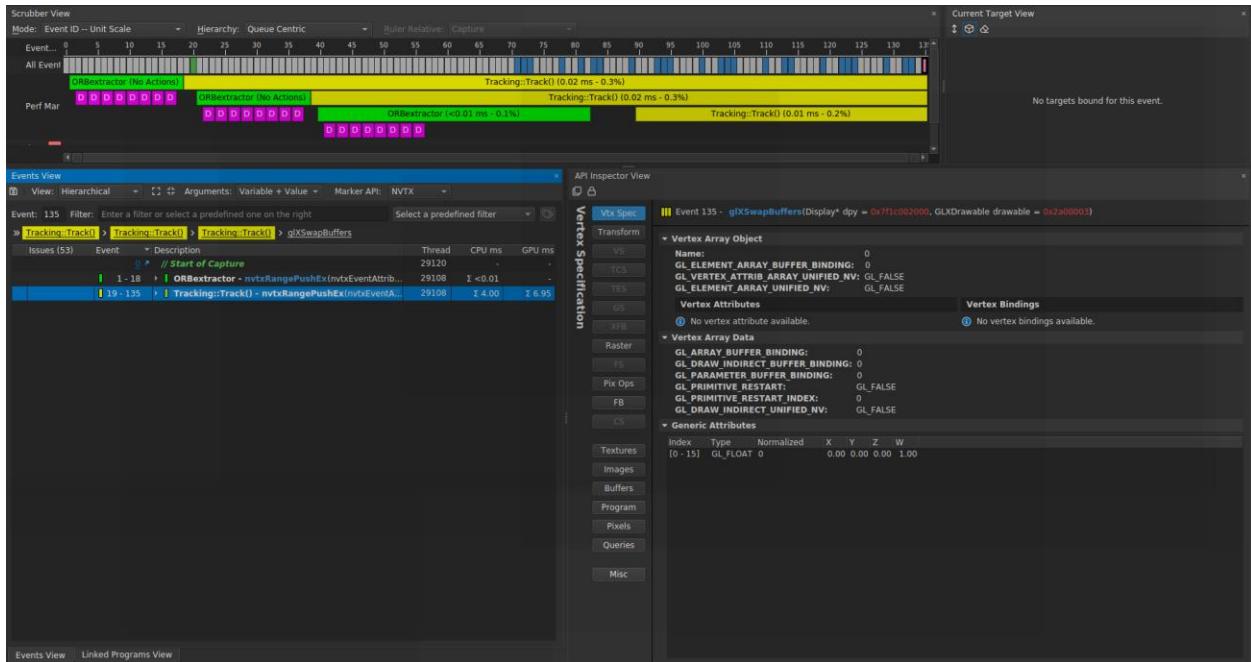
SLAM application in Standalone mode

You can launch the tool by command:

```
nv-nsight-gfx-for-l4t
```

Configure the connection to application





Extensions

Enable root login over SSH

1. Login to your server as root.
2. As the root user, edit the `sshd_config` file found in `/etc/ssh/sshd_config`:
`vim /etc/ssh/sshd_config` (*For details on working with Vim check out our !*)
3. Add the following line to the file, you can add it anywhere but it's good practice to find the block about authentication and add it there.

```
PermitRootLogin yes
```

4. Save and exit the file.

5. Restart the SSH server:

```
systemctl restart sshd
```